Google 288
Algorithm Primer

Read this in three days, before you learn the code solutions

---

43 patterns · 228 questions · 738 follow-ups

Every algorithm used by the easy, medium and hard Google study path, explained once, with a short template and the traps to avoid.

How to use it. Each pattern page says how to spot it, the idea in plain words, a short template with clear names, the cost, the traps, the follow-ups to expect, and the list of study-path questions that use it, with one line telling you what the algorithm does there. Some of them are starred, 150 in the whole book, chosen by how recently each was reported to be asked: learn those first, in the star order. Read a pattern, close the book, and write the template from memory. Then the long solutions in the app will read as “this is pattern 24 with a twist” instead of a wall of code.

★ A green star means every starred question of that pattern has a solution on PracHub or HackTheRounds, so you can check your answer against it. learn N is how many questions of that pattern to learn (150 in all); they carry big orange stars in the book.

Contents

The three day plan 4

Priority order: what Google asked most recently 4

Python toolbox and complexity 5

Follow-up families that repeat everywhere 6

Pattern finder: from the words of the question to the section 6

Patterns, most questions to learn first

☐ 15. Dict of per-key state (update, expire, rank) learn 7 19

☐ 16. Sweep over events in time order ★ learn 7 19

☐ 32. Trie: nested dicts, one character per level ★ learn 7 30

☐ 1. Union-find: merge the pairs, then ask who is together ★ learn 6 9

☐ 3. A heap of what is available: pop the best, put back the changed one ★ learn 6 10

☐ 4. Backtracking: choose, recurse, undo ★ learn 6 11

☐ 12. Recurrence on a position or a range (memo or table) ★ learn 6 16

☐ 5. Best-first search with a heap: Dijkstra and its minimax twin ★ learn 5 11

☐ 19. Greedy: scan once, carry the best so far, act only when forced ★ learn 5 21

☐ 21. Prefix sums, tier 1: precompute once, query in O(1) learn 5 22

☐ 27. Group by a canonical key learn 5 26

☐ 29. Nested structures: a stack of contexts, or recursion with a cache ★ learn 5 29

☐ 2. Frontier search (BFS or DFS), tier 1: spread, count, measure distance ★ learn 4 9

☐ 17. Binary search on a yes/no condition ★ learn 4 20

☐ 18. Sorted intervals: merge while they touch, then read gaps or search ★ learn 4 21

☐ 25. Sliding window: grow on the right, shrink on the left learn 4 25

☐ 30. Tree recursion, tier 1: visit a node, ask children, report upward ★ learn 4 29

☐ 40. Parent pointers: store who is above each node and climb ★ learn 4 35

☐ 8. Top-K by frequency: Counter plus nsmallest ★ learn 3 13

☐ 11. Rolling state: the next state comes only from the previous one learn 3 15

☐ 13. Grid DP: each cell from its neighbours above and to the left ★ learn 3 16

☐ 14. Counting or reaching an amount: the knapsack loop learn 3 17

☐ 20. Count pairs by looking up each partner learn 3 22

☐ 23. bisect on a sorted list: position, range count, keep sorted ★ learn 3 23

☐ 24. Sets and counters across two collections learn 3 24

☐ 26. Two pointers moving toward each other (or along two lists) learn 3 25

☐ 33. A dict for lookup plus an ordered container for who is next ★ learn 3 31

☐ 36. Simulate a machine: a class with state and one branch per event ★ learn 3 33

☐ 37. Digits and characters: peel one at a time and carry ★ learn 3 33

☐ 38. Split into pieces, fix each piece, join (and word wrap) ★ learn 3 34

☐ 39. Walk two strings at once with two indices learn 3 35

☐ 6. Topological sort (Kahn) and cycle detection ★ learn 2 12

☐ 7. Frontier search, tier 2: free steps, early exit, neighbours worked out on the fly ★ learn 2 13

☐ 22. Runs and groups: extend while it fits, restart when it breaks learn 2 23

☐ 28. Prefix sums, tier 2: running total plus a dict of earlier totals ★ learn 2 26

☐ 31. Tree DP: build a parent order once, add up from the leaves learn 2 30
☐ 9. Games: win or lose with a cache, and guessing with feedback ★ learn 1 14
☐ 10. K-way merge: a heap of the next item from each source ★ learn 1 14
☐ 34. Build or transform a grid with index arithmetic ★ learn 1 32
☐ 35. Look at the input once and keep only a small summary ★ learn 1 32
☐ 41. Random pick without repeats: swap with the last and pop ★ learn 1 36
☐ 42. Tree recursion, tier 2: pass a rule down, return a node or yes/no up ★ learn 1 36
☐ 43. Build and serialize: write out a structure, read it back with the same walk ★ learn 1 37

Appendix: algorithms hiding inside single questions 38

Coverage by section and source 40

After the 150 starred questions: what to do next 44

Behavioural questions: sample answers from your resume 45

The three day plan

Budget: 5 to 6 hours a day. Read a pattern for 10 minutes, then write its template from memory for 5 minutes. Do not copy; re-read only what you could not rebuild. Anything you cannot rebuild goes on a list for the end of the day.

Day 1: Recursion, dynamic programming, heaps and graphs. Backtracking, games, dynamic programming, heaps, top-k, BFS and DFS, topological sort, union-find and Dijkstra. Patterns 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14.

By the end of the day you can: Write the choose, recurse, undo skeleton and a memoised state function. Write BFS, Kahn, union-find and Dijkstra with a heap from memory. For any new problem, name the state of a DP or the frontier of a search before coding.

Day 2: Counting, scanning and searching. Hash maps, prefix sums, two pointers, windows, greedy, binary search, sorted intervals and sweeps. Patterns 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28.

By the end of the day you can: Explain in one sentence why a dict lookup replaces an inner loop. Write a prefix sum, a sliding window and a binary search on an answer from memory. Merge intervals and run a sweep over events without looking.

Day 3: Strings, design and trees. Strings and digits, grids, simulations and caches, parent pointers, random structures, nested parsing, streaming, trees, serialisation and tries. Patterns 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43.

By the end of the day you can: Write an LRU cache, a trie and a swap-and-pop random set from memory. State when a tree function returns a value and when it updates a global answer. Describe the rerooting formula and a length-prefixed encoding.

Day 4 (optional): open the app, pick one question per pattern, and write the main solution without notes. If you stall, the stall is the pattern you need to reread, not the question.

Priority order: what Google asked most recently

The days follow the order most worth your time: Day 1 holds the patterns with the highest study-path score (recursion, dynamic programming, heaps and graphs), Day 2 the patterns with the most recent reports (counting, scanning, searching, intervals), Day 3 strings, design and trees. Day 1 leans on small tools from Day 2 (dicts, sorting, deque, binary search), so read the toolbox chapter first and look up any tool you do not know in the Day 2 patterns as you meet it. Inside each day the patterns are ordered by that same ranking (most reported first), except that a pattern is never placed before the one it builds on (tier 1 before tier 2, BFS before Dijkstra). If time is short, ignore the day order and study patterns in the order of the table below. Rank uses questions reported in the last 90 days first, then all dated questions, then the study-path score.

Rank	#	App	Pattern	Last 90 days	Dated	Questions	Day
1	15	10	Dict of per-key state (update, expire, rank)	4	7	8	Day 2
2	29	27	Nested structures	4	5	5	Day 3
3	1	9	Union-find	3	7	7	Day 1
4	13	26	Grid DP	3	4	5	Day 1
5	16	2	Sweep over events in time order	2	8	9	Day 2
6	2	1	Frontier search (BFS or DFS), tier 1	2	6	6	Day 1
7	3	4	A heap of what is available	2	6	8	Day 1
8	17	11	Binary search on a yes/no condition	2	5	7	Day 2
9	30	3	Tree recursion, tier 1	2	5	8	Day 3
10	4	7	Backtracking	2	5	7	Day 1
11	18	16	Sorted intervals	2	4	6	Day 2
12	19	12	Greedy	2	4	5	Day 2
13	31	31	Tree DP	2	4	5	Day 3
14	20	30	Count pairs by looking up each partner	2	2	3	Day 2
15	32	17	Trie	1	8	8	Day 3
16	12	5	Recurrence on a position or a range (memo or table)	1	5	8	Day 1
17	5	6	Best-first search with a heap	1	5	8	Day 1
18	21	15	Prefix sums, tier 1	1	4	6	Day 2
19	33	18	A dict for lookup plus an ordered container for who is next	1	4	5	Day 3
20	22	8	Runs and groups	1	4	4	Day 2
21	23	14	bisect on a sorted list	1	3	6	Day 2
22	6	29	Topological sort (Kahn) and cycle detection	1	3	4	Day 1
23	7	28	Frontier search, tier 2	1	3	4	Day 1
24	24	13	Sets and counters across two collections	1	3	4	Day 2
25	25	40	Sliding window	1	2	4	Day 2
26	34	42	Build or transform a grid with index arithmetic	1	2	3	Day 3

Rank	#	App	Pattern	Last 90 days	Dated	Questions	Day
27	35	25	Look at the input once and keep only a small summary	0	6	6	Day 3
28	8	23	Top-K by frequency	0	5	5	Day 1
29	9	39	Games	0	5	5	Day 1
30	36	22	Simulate a machine	0	5	5	Day 3
31	37	24	Digits and characters	0	4	4	Day 3
32	38	43	Split into pieces, fix each piece, join (and word wrap)	0	4	5	Day 3
33	39	21	Walk two strings at once with two indices	0	3	6	Day 3
34	26	20	Two pointers moving toward each other (or along two lists)	0	2	6	Day 2
35	10	35	K-way merge	0	2	4	Day 1
36	40	41	Parent pointers	0	2	4	Day 3
37	41	37	Random pick without repeats	0	2	2	Day 3
38	11	19	Rolling state	0	1	5	Day 1
39	27	34	Group by a canonical key	0	1	5	Day 2
40	28	33	Prefix sums, tier 2	0	1	3	Day 2
41	42	36	Tree recursion, tier 2	0	1	3	Day 3
42	14	32	Counting or reaching an amount	0	0	4	Day 1
43	43	38	Build and serialize	0	0	3	Day 3

Python toolbox and complexity

Pick the target complexity from the input size. Interview inputs are chosen so one family of algorithms fits.

n up to	Target	Typical tools
20 to 25	$O(2^n)$, $O(n!)$	backtracking, bitmask DP
1,000	$O(n^2)$	table DP, all pairs, nested loops
100,000	$O(n \log n)$	sort, heap, binary search, balanced structure
1,000,000 and up	$O(n)$ or $O(\log n)$	one pass, hashing, two pointers, math

Costs of the tools you will use all day.

Structure	Fast	Slow (avoid in loops)
`list`	append, pop, index $O(1)$	`insert(0)`, `pop(0)`, `x in list`: $O(n)$
`dict`, `set`	get, set, delete, in: $O(1)$ average	iteration order is insertion order only
`deque`	append, popleft, appendleft $O(1)$	indexing in the middle $O(n)$
`heapq`	push, pop $O(\log n)$; peek `h[0]` $O(1)$	search or delete a middle item $O(n)$
sorted list + `bisect`	search $O(\log n)$	`insort` $O(n)$ (shifts items)
`str`	index, slice $O(k)$	`s += t` in a loop: build a list and join

Standard library you are expected to know.

- `Counter(items)`: counts; `most_common(k)`; subtract with `c[x] -= 1`. `defaultdict(list)` or `defaultdict(int)` removes the "key exists?" check.
- `deque`: BFS queue, sliding window, 0-1 BFS. `OrderedDict.move_to_end` and `popitem(last=False)` give an LRU in a few lines.
- `heapq`: `heappush`, `heappop`, `heapify`, `heapreplace`, `nsmallest(k, xs, key=...)`. It is a min-heap: push `-x` for a max-heap. Push tuples `(priority, tiebreak, item)` so items never get compared.
- `bisect_left(a, x)` first index with `a[i] >= x`; `bisect_right(a, x)` first index with `a[i] > x`; `insort`.
- `lru_cache(maxsize=None)` from functools memoises a function of hashable arguments (use tuples, not lists). From itertools: `combinations`, `permutations`, `accumulate` (prefix sums). `zip(*matrix)` transposes.
- `sorted(xs, key=lambda x: (-count[x], x))`: sort by several criteria; sort is stable. Bits: `x & (x-1)` clears the lowest set bit, `1 << i` is bit i, `x >> i & 1` reads it.
- Recursion depth is about 1000 in Python: for deep trees or chains use an explicit stack or a queue (BFS), or raise the recursion limit and say why it is safe.

The loop to run on every question. (1) Restate it and ask about duplicates, ties, empty input, negative numbers, the size of the input. (2) Say a brute force and its cost. (3) Name the pattern in this book and the state or data structure it needs. (4) State the complexity before coding. (5) Code with clear names. (6) Test on an empty case, a one item case, a tie, and a case that breaks your first idea. (7) Be ready for the follow-ups listed with each pattern.

Follow-up families that repeat everywhere

Across the questions in this book the same follow-ups return in different clothes. Learn one answer to each family and you can answer most of them.

The input is a stream or does not fit in memory. Keep a small summary instead of the data: a running max, a size-k heap, a count dict, a set of ids from the smaller file. For two huge inputs partition both by a hash of the key and process one partition at a time. Many machines: map (emit key and value), shuffle by key, reduce.

Make it thread safe. One lock around each method is correct and simple; finer locks (per key or per bucket) give more concurrency but need a lock order to avoid deadlock. Read-mostly data can use copy-on-write. Say what is atomic.

Items can also be deleted or updated. A plain sorted list or prefix sum breaks: move to a Fenwick tree or a segment tree, a sorted container with $O(\log n)$ updates, lazy deletion in a heap, or swap-and-pop in a list with a position dict. Union-find cannot split groups: process offline in reverse or rebuild.

Return the path or the items, not only the number. Store a parent or choice for each state when you first improve it, then walk back from the goal: `parent[v] = u` in BFS and Dijkstra, the chosen split in DP, the chosen index in a greedy scan.

Weights, costs or priorities are added. BFS becomes Dijkstra (non-negative weights, or 0-1 BFS for 0 and 1), greedy may stop being correct (re-check the exchange argument), plain counts become sums stored in the same tables.

Recursion is too deep or too slow. Convert to an explicit stack or a queue; add a cache to turn exponential into polynomial; compute subtree values by sweeping a parent-first order backwards.

Do it in place, or in $O(1)$ extra space. Reuse the input as the table (swap each value into its own slot, encode old and new state in different bits), or use two pointers instead of a set after sorting.

Online arrivals, many queries, repeated calls. Precompute once (prefix sums, adjacency, a trie, a sorted list) and make each query cheap; for arrivals that cannot be sorted ahead use a heap or a balanced structure.

The input is huge numerically. Use modular arithmetic, matrix exponentiation or fast doubling for recurrences, a closed form for counting, or strings for numbers beyond machine size.

Prove or justify. Greedy: name the exchange argument and a counter-example for the neighbouring problem. Binary search: show the predicate is monotone. DP: state the meaning of the state. Dijkstra: weights are non-negative so a popped node is final.

Pattern finder: from the words of the question to the section

Read the question, match the words, open the section. The first number is the pattern's place in this book; App is its section number in the study workspace.

#	App	Pattern	Words that give it away
1	9	Union-find	connected components, groups, are A and B related, add land one cell at a time, points within a radius.
2	1	Frontier search (BFS or DFS), tier 1	number of islands, flood fill, rotten oranges, shortest path in an unweighted graph, distance to nearest X.
3	4	A heap of what is available	minimum rooms, streaming median, reusable locker, schedule on identical CPUs, no two adjacent the same.
4	7	Backtracking	all permutations / combinations, generate parentheses, count paths of length k, can this hand win, assign people to tasks.
5	6	Best-first search with a heap	shortest path with weights, earliest arrival with timetables, minimise the maximum height on a path, swim in rising water.
6	29	Topological sort (Kahn) and cycle detection	course schedule, build order, are these race results consistent, validate dependencies.
7	28	Frontier search, tier 2	edge direction flips cost 1, reach with blocked nodes, shortest directed cycle through v, radios within a radius.
8	23	Top-K by frequency	k most frequent words, top chatters, task scheduler with cooldown, bottom k.
9	39	Games	who wins with optimal play, merge stacks game, guess the secret word, Wordle strategy.
10	35	K-way merge	merge k sorted lists, k smallest pair sums, sort a nearly sorted array, top-k subset sums.
11	19	Rolling state	climbing stairs, house robber, paint houses, fibonacci, dice with limits.
12	5	Recurrence on a position or a range (memo or table)	edit distance, longest palindromic subsequence, word break, split into expressions, count decompositions.
13	26	Grid DP	unique paths, largest square, longest run of one colour, paths with checkpoints.

#	App	Pattern	Words that give it away
14	32	Counting or reaching an amount	coin change, partition equal subset sum, target sum with plus and minus, unbounded items that fill a capacity.
15	10	Dict of per-key state (update, expire, rank)	remember the last time we saw X, top users, within the last 10 seconds, design a logger / counter.
16	2	Sweep over events in time order	minimum rooms, maximum overlap, skyline, who is online at time t, servers needed.
17	11	Binary search on a yes/no condition	minimum speed / capacity / time such that ..., largest length such that ..., first bad version, median of two sorted arrays.
18	16	Sorted intervals	merge intervals, do bookings conflict, free time, insert an interval, union of two lists.
19	12	Greedy	minimum number of ..., maximum profit, largest number you can form, cover all holes, jumps.
20	30	Count pairs by looking up each partner	pairs with difference k, two ways to write as a sum, rectangle / square from points.
21	15	Prefix sums, tier 1	sum of a range, balanced split, pivot index, equal halves, area up to a height.
22	8	Runs and groups	longest contiguous run, group consecutive items, streak, one replacement allowed.
23	14	bisect on a sorted list	count in a range, k-th largest as items stream in, nearest value, longest increasing subsequence.
24	13	Sets and counters across two collections	appears in both, not in the other list, duplicates, pair each x with 2x.
25	40	Sliding window	longest/shortest subarray or substring such that ..., at most k distinct, minimum window, budget of changes.
26	20	Two pointers moving toward each other (or along two lists)	palindrome, reverse, sorted array and a target pair, merge two sorted lists, triplets that sum to zero.
27	34	Group by a canonical key	anagrams, same shape, shifted words, group strings that are equivalent.
28	33	Prefix sums, tier 2	number of subarrays with sum k, multiple of k, maximum subarray, equal endpoints.
29	27	Nested structures	decode string 3[a2[c]], evaluate expression, nested list weight, expand placeholders.
30	3	Tree recursion, tier 1	height, diameter, right side view, level order, delete leaves, evaluate expression tree.
31	31	Tree DP	total size of a folder, sum of distances from every node, time to inform everyone, cheapest cut.
32	17	Trie	autocomplete, words with a prefix, longest match tokenizer, URL router, dictionary substrings.
33	18	A dict for lookup plus an ordered container for who is next	LRU cache, waitlist, round robin scheduler, remove the cheapest, undo / redo.
34	42	Build or transform a grid with index arithmetic	rotate a matrix, range sum in a 2D matrix, draw a diamond, fill in snake order, block minimum.
35	25	Look at the input once and keep only a small summary	file does not fit in memory, stream, minimise the number of oracle calls, two logs.
36	22	Simulate a machine	design tic-tac-toe, elevator system, iterator over lists, game of life, scorekeeper.
37	24	Digits and characters	add two big numbers, palindrome number, monotone digits, caesar cipher, k-th permutation.
38	43	Split into pieces, fix each piece, join (and word wrap)	snake_case to camelCase, format a number, justify text, wrap at width.
39	21	Walk two strings at once with two indices	is a subsequence, lexicographic order in a custom alphabet, can board A become board B, similar sentences.
40	41	Parent pointers	lowest common ancestor with parent links, validate a tree, employee hierarchy, is this a root.
41	37	Random pick without repeats	random set with O(1) delete, shuffle, unique random generator, no repeat within k songs.
42	36	Tree recursion, tier 2	valid BST, lowest common ancestor, in-order successor, mode in a tree.
43	38	Build and serialize	serialize / deserialize, construct tree from preorder and inorder, print a directory tree, encode a list of strings.

DAY 1

Recursion, dynamic programming, heaps and graphs

Patterns 1 to 14 (14 patterns)

Backtracking, games, dynamic programming, heaps, top-k, BFS and DFS, topological sort, union-find and Dijkstra.

1. Union-find: merge the pairs, then ask who is together ★

Union-find and connectivity (app section 9) 7 questions: 0 easy, 6 medium, 1 hard

Spot it: "connected components", "groups", "are A and B related", "add land one cell at a time", "points within a radius".

Idea: Each element points to a parent; the root is the group name. find(x) follows parents (flattening the chain as it goes), union(a, b) links the two roots. Pairs may come from edges, from neighbouring land cells, or from points within a radius (bucket points into cells of size r and test only the nine nearby cells). Offline path problems: sort edges from best to worst and union until two nodes meet; the last edge is the threshold.

```
parent = list(range(n))                    # items are 0..n-1 (5 of the 8 questions)
# parent = {}                               # names, tuples, cells that appear later:
#                                           # add parent.setdefault(x, x) as the first line of find
def find(x):
    while parent[x] != x:
        parent[x] = parent[parent[x]]      # path halving
        x = parent[x]
    return x
def union(a, b):
    ra, rb = find(a), find(b)
    if ra == rb: return False              # already together
    parent[ra] = rb;  return True

groups = n                                 # with the dict form: count only items you added
for a, b in pairs:
    if union(a, b): groups -= 1
```

Cost: Close to O(1) per operation in practice: the path halving in find keeps chains short, and no size bookkeeping is needed. **Traps:** Pick the style from the input: ids 0..n-1 use the list (start groups at n); names, tuples or items that appear later use the dict with setdefault (and `find` on an unseen item quietly creates it). Unfriending breaks plain union-find (use offline reverse processing or rebuild). Count groups by decrementing only on a real merge. Keep extra data (boss, size) at the root and carry it over when merging.

Follow-ups: Land can be removed. Many (source, target) queries (build once). Report the largest group size (track size at the root). Directed reachability is not union-find (use BFS or topological order).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Count Islands in a Grid and Track the Count as Land Cells Are Added**: each new land cell: groups += 1, then union it with every land neighbour (groups -= 1 per real merge); report groups after each addition *[PracHub (solution)]*

H ☐ ★**2 Friend Graph with Timeline (Friend / Unfriend Events)**: pairs = friend events sorted by time; the answer is the time of the union that makes groups == 1 *[1Point3Acres (notes), PracHub (solution) as "Earliest Time a Friendship Network Is Fully Connected with Friend and Unfriend Events"]*

M ☐ ★**3 2D Points Connected Components by Radius**: pairs = points within distance r, found by bucketing points into cells of size r and checking only the 9 nearby cells; answer = groups *[1Point3Acres (notes), PracHub (solution) as "Count Clusters of 2D Points Within a Radius"]*

M ☐ ★**4 Find A Threshold-Limited Path With Minimum Required Safety**: pairs = edges sorted from the highest safety down; union until find(s) == find(t); the answer is the safety of that last edge *[PracHub (solution)]*

M ☐ ★**5 Count Clusters of 2D Points Within a Radius**: identical to the line above: cell buckets give the candidate pairs, answer = groups *[PracHub (solution), 1Point3Acres (notes) as "2D Points Connected Components by Radius"]*

M ☐ ★**6 Determine Reporting Relationships**: pairs = peer links; keep the boss at the root and carry it over when two groups merge *[PracHub (solution)]*

M ☐ **Determine whether two nodes are related**: pairs = the given edges; the query is find(a) == find(b) *[PracHub (solution)]*

2. Frontier search (BFS or DFS), tier 1: spread, count, measure distance ★

Graph traversal: BFS and DFS (app section 1) 6 questions: 0 easy, 5 medium, 1 hard

Spot it: "number of islands", "flood fill", "rotten oranges", "shortest path in an unweighted graph", "distance to nearest X".

Idea: A frontier search has starts, a rule for which neighbours are usable, and a seen set. DFS or BFS both visit everything reachable. Use BFS when you need distances (layer by layer, each layer is one more step). Multi-source BFS puts every start in the queue at distance 0, so you get the distance to the nearest start for all cells in one pass. Counting components: loop over all cells and start a search from each unseen land cell.

```
def nearest_distance(grid, starts):
    rows, cols = len(grid), len(grid[0])
    dist = {s: 0 for s in starts}
    queue = deque(starts)
    while queue:
        r, c = queue.popleft()
        for nr, nc in ((r+1, c), (r-1, c), (r, c+1), (r, c-1)):
            if 0 <= nr < rows and 0 <= nc < cols and (nr, nc) not in dist \
                    and grid[nr][nc] != '#':
                dist[(nr, nc)] = dist[(r, c)] + 1
                queue.append((nr, nc))
```

```
    return dist                                 # missing key = unreachable
```

Cost: O(V + E), for a grid O(rows * cols). **Traps:** Mark cells seen when you enqueue, not when you dequeue (otherwise duplicates). Unreached cells stay unreachable (-1). To rebuild a path store parent[v] = u at first discovery.

Follow-ups: Huge grid (stream rows, or union-find). Cells change over time (rebuild only the affected area). Return every shortest path or count them (add counts per layer). Weighted edges (Dijkstra).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Number of Islands**: loop over every cell; each unseen land cell starts a new frontier_search (share one seen set); islands += 1 per start *[1Point3Acres (notes), PracHub (solution) as "Solve Number of Islands"]*

M ☐ ★**2 Rotten Oranges / Multi-Source BFS (taxis)**: starts = every rotten orange; usable = fresh orange; answer = rounds taken, or -1 if a fresh orange stays unreached *[1Point3Acres (notes), PracHub (solution) as "Solve Rotting Oranges and Bus Walk Reachability", Aced as "Rotting Oranges"]*

H ☐ ★**3 Infection Spread on a Grid (Cellular Automaton)**: starts = every infected cell; usable = healthy cell; the answer is the largest distance reached (number of rounds) *[PracHub (solution)]*

M ☐ ★**4 Find shortest relationship path using BFS**: one start; also store parent[v] = u when you first reach v; stop at the target and walk parent back to rebuild the path *[PracHub (solution), Aced (solution) as "Degrees of Friendship"]*

M ☐ **Compute distance to nearest taxi in grid**: starts = every taxi cell; usable = not a wall; the dist table is the answer, cells never reached stay -1 *[PracHub (solution)]*

M ☐ **Solve Grid and Counting Problems**: the island count again (queue version); the second part of the question is a short counting formula, no search *[PracHub (solution)]*

3. A heap of what is available: pop the best, put back the changed one ★

Heaps and top-k (app section 4) 8 questions: 0 easy, 7 medium, 1 hard

Spot it: "minimum rooms", "streaming median", "reusable locker", "schedule on identical CPUs", "no two adjacent the same".

Idea: Keep a heap of the things you can pick right now, keyed by what makes one better. Pop the best, use it, and push it back with its new state (a room with a new end time, a task with a smaller count). Two heaps hold a stream median: a max-heap for the lower half, a min-heap for the upper half, rebalanced so sizes differ by at most 1. Python only has a min-heap, so negate values for a max-heap.

```
free_at = []                                # heap of end times of rooms in use
for start, end in sorted(meetings):
    if free_at and free_at[0] <= start:
        heapq.heapreplace(free_at, end)     # reuse the room that frees first
    else:
        heapq.heappush(free_at, end)        # open a new room
rooms = len(free_at)

low, high = [], []                          # max-heap (negated), min-heap
def add(x):
    heapq.heappush(low, -x)
    heapq.heappush(high, -heapq.heappop(low))
    if len(high) > len(low): heapq.heappush(low, -heapq.heappop(high))
def median(): return -low[0] if len(low) > len(high) else (-low[0] + high[0]) / 2
```

Cost: O(log n) per push and pop. **Traps:** Tuples compare left to right: add a tie-breaker (an id) so unorderable objects never compare. Lazy deletion for stale entries. State why greedy is right (the soonest-free resource is the best one to reuse).

Follow-ups: Return which room each meeting uses. Rank by busy time. Online arrivals. Sliding median (two heaps with lazy deletion). Songs by the same artist at least k apart (hold artists in a cooldown queue).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Trip Expense Settlement**: a heap of creditors and a heap of debtors; the biggest of each settles min(amount) and the remainder goes back *[HackTheRounds #42 (solution), PracHub (solution) as "Compute minimal transfers to settle group expenses", Aced]*

M ☐ ★**3 Implement a Reusable Locker Allocation System**: heap of freed locker numbers; take the smallest freed one before handing out a new number *[PracHub (solution)]*

M ☐ ★**4 Streaming Median (LC 295) with Follow-ups**: two heaps: a max-heap for the lower half and a min-heap for the upper half; keep them balanced *[1Point3Acres (notes), PracHub (solution) as "Track the Running Median of an Integer Stream"]*

H ☐ ★**5 Huffman Tree Construction**: pop the two smallest weights, push their sum as a new node; repeat until one node remains *[1Point3Acres (notes), PracHub (solution) as "Implement binary encode/decode with frequency tree"]*

M ☐ ★**6 Compute minimum number of rooms needed**: heap of end times; a meeting reuses the room that frees first if it ended, otherwise push a new room *[PracHub (solution), Aced as "Given a list of car rental requests with start and end times, assign cars to rentals such ", 1Point3Acres (notes) as "Meeting Rooms / Interval Scheduling Variants"]*

M ☐ **Simulate room assignments for scheduled meetings**: a heap of free room numbers and a heap of (end time, room); a delayed meeting takes the earliest-ending room and shifts later *[PracHub (solution)]*
M ☐ **Design server allocation for multi-type nodes**: a heap per type of (flexibility, server); skip servers already allocated (lazy deletion) *[PracHub (solution)]*
M ☐ ★2 ★ **Playlist Shuffle No Adjacent Artists**: pop the artist with the most songs left, hold it for one turn, then push it back with a smaller count *[HackTheRounds #43 (solution)]*

4. Backtracking: choose, recurse, undo ★

Recursion and backtracking (app section 7) 7 questions: 0 easy, 6 medium, 1 hard

Spot it: "all permutations / combinations", "generate parentheses", "count paths of length k", "can this hand win", "assign people to tasks".

Idea: Build the answer one choice at a time. At each step loop over the legal choices, apply one, recurse, then undo it so the next choice starts from the same state. Prune as early as possible: only add a close bracket when closed < opened, always try the smallest remaining tile first for Mahjong. If you only need the count, return numbers instead of storing every solution. Matching problems use an augmenting path: if the preferred slot is taken, recursively ask its owner to move.

```
def parentheses(n):
    out = []
    def build(current, opened, closed):
        if len(current) == 2 * n:
            out.append(current);  return
        if opened < n:  build(current + '(', opened + 1, closed)
        if closed < opened:  build(current + ')', opened, closed + 1)
    build('', 0, 0)
    return out

def permutations(items):
    out, chosen, used = [], [], [False] * len(items)
    def build():
        if len(chosen) == len(items): out.append(chosen[:]);  return
        for i, x in enumerate(items):
            if used[i]: continue
            used[i] = True;  chosen.append(x)
            build()
            chosen.pop();  used[i] = False        # undo
    build()
    return out
```

Cost: Exponential: n! permutations, 2^n subsets, Catalan(n) bracket strings. Pruning cuts the constant. **Traps:** Copy the list when you store it (`chosen[:]`). Undo exactly what you did. Skip equal neighbours (after sorting) to avoid duplicate results. Use a bitmask for visited sets when n is small.

Follow-ups: Count only. Iterative version with an explicit stack. Next lexicographic permutation. k-th permutation directly (factorial number system). Priorities or capacities in matching.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Mahjong Winning Hand Detection (14 Tiles)**: take the smallest tile; try it as a triple, then as the start of a run; undo the counts after each try *[1Point3Acres (notes), PracHub (solution) as "Determine if a 14-tile hand is winning"]*
M ☐ ★2 **Permutations**: the pattern template above *[Aced (solution), PracHub (solution) as "Generate all permutations of an array"]*
M ☐ ★3 **Count Prefix-Matching Simple Paths in a Labeled Graph**: extend the path along edges whose label matches the next query letter, with a bitmask of visited nodes; count per depth *[PracHub (solution)]*
H ☐ ★4 **Matrix Flower Placement with House Adjacency**: place row by row; reject a row early when the houses two rows above are still not next to a flower *[1Point3Acres (notes), PracHub (solution) as "Solve Flower Placement and Directory Deletion"]*
M ☐ ★5 **Generate Parentheses**: choices are '(' and ')'; allowed when opened < n, and ')' only when closed < opened; complete at length 2n *[Aced (solution), PracHub (solution) as "Generate All Valid Combinations of Balanced Parentheses"]*
M ☐ ★6 **Determine if a 14-tile hand is winning**: the same try-triple-then-run recursion over a count array, after choosing the pair *[PracHub (solution), 1Point3Acres (notes) as "Mahjong Winning Hand Detection (14 Tiles)"]*
M ☐ **Match people to questions using tags and priorities**: try each question of a person; if it is taken, recursively ask its owner to move (augmenting path); undo by not recording on failure *[PracHub (solution)]*

5. Best-first search with a heap: Dijkstra and its minimax twin ★

Dijkstra and shortest paths (app section 6) 8 questions: 0 easy, 4 medium, 4 hard

Spot it: "shortest path with weights", "earliest arrival with timetables", "minimise the maximum height on a path", "swim in rising water".

Idea: Always expand the unfinished node with the smallest cost so far, kept in a heap. When a node is popped its cost is final (weights non-negative). The cost does not have to be a sum: use max(cost, edge) for "minimise the worst step" (minimax), or the

arrival time at the next station for timetables (you may leave only at or after you arrive). For maximising a quantity, push negated values.

```
def dijkstra(graph, source, combine=lambda cost, w: cost + w):
    best = {source: 0};  heap = [(0, source)]
    while heap:
        cost, u = heapq.heappop(heap)
        if cost > best.get(u, float('inf')): continue    # stale entry
        for v, w in graph[u]:
            new_cost = combine(cost, w)
            if new_cost < best.get(v, float('inf')):
                best[v] = new_cost
                heapq.heappush(heap, (new_cost, v))
    return best

# minimax: combine=lambda cost, w: max(cost, w)
```

Cost: O((V + E) log V). **Traps:** Skip stale heap entries. Negative edges break it (use Bellman-Ford). Track parent[v] to rebuild the path. With unit weights plain BFS is enough; with only 0 and 1 use 0-1 BFS.

Follow-ups: Return the path. Fewest transfers. Several travellers (run from each, combine at meeting nodes). Weights change over time. Layover rules (shift the departure test).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Path Minimizing the Maximum Cell (Swim in Rising Water)**: the same minimax loop with a given start and end cell (Swim in Rising Water) *[1Point3Acres (notes), PracHub (solution) as "Compute earliest arrival in time-grid"]*

M ☐ ★**2 Find the shortest path in a graph.**: combine = cost + w; also store parent[v]; stop at the target and rebuild the path *[Aced, PracHub (solution) as "Compute shortest path in cyclic graph"]*

M ☐ ★**3 Given a train schedule with departure city, arrival city, start time, end time, and duration, and given a source and destination, find the sequence of trains that gets you there at the earliest possible time.**: cost = arrival time; neighbors only offers trains that leave at or after the current time; combine = the train's arrival time *[Aced, PracHub (solution) as "Reachability Between Stations in a Timed Train Schedule"]*

H ☐ ★**4 Find minimum threshold enabling grid path**: the same minimax loop: the smallest threshold that lets you walk from corner to corner *[PracHub (solution)]*

M ☐ ★**5 Given a matrix with a rat, bread, cat, and obstacles, find a path from the rat to the bread that maximizes the minimum distance from cat.**: maximise the minimum distance to the cat along the path: push negated values, combine = min(value, cat distance of the cell) *[Aced, PracHub (solution) as "Maximize the minimum value along a grid path"]*

H ☐ **Minimize maximum height along a grid path**: the same minimax loop on a height grid *[PracHub (solution), 1Point3Acres (notes) as "Path Minimizing the Maximum Cell (Swim in Rising Water)"]*

H ☐ **Minimize travel time with optional meeting point**: three runs: from a, from b, and from the destination on the reversed graph; answer = min over meeting points m of max(da, db) + dd *[PracHub (solution)]*

H ☐ **Compute minimax grid path and network delay**: combine = max(cost, height of the next cell); the answer is the cost at the bottom-right cell *[PracHub (solution)]*

6. Topological sort (Kahn) and cycle detection ★

Topological sort and cycles (app section 29) 4 questions: 0 easy, 4 medium, 0 hard

Spot it: "course schedule", "build order", "are these race results consistent", "validate dependencies".

Idea: Compute the in-degree of every node. Put all in-degree 0 nodes in a queue, repeatedly remove one, append it to the order, and decrease its neighbours' in-degrees (adding new zeros). If the order contains every node the graph has an order; otherwise the remaining nodes lie on or behind a cycle. Layering by rounds gives the minimum number of semesters.

```
def topo_order(nodes, edges):              # edge (a, b): a must come before b
    indegree = {v: 0 for v in nodes};  after = defaultdict(list)
    for a, b in edges:
        after[a].append(b);  indegree[b] += 1
    queue = deque(v for v in nodes if indegree[v] == 0)
    order = []
    while queue:
        v = queue.popleft();  order.append(v)
        for w in after[v]:
            indegree[w] -= 1
            if indegree[w] == 0: queue.append(w)
    return order if len(order) == len(indegree) else None
```

Cost: O(V + E). **Traps:** Nodes that only appear in edges still need an in-degree entry. Direction of the edge (prerequisite to course). Duplicate edges inflate the in-degree consistently, which is fine.

Follow-ups: Report one actual cycle (DFS with colours). All missing references. Semesters (process the queue layer by layer). Recursive DFS overflow on long chains (use Kahn).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front

of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Course Scheduling and Dependency Variants**: build the vertex set from the edges too; return the order or None *[1Point3Acres (notes), PracHub (solution) as "Produce a Deterministic Course Plan"]*

M ☐ ★2 **Course Schedule**: edge = prerequisite -> course; return the order, or an empty list when there is a cycle *[Aced (solution), PracHub (solution) as "Course Schedule Feasibility"]*

M ☐ **Solve tree, graph, sliding-window problems**: the course-schedule part is can-finish: len(order) == n; the zigzag and window parts are separate small loops *[PracHub (solution)]*

M ☐ **Validate course catalog dependencies**: first reject prerequisites naming courses that do not exist; then the same peel; report the missing names or the cycle *[PracHub (solution)]*

7. Frontier search, tier 2: free steps, early exit, neighbours worked out on the fly ★

Graph traversal: BFS and DFS (app section 28) 4 questions: 0 easy, 4 medium, 0 hard

Spot it: "edge direction flips cost 1", "reach with blocked nodes", "shortest directed cycle through v", "radios within a radius".

Idea: Same BFS, but the graph is not given as lists: neighbours(u) is computed when needed (all radios within a radius, using squared distances). A mix of free and paid moves is 0-1 BFS: a deque where free steps use appendleft and paid steps use append, so the deque stays sorted by distance. Early exit: stop when the goal is popped. Shortest cycle through v: start at v and return dist[u] + 1 the moment a neighbour of u equals v.

```
def zero_one_bfs(graph, source):
    dist = {source: 0};  queue = deque([source])
    while queue:
        u = queue.popleft()
        for v, weight in graph[u]:          # weight is 0 or 1
            if dist[u] + weight < dist.get(v, float('inf')):
                dist[v] = dist[u] + weight
                if weight == 0: queue.appendleft(v)
                else:           queue.append(v)
    return dist
```

Cost: O(V + E) for BFS and 0-1 BFS; O(V^2) when neighbours are computed from all other nodes. **Traps:** Compare squared distances to avoid square roots. A blocked source means everything is unreachable. A node can be improved after first discovery in 0-1 BFS, so check strictly smaller.

Follow-ups: Arbitrary non-negative weights (Dijkstra). Shortest cycle in the whole graph (BFS from every node). Nodes blocked over time. Queries on the same graph (precompute adjacency).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Shortest Path Nodes and Escaping a Pursuer**: run it from a, from t and from b; a node is on a shortest a-to-t path when da + dt == da[t]; compare with the pursuer's distance db *[1Point3Acres (notes), PracHub (solution) as "Shortest Path, Nodes on Any Shortest Path, and Whether Alice Can Evade Bob"]*

M ☐ ★2 **Minimum Direction Violations in a Directed Graph**: an edge in its own direction costs w = 0, the same edge walked backwards costs w = 1; free steps use appendleft *[1Point3Acres (notes), PracHub (solution) as "Minimize Direction Violations in a Directed Road Network"]*

M ☐ **Compute shortest paths with blocked nodes**: neighbors skip every blocked node; if the source itself is blocked the answer is all -1; unreachable nodes stay -1 *[PracHub (solution)]*

M ☐ **Shortest Directed Cycle Through a Node**: start = v; the goal is an edge that leads back to v, so return dist[u] + 1 the moment a neighbour equals v *[1Point3Acres (notes), PracHub (solution) as "Find minimum cycle cost per node"]*

8. Top-K by frequency: Counter plus nsmallest ★

Heaps and top-k (app section 23) 5 questions: 0 easy, 5 medium, 0 hard

Spot it: "k most frequent words", "top chatters", "task scheduler with cooldown", "bottom k".

Idea: Count with Counter, then take the k best with a size-k heap: heapq.nsmallest(k, items, key=lambda w: (-count[w], w)) gives most frequent first with ties by name. This costs O(d log k) for d distinct items instead of sorting everything. For a cooldown scheduler, only the biggest count matters: answer = max(len(tasks), (biggest - 1) * (n + 1) + number of tasks tied for the biggest).

```
counts = Counter(word for line in lines for word in line.split()[:1])
top = heapq.nsmallest(k, counts, key=lambda w: (-counts[w], w))

def least_interval(tasks, n):
    counts = Counter(tasks).values()
    biggest = max(counts)
    tied = sum(1 for c in counts if c == biggest)
    return max(len(tasks), (biggest - 1) * (n + 1) + tied)
```

Cost: O(N + d log k) time, O(d) space. **Traps:** State the tie-break rule. A stream with too many distinct keys needs an approximate counter (Count-Min sketch, Misra-Gries). The log may not fit on one machine: count per machine and merge, or partition by hash.

Follow-ups: Unbounded stream and bounded memory. top_k must be O(k) per call (maintain buckets by count). Different cooldowns or priorities. Print an actual schedule (max-heap of counts plus cooldown queue).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Top-K Frequent in Large Logs**: count the first word of every line (the speaker); return the names of the k most frequent *[1Point3Acres (notes), PracHub (solution) as "Return the Top K Most Frequent Elements", Aced as "Top k frequent elements"]*

M ☐ ★2 **Rearrange Tasks With Cooldown**: only the biggest count matters: answer = max(len(tasks), (biggest - 1) * (n + 1) + how many tasks share the biggest count) *[PracHub (solution), Aced (solution) as "Design a task scheduler in Python."]*

M ☐ ★3 **Implement frequency + distance top-K queries**: count the point tuples; the tie-break is distance from the origin, then x, then y *[PracHub (solution)]*

M ☐ **Find top/bottom-k words in list or stream**: the same Counter; run nsmallest once with -count and once with +count for the bottom-k *[PracHub (solution)]*

M ☐ **Solve chat, grid paths, and car rentals**: the top-K part is the pattern template above; the other parts are a rolling-column path count and a greedy car assignment *[PracHub (solution)]*

9. Games: win or lose with a cache, and guessing with feedback ★

Recursion and backtracking (app section 39) 5 questions: 0 easy, 1 medium, 4 hard

Spot it: "who wins with optimal play", "merge stacks game", "guess the secret word", "Wordle strategy".

Idea: Describe the whole game position as a hashable state (a sorted tuple of stacks). A position is winning if some move leads to a losing position for the opponent; cache the answer per state so each state is solved once. For guessing games, choose the guess that splits the remaining candidates into the most even feedback groups (minimise the size of the largest group), then keep only candidates consistent with the feedback received.

```
@lru_cache(maxsize=None)
def wins(state):                          # state: sorted tuple of stack sizes
    for i in range(len(state)):
        for j in range(i + 1, len(state)):
            if state[i] + state[j] <= cap:
                rest = list(state);  merged = rest[i] + rest[j]
                del rest[j];  del rest[i]
                rest.append(merged)
                if not wins(tuple(sorted(rest))): return True
    return False                          # no move, or every move loses

def best_guess(words, candidates):
    def worst_group(guess):
        groups = Counter(matches(guess, w) for w in candidates)
        return max(groups.values())
    return min(words, key=worst_group)
```

Cost: Number of distinct states times moves per state; guessing is O(words * candidates). **Traps:** Normalise the state (sort) so equal positions share one cache entry. Say who moves next and what "no move" means (a loss in normal play). Terminal states need a base case.

Follow-ups: Count winning first moves. Misere play. A sum of independent games (Sprague-Grundy values). Weighted answer likelihoods. How big is the real state space?

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

H ☐ ★1 **Find secret word with match-count feedback**: guess the word that minimises the size of the biggest feedback group; keep the words with the same match count *[PracHub (solution)]*

M ☐ **Choose optimal guesses for green-only Wordle**: bitmask per (position, letter) of the words that have it; the best guess splits the words into the most even groups *[PracHub (solution), 1Point3Acres (notes) as "Wordle-Style 5-Letter Minimum-Guess Strategy"]*

H ☐ **Determine winner in stack-merging game**: state = the sorted tuple of stack sizes; a move merges two whose sum stays within the cap *[PracHub (solution), 1Point3Acres (notes) as "Babylon Tile-Stack Merge Game"]*

H ☐ **Wordle-Style 5-Letter Minimum-Guess Strategy**: ask about the letter at a position that splits the remaining words most evenly (min of yes-count and no-count is largest) *[1Point3Acres, PracHub (solution) as "Implement a Word Guessing Game"]*

H ☐ **Babylon Tile-Stack Merge Game**: state = the sorted tuple of stacks; a move merges two stacks that share a height or a colour; you win if some merge leaves the opponent without a winning reply *[1Point3Acres (notes), PracHub (solution) as "Determine winner in stack-merging game"]*

10. K-way merge: a heap of the next item from each source ★

Heaps and top-k (app section 35) 4 questions: 1 easy, 1 medium, 2 hard

Spot it: "merge k sorted lists", "k smallest pair sums", "sort a nearly sorted array", "top-k subset sums".

Idea: The smallest unseen item overall must be the head of one of the sources. Keep the current head of each source in a heap, pop the smallest, and push its successor (the next item of the same source). For pair sums of two sorted arrays, seed with (a[i] + b[0], i, 0); the successor of (i, j) is (i, j+1). For a k-sorted array, keep a heap of k + 1 items and always emit the smallest.

```
def merge_k(lists):
    heap = [(lst[0], i, 0) for i, lst in enumerate(lists) if lst]
    heapq.heapify(heap)
    out = []
    while heap:
        value, which, pos = heapq.heappop(heap)
        out.append(value)
        if pos + 1 < len(lists[which]):
            heapq.heappush(heap, (lists[which][pos + 1], which, pos + 1))
    return out
```

Cost: O(N log k) for N items in total from k sources. **Traps:** Include the source id so equal values never compare lists or nodes. Seed from the largest end if you want the largest items. Do not merge lists one by one into a growing result (O(N k)).

Follow-ups: Lists too large for memory (streams or files, same algorithm). Not sorted (sort first or use a selection algorithm). Three arrays. Return the subsets themselves.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

H ☐ ★**1 Merge k sorted linked lists.**: seed with each list head (value, list number, node); the successor of a node is node.next *[Aced (solution), PracHub (solution) as "Merge K Individually Sorted Integer Arrays Into One Sorted Array"]*

E ☐ **Find k pairs with smallest sums**: seed with (a[i] + b[0], i, 0); the successor of pair (i, j) is (i, j + 1) *[PracHub (solution)]*

M ☐ **Sort a nearly sorted (or K sorted) array.**: keep a heap of the first k + 1 items; pop the smallest and push the next input item (heapreplace) *[Aced (solution), PracHub (solution) as "Sort a nearly sorted array"]*

H ☐ **Return top-k subset sums in descending order**: heap of (sum, index); each popped state pushes two successors: add the next item, or swap the last item for the next *[PracHub (solution)]*

11. Rolling state: the next state comes only from the previous one

Dynamic programming and memoization (app section 19) 5 questions: 2 easy, 3 medium, 0 hard

Spot it: "climbing stairs", "house robber", "paint houses", "fibonacci", "dice with limits".

Idea: If dp[i] depends only on dp[i-1] (and dp[i-2]), keep just those few variables instead of the full table. Climbing stairs: (a, b) = (b, a + b). House robber: keep (best if you take this house, best if you skip it). Paint house with k colours: keep a row of costs and, for each colour, add the smallest previous value that is not the same colour (store the two smallest).

```
def rob(money):
    take, skip = 0, 0
    for x in money:
        take, skip = skip + x, max(take, skip)
    return max(take, skip)

def climb(n):
    a, b = 1, 1
    for _ in range(n - 1):
        a, b = b, a + b
    return b
```

Cost: O(n) time, O(1) space (O(k) for a row of colours). **Traps:** Initialise with the true base cases. Update both variables in one tuple assignment. Circular houses: run twice, once without the first house and once without the last.

Follow-ups: n up to 10^18 (matrix exponentiation, or fast doubling). Steps 1 to k (sliding window sum). Broken steps. Return which houses (keep parent choices). Closed form (Binet) and its floating point limits.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Given n houses in a line with money in each, find the maximum amount a robber can steal, without stealing from two adjacent houses.**: state = (best if you take this house, best if you skip it): take, skip = skip + x, max(take, skip) *[Aced (solution), PracHub (solution) as "Maximize sum without choosing adjacent elements"]*

E ☐ **Fibonacci Numbers**: the same pair update starting from (0, 1) *[Aced (solution), PracHub (solution) as "Compute Fibonacci Numbers"]*

M ☐ **Maximize sum without choosing adjacent elements**: exactly the house robber update on the array *[PracHub (solution)]*

E ☐ ★**2** ★ **Climbing Stairs**: state = the last two counts: a, b = b, a + b, repeated n - 1 times *[Aced (solution)]*

M ☐ ★**3** ★ **Write SQL code to publish the Fibonacci series.**: WITH RECURSIVE keeps (n, a, b); every new row is (n + 1, b, a + b): the same rolling pair, written

as a query *[Aced (solution)]*

12. Recurrence on a position or a range (memo or table) ★

Dynamic programming and memoization (app section 5) 8 questions: 0 easy, 6 medium, 2 hard

Spot it: "edit distance", "longest palindromic subsequence", "word break", "split into expressions", "count decompositions".

Idea: Define best(i) for a suffix or best(i, j) for a range in words, write how it follows from smaller states, and either memoise the recursion or fill a table in an order where smaller states come first. Strings: dp[i][j] for prefixes of two strings (equal letters copy the diagonal, otherwise 1 + the cheapest neighbour). Ranges: best(l, r) tries every split point m. Digit DP walks the number from the top with a tight flag and a state summary.

```
def edit_distance(a, b):
    previous = list(range(len(b) + 1))
    for i in range(1, len(a) + 1):
        current = [i] + [0] * len(b)
        for j in range(1, len(b) + 1):
            if a[i - 1] == b[j - 1]: current[j] = previous[j - 1]
            else: current[j] = 1 + min(previous[j], current[j - 1], previous[j - 1])
        previous = current
    return previous[-1]

@lru_cache(maxsize=None)
def ways(i):                                # decompositions of text[i:]
    if i == len(text): return 1
    return sum(ways(i + len(w)) for w in dictionary if text.startswith(w, i))
```

Cost: States times transitions, for example O(n * m) for edit distance and O(n^2 * split) for range DP. **Traps:** Write the state and recurrence in plain words before coding. Base cases. For palindromic subsequences the table is indexed by (i, j) with i <= j. Memoised recursion can hit the recursion limit on long inputs.

Follow-ups: Return the actual solution (keep choices). Count modulo a prime. Count in a range [A, B] (count up to B minus up to A-1). Faster version (Fenwick tree, sparse table). Different costs per operation.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Target Expressions from Ordered Digits**: best(l, r) = every value reachable from digits l..r; split at m and combine the two sides with +, - and * *[1Point3Acres (notes), PracHub (solution) as "Insert Operators to Reach a Target"]*

M ☐ ★2 **Count Word Decompositions into Chemical Element Symbols**: bottom-up ways[i]: add ways[i+1] if the one-letter symbol exists and ways[i+2] if the two-letter symbol exists *[1Point3Acres (notes), PracHub (solution) as "Count Spellings Using Chemical Element Symbols"]*

H ☐ ★3 **Edit distance**: table over two strings, two rows: equal letters copy the diagonal, otherwise 1 + the smallest of the three neighbours *[Aced (solution), PracHub (solution) as "Compute Edit Distance"]*

H ☐ ★4 **Count Perfect Wake Numbers up to N**: digit DP: cache best(remaining digits, used-digit mask, last digit, falling?) and walk along N with a tight flag *[1Point3Acres (notes), PracHub (solution) as "Count Good Numbers up to a Limit"]*

M ☐ ★5 **Fence Painting Minimum Operations (Gemini Team)**: range (l, r): paint one stripe at the lowest height, then solve the pieces on both sides; a sparse table finds the lowest point fast *[1Point3Acres (notes), PracHub (solution) as "Minimize Fence Painting Operations"]*

M ☐ ★6 **Find the longest palindromic subsequence using dynamic programming.**: table over (i, j): equal ends add 2 to the inside; otherwise the better of dropping either end *[Aced (solution), PracHub (solution) as "Compute longest palindromic subsequence"]*

M ☐ **Maximize 1D deliveries within distance and minimize total distance**: table over (orders used, drivers used): skip an order, skip a driver, or match them when they are within D (count, then minus distance) *[PracHub (solution)]*

M ☐ ★ **Jump Game with Profit**: best[i] = max over every j before i that can reach i (j + arr[j] >= i) of best[j] + profit[i] *[1Point3Acres (notes)]*

13. Grid DP: each cell from its neighbours above and to the left ★

Dynamic programming and memoization (app section 26) 5 questions: 0 easy, 3 medium, 2 hard

Spot it: "unique paths", "largest square", "longest run of one colour", "paths with checkpoints".

Idea: Fill the table row by row; each cell depends on cells already filled. Number of paths: ways[r][c] = ways[r-1][c] + ways[r][c-1] (zero on blocked cells or missed checkpoints). Same-colour squares: if the cell and its three neighbours match, side[r][c] = 1 + min of the three neighbouring sides. Longest run in four directions: keep four run lengths per cell from the previous row and the left cell.

```
squares = 0
side = [[0] * cols for _ in range(rows)]
for r in range(rows):
    for c in range(cols):
        if r and c and grid[r][c] == grid[r-1][c] == grid[r][c-1] == grid[r-1][c-1]:
            side[r][c] = 1 + min(side[r-1][c], side[r][c-1], side[r-1][c-1])
        else:
            side[r][c] = 1
        squares += side[r][c]               # a square of side s counts s squares
```

Cost: O(rows * cols) time, O(cols) space if you keep only the previous row. **Traps:** Edges (first row and column). Diagonal moves

need three predecessors. A checkpoint resets rows that do not pass through it.

Follow-ups: Largest square only. Allow one different cell. Grid arrives row by row (keep one row). Return where the run is.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Grid Path Count with Diagonal Moves and Checkpoints**: advance column by column keeping the number of ways in each row; new[r] = old[r-1] + old[r] + old[r+1]; zero out rows that miss a checkpoint *[1Point3Acres (notes), PracHub (solution) as "Count Up-Right/Right/Down-Right Grid Paths Through Ordered Checkpoints"]*

M ☐ ★**2 Longest Same-Color Grid Sequence**: keep four run lengths per cell (horizontal, vertical, two diagonals) from the previous row and the left cell; the answer is the largest *[1Point3Acres (notes), PracHub (solution) as "Find the Longest Straight Run of One Color"]*

H ☐ ★**3 Find the Longest Straight Run of One Color**: the same four run lengths; extend a run when the neighbour in that direction has the same colour *[PracHub (solution), 1Point3Acres (notes) as "Longest Same-Color Grid Sequence"]*

M ☐ **Count Same-Color Squares in Matrix**: the pattern template above: a square ends here if the three neighbours match; its size is 1 + min of their sizes *[1Point3Acres (notes), PracHub (solution) as "Count same-color squares in a character grid"]*

H ☐ **Count same-color squares in a character grid**: exactly the same recurrence on characters *[PracHub (solution), 1Point3Acres (notes) as "Count Same-Color Squares in Matrix"]*

14. Counting or reaching an amount: the knapsack loop

Dynamic programming and memoization (app section 32) 4 questions: 0 easy, 4 medium, 0 hard

Spot it: "coin change", "partition equal subset sum", "target sum with plus and minus", "unbounded items that fill a capacity".

Idea: dp[t] is the answer for total t. For each item, update dp across totals. Use items unlimited times: loop t upward (dp[t] = min(dp[t], dp[t - c] + 1)). Use each item once: loop t downward so an item cannot be used twice in one pass. For plus/minus signs, rewrite as a subset sum of (total + target) / 2. A bitset (Python int) makes subset sum fast: reach |= reach << x.

```
def coin_change(coins, amount):
    best = [0] + [float('inf')] * amount
    for t in range(1, amount + 1):
        for c in coins:
            if c <= t: best[t] = min(best[t], best[t - c] + 1)
    return -1 if best[amount] == float('inf') else best[amount]

def can_partition(nums):
    total = sum(nums)
    if total % 2: return False
    reach = 1                              # bit t is set if total t is reachable
    for x in nums: reach |= reach << x
    return (reach >> (total // 2)) & 1 == 1
```

Cost: O(amount * number of coins), or O(n * sum) for subset sum. **Traps:** Count combinations (loop items outside) versus permutations (loop totals outside). Direction of the inner loop decides reuse. Negative numbers break the plain table.

Follow-ups: Count the ways (LeetCode 518). Return which coins. Minimise the difference of the two parts. Does greedy work (only for special coin systems)?

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**3 Coin Change**: fewest coins: start with infinity, best[t] = min(best[t], best[t - c] + 1); -1 when it stays infinity *[Aced (solution), PracHub (solution) as "Find the Minimum Number of Coins for an Amount"]*

M ☐ **Explain how to find a target sum in an array.**: turn +/- signs into a subset sum of (total + S) / 2, items used once (sweep downwards) *[Aced (solution), PracHub (solution) as "Count expressions reaching a target sum"]*

M ☐ ★**1** ★ **Reverse Coin Change**: given the ways table, find the coin values: the first amount whose count is wrong must be a coin; add it and rebuild *[HackTheRounds #633 (solution)]*

M ☐ ★**2** ★ **Partition an array into two sub-arrays with equal sum.**: subset sum to total / 2, written with a bitset: reach |= reach << x *[Aced (solution)]*

DAY 2

Counting, scanning and searching

Patterns 15 to 28 (14 patterns)

Hash maps, prefix sums, two pointers, windows, greedy, binary search, sorted intervals and sweeps.

15. Dict of per-key state (update, expire, rank)

Hash maps and sets (app section 10) 8 questions: 1 easy, 7 medium, 0 hard

Spot it: "remember the last time we saw X", "top users", "within the last 10 seconds", "design a logger / counter".

Idea: Keep one dictionary whose key is the thing you track (message, user, word) and whose value is the small state about it: a timestamp, a count, a set, a deque. Every operation is a lookup, an update, or an expiry check on that one entry, so each call costs O(1) on average. To rank, keep a second structure (buckets by count, or sort at read time) instead of re-sorting on every update.

```
last_shown = {}
def should_print(now, message, window=10):
    if message in last_shown and now - last_shown[message] < window:
        return False
    last_shown[message] = now
    return True

hits = deque()                          # sliding window of hit times
def hit(now):  hits.append(now)
def count(now, window=300):
    while hits and hits[0] <= now - window:
        hits.popleft()                  # expire from the old end
    return len(hits)
```

Cost: O(1) per call; expiry is amortised O(1) because each item is removed once. **Traps:** Decide first whether a rejected call refreshes the timestamp (usually it does not). Timestamps can repeat. Use < or <= consistently at the window edge.

Follow-ups: Out of order events (keep sorted, or a heap). Bounded memory (evict expired keys lazily or in a background sweep). Thread safety (one lock, or a lock per key). Millions of hits per second (bucket by second instead of storing every hit).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Active Communication Top Users**: a set of seen pairs, a Counter of partners per user, and buckets by count to read the top users in O(1) *[1Point3Acres (notes), PracHub (solution) as "Most Active Users in a Live Communication Stream"]*

E ☐ ★2 **Robot Status Message Deduplication Window**: the pattern template above: show a message only if it was not shown in the last 10 seconds *[1Point3Acres (notes), PracHub (solution) as "Filter Repeated Robot Status Messages"]*

M ☐ ★3 **Next-Word Predictor**: defaultdict(Counter): counts[a][b] += 1 for neighbouring words; remember the best next word per word *[1Point3Acres (notes), PracHub (solution) as "Build a Next-Word Predictor"]*

M ☐ ★5 **Hit Counter**: one deque of hit times; drop the ones older than 300 seconds; the answer is the length *[HackTheRounds #39 (solution), PracHub (solution) as "Design a rolling hit counter API"]*

M ☐ ★6 **Filter Repeated Robot Status Messages**: the same rule applied to a stream of (timestamp, message) pairs; keep what is accepted *[PracHub (solution), 1Point3Acres (notes) as "Robot Status Message Deduplication Window"]*

M ☐ ★7 **Website Activity Analytics: Unique Users + Session Time**: dict activity -> set of users for unique counts; dict user -> sorted events to add up login-to-logout session time *[1Point3Acres (notes), PracHub (solution) as "User Activity Analytics: Unique Users per Activity and Average Time on Site"]*

M ☐ **Design a logger that returns top-K messages**: dict message -> list of timestamps; top_k counts each list inside the time range with bisect, then takes the k largest *[PracHub (solution)]*

M ☐ ★4 ★ **Batch and Streaming Log Deduplication**: dict message -> first (or latest) timestamp depending on the policy; sort the result at the end *[1Point3Acres (notes)]*

16. Sweep over events in time order ★

Intervals and scheduling (app section 2) 9 questions: 0 easy, 6 medium, 3 hard

Spot it: "minimum rooms", "maximum overlap", "skyline", "who is online at time t", "servers needed".

Idea: Turn every interval into two events (+1 at start, -1 at end), sort all events by time, and sweep while keeping a running count; the maximum count is the answer. Break ties explicitly (ends before starts if a touching meeting may reuse the room). In two dimensions, put +1/-1 at the four corners of a rectangle in a difference array and take running sums along both axes. A skyline sweep keeps a max-heap of active heights and records a point when the top changes.

```
events = []
for start, end in meetings:
    events.append((start, 1))
    events.append((end, -1))                # -1 sorts before +1 at equal time
events.sort()
rooms = best = 0
for time, delta in events:
    rooms += delta
    best = max(best, rooms)

diff = [[0] * (w + 1) for _ in range(h + 1)]   # rectangles on a grid
for x1, y1, x2, y2 in rects:
    diff[y1][x1] += 1;  diff[y1][x2] -= 1
    diff[y2][x1] -= 1;  diff[y2][x2] += 1
```

Cost: O(n log n) for the sort, O(n) for the sweep. **Traps:** Tie-break at equal times. Meetings that cross midnight must be split in two. Merge one job's own intervals first if it should count once. Simulations (two queues at a gate) are sweeps with a rule per second.

Follow-ups: Return which room or server each interval gets (heap of free rooms). Coordinates up to 10^9 (coordinate compression or a sparse map). Online arrivals (cannot sort ahead; use a heap of end times).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Train / Flight Schedule Reachability**: sort legs by departure; a leg is usable if you have already reached its origin before it leaves; keep the earliest arrival per place *[1Point3Acres (notes), PracHub (solution) as "Reachability Between Stations in a Timed Train Schedule"]*

M ☐ ★**2 Meeting Rooms III Variant**: split meetings that cross midnight into two pieces first, then the same sweep *[HackTheRounds #630 (solution), PracHub (solution) as "Compute minimum rooms for time intervals"]*

M ☐ ★**3 Meeting Rooms / Interval Scheduling Variants**: the pattern template above: the largest running count is the number of rooms *[1Point3Acres (notes), PracHub (solution) as "Assign Meetings to the Minimum Number of Rooms"]*

H ☐ ★**4 Count Overlapping Rectangle Updates on a Grid**: the same idea in two dimensions: a difference array (+1, -1, -1, +1 at the corners) and two running-sum passes *[PracHub (solution)]*

M ☐ ★**5 Solve Rooms and Top-K Streams**: sort starts and ends separately and walk two pointers: a start before the earliest end needs a new room; the stream part is a size-K heap *[PracHub (solution)]*

M ☐ ★**6 Gate Entry / Exit Time Resolution with Tie-Break Rules**: the same two-queue simulation with the tie-break rules from the statement *[1Point3Acres (notes), PracHub (solution) as "Compute Turnstile Crossing Times"]*

H ☐ ★**7 Compute Turnstile Crossing Times**: sort people by arrival time; two queues (entering, exiting) with a rule for who goes first at each second *[PracHub (solution)]*

M ☐ **Compute servers needed for daily recurring jobs**: merge each job's own intervals first so one job counts once, then sweep all of them together *[PracHub (solution)]*

H ☐ **Compute city skyline outline**: events at each building's left and right edge; a max-heap of active heights; record a point when the top height changes *[PracHub (solution)]*

17. Binary search on a yes/no condition ★

Binary search (app section 11) 7 questions: 0 easy, 4 medium, 3 hard

Spot it: "minimum speed / capacity / time such that ...", "largest length such that ...", "first bad version", "median of two sorted arrays".

Idea: Find a quantity x and a test ok(x) that is monotone: false up to a threshold and true after it (or the reverse). Then binary search the threshold in O(log range) calls to ok. Write ok(x) as a simple linear scan or sum. To find the largest x that works, search the first x that fails and subtract one.

```
def first_true(low, high, ok):              # smallest x in [low, high] with ok(x)
    while low < high:
        mid = (low + high) // 2
        if ok(mid): high = mid
        else:       low = mid + 1
    return low

speed = first_true(1, max(piles),           # minimum eating speed
        lambda v: sum(-(-p // v) for p in piles) <= hours)
```

Cost: O(log(range) * cost of ok). **Traps:** Choose bounds that surely contain the answer (and say why). Use `mid = (low + high) // 2` with `high = mid` to avoid an infinite loop. Median of two arrays: binary search how many items come from the shorter array.

Follow-ups: Unknown length or huge file (double the index until the condition flips). Rate limited oracle (parallel search: more calls, fewer rounds). k sorted arrays or the k-th element (generalise the partition).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

H ☐ ★**1 Median of Two Sorted Arrays (LC 4) with a Follow-up Variant**: x = how many items come from the shorter array's left side; ok(x) = the left part's largest <= the right part's smallest *[1Point3Acres (notes), PracHub (solution) as "Median of All Values Across Several Sorted Integer Arrays"]*

M ☐ ★**2 Cutting Ribbons (LC 1891)**: x = the piece length; ok(x) = the ribbons give at least k pieces; you want the LARGEST ok, so search the first x that fails and subtract 1 *[1Point3Acres (notes), PracHub (solution) as "Find maximum cable length by cutting"]*

M ☐ ★**3 First Bad Version with Parallel Search Follow-up**: x = the version number; ok(x) = is_bad(x) *[1Point3Acres (notes), PracHub (solution) as "Find the First Bad Version in a Monotonic Release History"]*

H ☐ ★**4 Find the median of two sorted arrays.**: the same partition search on the shorter array *[Aced, PracHub (solution) as "Find median of two sorted arrays"]*

M ☐ **Solve Banana Speed and Interval Merge**: x = the eating speed; ok(x) = total hours (ceil of pile / x summed) <= h; first_true over 1..max pile *[PracHub (solution)]*

H ☐ **Find longest duplicated substring**: x = the substring length; ok(x) = some substring of that length appears twice (rolling hash); search the largest ok *[PracHub*

(solution)]

M ☐ ★ **Given a list with response time for each server, calculate how long time it will take to process n requests.**: x = the elapsed time; ok(x) = sum(x // time_i) >= n requests *[Aced (solution)]*

18. Sorted intervals: merge while they touch, then read gaps or search ★

Intervals and scheduling (app section 16) 6 questions: 1 easy, 3 medium, 2 hard

Spot it: "merge intervals", "do bookings conflict", "free time", "insert an interval", "union of two lists".

Idea: Sort by start. Walk once: if the next interval starts at or before the current end, extend the end with max; otherwise close the current one and start a new one. After merging, intervals are disjoint and sorted, so the gaps are the spaces between them and a point query is a bisect over the starts. A booking conflicts if the previous interval ends after the new start or the next one starts before the new end.

```
intervals.sort()
merged = [list(intervals[0])]
for start, end in intervals[1:]:
    if start <= merged[-1][1]:              # touches or overlaps
        merged[-1][1] = max(merged[-1][1], end)
    else:
        merged.append([start, end])

starts = [s for s, _ in merged]
i = bisect_right(starts, x) - 1             # the only candidate containing x
inside = i >= 0 and merged[i][1] >= x
```

Cost: O(n log n) to sort and merge, O(log n) per query. **Traps:** Decide whether touching intervals ([1,3] and [4,5], or [1,3] and [3,4]) merge. Take max of the ends, not the new end. Closed versus half-open intervals change every comparison.

Follow-ups: Insert one interval into a merged list. Remove a range. Several blocked windows (subtract each). Double booking allowed but never triple (counter dict of boundaries). Too many intervals for memory (stream and merge in chunks).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Union of Sorted Interval Lists**: walk both sorted lists with two indices, always taking the smaller start next, then the same merge test *[1Point3Acres (notes), PracHub (solution) as "Compute the Union of Two Sorted Interval Lists"]*

M ☐ ★**2 Merge Intervals**: the pattern template above *[Aced (solution), PracHub (solution) as "Merge overlapping intervals"]*

E ☐ ★**3 Check an Interval Against Existing Intervals for Conflicts**: bisect on the starts; a new booking conflicts if the previous one ends after its start or the next one starts before its end *[PracHub (solution)]*

M ☐ ★**4 Find Containing Range**: merge once, keep the starts in a list, and bisect to find the one interval that could contain x *[PracHub (solution)]*

H ☐ **Find safe travel intervals between planet influences**: turn each planet into (centre - radius, centre + radius), merge, and read the gaps (including the infinite ends) *[PracHub (solution)]*

H ☐ **Implement a function that processes a stream of interval data.**: the Range Module add: bisect, merge the block, splice in one interval; intervals() returns the pairs *[Aced, PracHub (solution) as "Insert and merge an interval"]*

19. Greedy: scan once, carry the best so far, act only when forced ★

Greedy (app section 12) 5 questions: 1 easy, 4 medium, 0 hard

Spot it: "minimum number of ...", "maximum profit", "largest number you can form", "cover all holes", "jumps".

Idea: Make the locally best choice and never revisit it. It is correct when you can argue an exchange: any optimal answer can be changed to use your choice without getting worse. Typical shapes: carry the farthest reach and jump only when you hit the end of the current jump; scan digits from the right fixing violations; sort then take as much of the best item as fits.

```
jumps = end = farthest = 0                  # minimum jumps
for i in range(len(a) - 1):
    farthest = max(farthest, i + a[i])
    if i == end:                            # forced to jump now
        jumps += 1
        end = farthest

holes.sort();  boards = 0;  covered_to = -1   # cover holes with boards
for h in holes:
    if h > covered_to:
        boards += 1
        covered_to = h + length - 1
```

Cost: O(n) or O(n log n) when a sort is needed. **Traps:** State why greedy is valid (exchange argument) and give a counter-example where it fails (general knapsack, coins like 1, 3, 4). Ties and empty input.

Follow-ups: Many trades or a fee (turn into DP). Return the chosen indexes. Weights not all 1 (greedy may stop working: ask what changes). Huge numbers given as strings.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Bounded Knapsack Variant**: sort items by value, highest first; take min(available count, remaining capacity) of each *[1Point3Acres (notes), PracHub (solution) as "Maximum Value from Unit-Weight Items with Limited Copies Under a Capacity"]*

M ☐ ★2 **Visible People in a Line (Last Person, Then Everyone)**: scan from the end carrying the tallest person seen between; count a person when no taller one blocks the view *[1Point3Acres (notes), PracHub (solution) as "Count visible people to the right"]*

M ☐ ★3 **Find the Largest Monotone Increasing Number**: scan the digits right to left; when a digit is smaller than the one before it, lower the earlier digit and fill 9s after it *[PracHub (solution)]*

E ☐ ★4 **Maximum Profit**: pick the cheapest and the dearest city with a fixed tie-break (min and max once) *[Aced (solution), PracHub (solution) as "Maximum Profit from One Stock Purchase and a Later Sale"]*

M ☐ ★5 **Min boards to cover roof holes**: the pattern template above *[PracHub (solution)]*

20. Count pairs by looking up each partner

Hash maps and sets (app section 30) 3 questions: 0 easy, 2 medium, 1 hard

Spot it: "pairs with difference k", "two ways to write as a sum", "rectangle / square from points".

Idea: For every item, the partner you need is determined by the rule (v + k, target - v, the opposite corner). Store what you have seen in a Counter or set, then look the partner up in O(1) instead of scanning all other items. When a pair would be found twice, either count only one direction or divide the total by 2.

```
counts = Counter(values)
pairs = 0
for v in counts:
    if k == 0:
        pairs += counts[v] >= 2          # the partner is itself
    elif v + k in counts:
        pairs += 1

points = set(map(tuple, grid_points))     # squares from two opposite corners
for (x1, y1), (x2, y2) in combinations(points, 2):
    if x1 != x2 and y1 != y2 and abs(x1 - x2) == abs(y1 - y2):
        total += (x1, y2) in points and (x2, y1) in points
total //= 2                               # each square is seen from both diagonals
```

Cost: O(n) for single partner lookups, O(n^2) for all pairs of points. **Traps:** The partner may equal the item itself (k = 0). Count distinct pairs, not index pairs. Bucket pairs by a signature (sum of coordinates, squared length) to find rectangles without checking all quadruples.

Follow-ups: Stage 2 with new rules (keep the code adaptable). Duplicates allowed (count with multiplicity). Beat O(n^2) pairs (bucket by signature).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Sum of Two Squares in Two Ways**: enumerate every pair (a, b) with a <= b once, add a*a + b*b to a Counter, answer = sums counted exactly twice *[HackTheRounds #632 (solution), PracHub (solution) as "Find numbers with exactly two sum-of-squares forms"]*

H ☐ ★2 **Maximum-Area Rectangle from Planar Points**: bucket every pair of points by (sum of coordinates, squared length); two pairs in one bucket are the diagonals of a rectangle *[1Point3Acres (notes), PracHub (solution) as "Largest-Area Rectangle From a Point Set When Sides Need Not Be Axis-Aligned"]*

M ☐ ★3 ★ **Progressive Array Partitioning**: Counter of the values; stage 1 asks whether all the counts are different: len(set(counts.values())) == len(counts) *[1Point3Acres (notes)]*

21. Prefix sums, tier 1: precompute once, query in O(1)

Arrays and prefix sums (app section 15) 6 questions: 1 easy, 5 medium, 0 hard

Spot it: "sum of a range", "balanced split", "pivot index", "equal halves", "area up to a height".

Idea: Let prefix[i] be the sum of the first i items. Then any range sum is prefix[r+1] - prefix[l], and a balance point is an index where the left total equals total minus left minus the item. The same trick works for counts of any property that adds up (equal-parity neighbours, ones in a string, area strips).

```
prefix = [0]
for x in a:
    prefix.append(prefix[-1] + x)
def range_sum(l, r): return prefix[r + 1] - prefix[l]

total = sum(a);  left = 0
for i, x in enumerate(a):                 # pivot index
    if left == total - left - x:
```

```
        return i
    left += x
```

Cost: O(n) to build, O(1) per query. **Traps:** Off by one between prefix[i] and prefix[i+1]. Use prefix[0] = 0 so a range starting at 0 works. For a removed element, look up the value that would balance both sides in a dict of first/last positions.

Follow-ups: Array changes between queries (Fenwick tree or segment tree). Circular array (double it). Remove one or k elements. Count the ways instead of finding one.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

E ☐ ★1 **Parity Alternating Subarray Query**: pre counts neighbouring pairs of equal parity; a range alternates when pre[r] - pre[l] == 0 *[HackTheRounds #35 (solution), PracHub (solution) as "Answer parity-alternation range queries"]*

M ☐ ★2 **Range Duplicate Check with Constant-Time Queries**: next[i] = next occurrence of the same value; reach[l] = smallest next over l..end; the range l..r is all distinct when reach[l] > r *[1Point3Acres (notes), PracHub (solution) as "O(1) Range Queries: Are All Elements of a Subarray Distinct?"]*

M ☐ ★3 **Find a Horizontal Cut That Bisects Rectangle Area**: events by height with a running width; accumulate area strip by strip until half the total, then solve inside that strip *[PracHub (solution)]*

M ☐ ★5 **Check pivot after removing one element**: prefix sums plus first/last position maps: look up the value whose removal would balance both sides *[PracHub (solution)]*

M ☐ ★4 ★ **Gold Chain Split into Two Equal Halves**: prefix sums with first/last position maps: find two cut points that leave equal halves around one removed link *[1Point3Acres (notes)]*

M ☐ ★ **Count the number of ways to split a string.**: positions of the ones; the answer multiplies the gap sizes around the k-th and 2k-th one (zero ones: choose two cut points) *[Aced (solution)]*

22. Runs and groups: extend while it fits, restart when it breaks

Arrays and prefix sums (app section 8) 4 questions: 0 easy, 4 medium, 0 hard

Spot it: "longest contiguous run", "group consecutive items", "streak", "one replacement allowed".

Idea: Walk the array once with a running length (or running sum). If item i fits with item i-1, extend the run; otherwise close it, record the best, and restart at 1. For "one change allowed", precompute for each index the run length ending there (left[i]) and starting there (right[i]); changing the element at i can join left[i-1] and right[i+1].

```
best = run = 1
for i in range(1, len(a)):
    run = run + 1 if a[i] >= a[i - 1] else 1
    best = max(best, run)

left = [1] * n;  right = [1] * n          # one change allowed
for i in range(1, n):
    if a[i] >= a[i - 1]: left[i] = left[i - 1] + 1
for i in range(n - 2, -1, -1):
    if a[i] <= a[i + 1]: right[i] = right[i + 1] + 1
best = max(left[i - 1] + 1 + right[i + 1] for i in range(1, n - 1)
           if a[i + 1] - a[i - 1] >= 0)
```

Cost: O(n) time, O(1) extra for one pass, O(n) for the left/right tables. **Traps:** Strict versus non-strict comparison. Remember to record the last run after the loop. A run of one element is valid. The joined run needs the two outer neighbours to be compatible.

Follow-ups: Return where the longest run is (store its start). Up to k modifications (then it is a sliding window). Many (l, r) queries (prefix arrays). Group by key per user.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Longest Non-Decreasing Subarray with One Replacement**: left[i] and right[i] run lengths; the best answer joins the run before and after one changed element *[1Point3Acres (notes), PracHub (solution) as "Longest Non-Decreasing Subarray After One Replacement"]*

M ☐ **Find longest increasing contiguous subarray**: fits = a[i] > a[i-1]; also keep a left run and a right run per index so one change can join two runs *[PracHub (solution)]*

M ☐ ★2 ★ **Piano Hand Movement Grouping**: start a new group when adding the note would make max - min exceed k *[1Point3Acres (notes)]*

M ☐ ★ **Sum of Good Subarrays (Adjacent Diff ±1)**: keep the running length and running sum of the current up-run and down-run; add both to the total at every index *[1Point3Acres (notes)]*

23. bisect on a sorted list: position, range count, keep sorted ★

Binary search (app section 14) 6 questions: 0 easy, 4 medium, 2 hard

Spot it: "count in a range", "k-th largest as items stream in", "nearest value", "longest increasing subsequence".

Idea: bisect_left(a, x) gives the first index with a[i] >= x and bisect_right gives the first index with a[i] > x. Their difference counts items equal to x; two calls count a range. insort keeps a list sorted while inserting. The tails list of the longest increasing subsequence is also maintained with bisect_left: replace the first tail that is at least x, or append.

```
def count_in_range(sorted_a, low, high):   # low <= a[i] <= high
    return bisect_right(sorted_a, high) - bisect_left(sorted_a, low)

tails = []                                 # longest increasing subsequence length
for x in values:
    i = bisect_left(tails, x)
    if i == len(tails): tails.append(x)
    else:               tails[i] = x
length = len(tails)

envelopes.sort(key=lambda e: (e[0], -e[1]))  # width up, height down
```

Cost: O(log n) per search, O(n) per insort (list shift), O(n log n) for LIS. **Traps:** Strict versus non-strict increase decides bisect_left versus bisect_right. For Russian dolls, sort heights descending on equal widths. bisect on prefixes: `prefix + chr(0x10FFFF)` is a safe upper bound.

Follow-ups: Deletions too (use a Fenwick tree or a sorted container). Return the actual chain (store parent links). Count smaller numbers after self (Fenwick tree or merge sort).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Sliding-Window Mean After Removing the Largest K Values**: keep the window in a sorted list; insort the new value, remove the old one with bisect_left, and track the sum of the top-K region *[PracHub (solution)]*

H ☐ ★**3 Find Maximum Envelope Nesting**: exactly the same sort and tails list *[PracHub (solution), 1Point3Acres (notes) as "Russian Doll Envelopes"]*

M ☐ **Streaming Insert + Find K-th Largest**: insort every new value; the k-th largest is sorted_list[-1 - k] *[1Point3Acres (notes), PracHub (solution) as "Design structure for insert and k-th largest"]*

M ☐ **Russian Doll Envelopes**: sort by width ascending and height descending, then the longest increasing run of heights using a tails list and bisect_left *[1Point3Acres (notes), PracHub (solution) as "Find Maximum Envelope Nesting"]*

H ☐ ★**2** ★ **Count of Smaller Numbers After Self**: go from the right, insort each number into a sorted list; the count is bisect_left of the number before inserting (a Fenwick tree is the fast version) *[HackTheRounds #36 (solution)]*

M ☐ ★ **Given positions of N kids and M ice cream sellers on a 1D line, find the smallest distance E such that every kid lies within the range [P - E, P + E] of at least one seller.**: for each kid bisect into the sorted sellers and take the nearer of the two neighbours; the answer is the largest of those distances *[Aced]*

24. Sets and counters across two collections

Hash maps and sets (app section 13) 4 questions: 1 easy, 2 medium, 1 hard

Spot it: "appears in both", "not in the other list", "duplicates", "pair each x with 2x".

Idea: Put the first collection into a set (membership) or a Counter (multiplicity), then stream the second collection past it and test each item in O(1). A set answers "is it there"; a Counter answers "how many are left", and you decrement it as you use items up. When the smaller side fits in memory, always load the smaller side.

```
seen = set()
for value in first_list:
    if value in seen:
        return True                # duplicate found
    seen.add(value)

remaining = Counter(values)        # pair each x with 2x
for x in sorted(remaining, key=abs):
    if remaining[x] > remaining[2 * x]:
        return []
    remaining[2 * x] -= remaining[x]
```

Cost: O(n + m) time, O(min(n, m)) extra space. **Traps:** Zero is its own double. Negative numbers need sorting by absolute value. Do not mutate the Counter while iterating its keys unless you iterate a sorted copy.

Follow-ups: Second input is a huge stream (load only the smaller set). No extra memory allowed (sort both and use two pointers, or use the array itself as the set by swapping). Three or more collections (intersect step by step).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

H ☐ ★**1 Smallest Missing Positive Integer in Linear Time and Constant Extra Space**: use the array itself as the set: swap each value x into slot x - 1, then the first slot with the wrong value is the answer *[PracHub (solution)]*

M ☐ ★3 **Min deletions to avoid overlap in first k**: Counter of the left part and of the right part; for each value in both, delete min(left count, right count) *[PracHub (solution)]*

M ☐ **Reconstruct original array from doubled shuffle**: Counter of the values; go from the smallest, pairing each x with 2x and using both counts up *[PracHub (solution)]*

E ☐ ★2 ★ **Find the Duplicates**: values that appear in both lists, each reported once (set intersection, or two pointers on sorted lists) *[Aced (solution)]*

25. Sliding window: grow on the right, shrink on the left

Two pointers and sliding window (app section 40) 4 questions: 0 easy, 3 medium, 1 hard

Spot it: "longest/shortest subarray or substring such that ...", "at most k distinct", "minimum window", "budget of changes".

Idea: Keep a window [left, right] and a summary of what is inside (a count dict, a sum, a deque of maxima). Move right forward one item at a time and add it. While the window is too bad, remove items from the left. For a longest window, record the size after shrinking; for a shortest window, record while the window is still valid, then shrink.

```
count = Counter();  left = 0;  best = 0
for right, ch in enumerate(s):             # longest with at most k distinct
    count[ch] += 1
    while len(count) > k:
        count[s[left]] -= 1
        if count[s[left]] == 0: del count[s[left]]
        left += 1
    best = max(best, right - left + 1)

maxima = deque()                           # window maximum
for i, x in enumerate(a):
    while maxima and a[maxima[-1]] <= x: maxima.pop()
    maxima.append(i)
    if maxima[0] <= i - size: maxima.popleft()
```

Cost: O(n): each index enters and leaves the window once. **Traps:** Delete keys whose count reaches 0 or `len(count)` lies. Works only when removing items from the left can fix the window (non-negative values, monotone condition). "Exactly k" = at most k minus at most k-1.

Follow-ups: Window by elapsed time instead of count (pop while the oldest is too old). Minimum or median in the window. Count windows instead of finding the best. Many queries with different parameters.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Shortest Subarray with K Distinct Integers**: state = a count dict; while it holds at least k distinct values, record the window and shrink from the left *[HackTheRounds #34 (solution), PracHub (solution) as "Find shortest subarray with at least k distinct"]*

H ☐ ★3 **Necklace Cut into Two Halves with Equal D/R Counts**: a fixed-size window of half the length slides around the doubled string, keeping a count of D *[1Point3Acres (notes), PracHub (solution) as "Find fair split of a two-color necklace"]*

M ☐ ★4 **Find the minimum window substring.**: state = how many of each needed letter are still missing; while nothing is missing, record the window and shrink *[Aced, PracHub (solution) as "Minimum Window Substring with Repeated Required Characters"]*

M ☐ ★2 ★ **Temperature Monitoring Data Structure**: a deque of the last n readings with a running total; a second deque of decreasing values gives the maximum *[1Point3Acres (notes)]*

26. Two pointers moving toward each other (or along two lists)

Two pointers and sliding window (app section 20) 6 questions: 4 easy, 2 medium, 0 hard

Spot it: "palindrome", "reverse", "sorted array and a target pair", "merge two sorted lists", "triplets that sum to zero".

Idea: Put one pointer at each end (or at the start of each list) and move whichever pointer the rule says. On a sorted array, if the pair sum is too small you must move the left pointer right, and if too large the right pointer left; this discards a whole row of pairs per step. For triplets, fix the first number and run the pair routine on the rest.

```
i, j = 0, len(a) - 1
while i < j:                               # pair with a given sum
    s = a[i] + a[j]
    if s == target: return i, j
    if s < target: i += 1
    else:          j -= 1

previous, current = None, head             # reverse a linked list
while current:
    current.next, previous, current = previous, current, current.next
```

Cost: O(n) after sorting, O(n log n) with the sort, O(1) space. **Traps:** Sort first when the problem is about values, not positions. Skip duplicates for unique triplets. Update all three pointers in one assignment when reversing links.

Follow-ups: Do it recursively. Reverse only a range or in groups of k. Allow one deletion (palindrome). Runs stored as (value,

count): advance by min of the two remaining counts.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

E ☐ ★1 **Reverse a linked list.**: three pointers prev, current, next: point current.next back at prev, then move all three one step *[Aced (solution), PracHub (solution) as "Reverse a linked list"]*

E ☐ ★2 **Sort a binary array in linear time.**: left skips 0s, right skips 1s; when left sits on a 1 and right on a 0, swap them *[Aced, PracHub (solution) as "Sort Three Categories In Place"]*

E ☐ **Reverse a Sentence**: swap from both ends to reverse the whole array, then reverse every word the same way *[Aced (solution), PracHub (solution) as "Reverse Words While Preserving Spaces"]*

M ☐ **Find median of merged RLE arrays**: walk both sorted run lists, always taking the smaller value, adding its count until you pass the middle position *[PracHub (solution)]*

M ☐ **Run-Length Vector Storage + Dot Product**: one pointer per run list; multiply the values and take min(remaining counts) of the two current runs *[1Point3Acres (notes), PracHub (solution) as "Implement sparse vector dot product", Aced as "Find the dot product of two sparse vectors."]*

E ☐ ★3 ★ **Is this a valid palindrome?**: compare the two ends; skip characters that are not letters or digits; stop when the pointers cross *[Aced (solution)]*

27. Group by a canonical key

Hash maps and sets (app section 34) 5 questions: 1 easy, 3 medium, 1 hard

Spot it: "anagrams", "same shape", "shifted words", "group strings that are equivalent".

Idea: Compute a key that is identical for every member of one group and different across groups, then put items in a dict of lists under that key. The art is choosing the key: sorted letters or a 26-slot count tuple for anagrams, the differences between neighbouring letters (mod 26) for shifted words, the keyboard row for typeable words.

```
groups = defaultdict(list)
for word in words:
    key = tuple(sorted(word))              # or a 26-count tuple
    groups[key].append(word)
return list(groups.values())

def shift_key(word):                       # shifted-word key
    return tuple((ord(b) - ord(a)) % 26 for a, b in zip(word, word[1:]))
```

Cost: O(total letters * log L) with sorted keys, O(total letters) with count tuples. **Traps:** Lists are not hashable: convert the key to a tuple or string. The count-tuple key avoids the log factor. Decide how to treat case, spaces and Unicode.

Follow-ups: Different keyboard or alphabet. Subsets of letters instead of full anagrams. Words spread over many machines (hash the key to pick the machine).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★2 **Group anagrams**: key = the sorted letters (or a tuple of 26 letter counts) *[Aced (solution), PracHub (solution) as "Group Anagrams"]*

M ☐ ★4 **Check if all substrings are anagrams of words**: key = the tuple of letter counts; keep the keys of the allowed words in a set and test every substring of length 3 or more against it *[PracHub (solution)]*

M ☐ ★5 **Given an array of strings, return groups of "buddy strings," where the strings have the same length and the distance between corresponding characters is the same.**: key = the differences between neighbouring letters, modulo 26 (a shifted word has the same differences) *[Aced, PracHub (solution) as "Group Strings Equivalent under Uniform Alphabet Shifts"]*

E ☐ ★1 ★ **Given an array of strings, return the words that can be typed using letters of the alphabet on only one row of American keyboard.**: key = the keyboard row of each letter; a word qualifies when its letters map to one row *[Aced]*

H ☐ ★3 ★ **Given two strings and two allowed operations, swapping positions of letters and converting all occurrences of one character to another, can the first string be transformed into the second?**: same letters used, and the sorted list of letter frequencies must match: compare set(s) == set(t) and the sorted Counter values *[Aced]*

28. Prefix sums, tier 2: running total plus a dict of earlier totals ★

Arrays and prefix sums (app section 33) 3 questions: 0 easy, 2 medium, 1 hard

Spot it: "number of subarrays with sum k", "multiple of k", "maximum subarray", "equal endpoints".

Idea: A subarray (j, i] has sum prefix[i] - prefix[j]. So for each i you do not scan all j: you ask the dict "have I seen prefix[i] - target before, and how many times?". For a maximum subarray the dict degenerates to one number, the smallest earlier prefix (this is Kadane). For multiples of k, store the first index of each remainder prefix % k.

```
seen = {0: 1};  total = ways = 0
for x in a:                                # count subarrays with sum k
    total += x
    ways += seen.get(total - k, 0)
    seen[total] = seen.get(total, 0) + 1
```

```
best, total, smallest = -inf, 0, 0          # maximum subarray sum
for x in a:
    total += x
    best = max(best, total - smallest)
    smallest = min(smallest, total)
```

Cost: O(n) time, O(n) space for the dict. **Traps:** Seed the dict with prefix 0 (the empty prefix). Update the dict after using it, otherwise an item pairs with itself. With a minimum length, update the "smallest earlier total" one step late.

Follow-ups: All positive numbers (then a sliding window works). Circular array. Return the subarray itself (store indices). Length limits (min or max length).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Find the maximum subarray sum.**: the dict is just one number: the smallest earlier total; the best subarray ends here with total - smallest (Kadane is the same idea) *[**Aced (solution)**, **PracHub (solution) as "Find the Maximum-Sum Contiguous Subarray"**]*

M ☐ ★**2 Solve array-sum and city-community problems**: count subarrays summing to k with the pattern template above; the second task is union-find over building cells *[**PracHub (solution)**]*

H ☐ ★ **Given an integer array nums and an integer k, return true if nums has a subarray of at least two elements whose sum is a multiple of k.**: key = total % k; store the first index; a hit at least 2 positions back means a subarray whose sum is a multiple of k *[**Aced (solution)**]*

DAY 3

Strings, design and trees

Patterns 29 to 43 (15 patterns)

Strings and digits, grids, simulations and caches, parent pointers, random structures, nested parsing, streaming, trees, serialisation and tries.

29. Nested structures: a stack of contexts, or recursion with a cache ★

Stacks and queues (app section 27) 5 questions: 0 easy, 5 medium, 0 hard

Spot it: "decode string 3[a2[c]]", "evaluate expression", "nested list weight", "expand placeholders".

Idea: Each open bracket starts a new context, each close bracket finishes it and combines the result into the one below. Keep a stack of (partial result, pending multiplier) and push on open, pop on close. For recursive definitions (templates referring to keys), resolve with a cache and a "currently resolving" set so a cycle is detected instead of looping forever.

```
stack, current, number = [], '', 0
for ch in s:
    if ch.isdigit():     number = number * 10 + int(ch)
    elif ch == '[':      stack.append((current, number));  current, number = '', 0
    elif ch == ']':
        previous, times = stack.pop();  current = previous + current * times
    else:                current += ch

def resolve(key, cache, visiting):
    if key in cache: return cache[key]
    if key in visiting: raise ValueError('cycle')
    visiting.add(key)
    cache[key] = expand(values[key], cache, visiting);  visiting.discard(key)
    return cache[key]
```

Cost: O(output size) for decoding, O(n) for parsing. **Traps:** Multi-digit numbers. The decoded string may be enormous (compute its length or the k-th character without building it). Unknown keys and unmatched brackets need a defined behaviour.

Follow-ups: Recursion depth (iterate with an explicit stack). Division and precedence in expressions. Variables and variadic functions. Inverse weighting (deeper counts less).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Template Variable Expansion with Follow-Ups**: resolve(name) with a cache and a 'being resolved' set; meeting a name that is being resolved is a cycle *[1Point3Acres (notes), PracHub (solution) as "Expand Recursive Key-Value Placeholders in a Template"]*

M ☐ ★2 **Nested List Weight Sum (LC 339) with Follow-Ups**: depth counter: '[' adds 1, ']' subtracts 1; each number adds value * depth (read digits and a minus sign as you go) *[1Point3Acres (notes), PracHub (solution) as "Compute a Depth-Weighted Sum of a Nested List"]*

M ☐ ★3 **Expand Recursive Key-Value Placeholders in a Template**: regex-find {key}, expand each key recursively with a cache and a visiting set to stop cycles *[PracHub (solution), 1Point3Acres (notes) as "Template Variable Expansion with Follow-Ups"]*

M ☐ ★4 **Decode String**: the pattern template above *[1Point3Acres (notes), PracHub (solution) as "Decode an encoded string"]*

M ☐ ★5 **Nested Function Expression Evaluator**: tokenise, then a recursive parse(i): a number returns itself; a name reads its arguments until ')' and applies the function *[1Point3Acres (notes), PracHub (solution) as "Evaluate a Lisp-Style Expression Language with let, add and mult Scoping"]*

30. Tree recursion, tier 1: visit a node, ask children, report upward ★

Binary trees and BSTs (app section 3) 8 questions: 3 easy, 4 medium, 1 hard

Spot it: "height", "diameter", "right side view", "level order", "delete leaves", "evaluate expression tree".

Idea: Write a function that takes a node and returns one number or object about the subtree (height, value, whether it is valid). Call it on the children first (post-order), combine, and update a global answer if the answer is not the return value. Diameter: return 1 + max(left, right) but update best with left + right. For per-level problems pass depth down and store the first or last node seen at each depth.

```
best = 0
def height(node):
    global best
    if node is None: return 0
    left, right = height(node.left), height(node.right)
    best = max(best, left + right)          # diameter in edges
    return 1 + max(left, right)

def right_view(node, depth, view):
    if node is None: return
    if depth == len(view): view.append(node.val)
    right_view(node.right, depth + 1, view)     # right first
    right_view(node.left, depth + 1, view)
```

Cost: O(n) time, O(height) stack space. **Traps:** Say whether lengths count nodes or edges. A chain of 100,000 nodes overflows Python recursion: convert to an explicit stack. Return and update are two different numbers.

Follow-ups: N-ary tree (loop over children). Is the tree balanced (return height or -1). Return the path itself (store the best endpoints). Delete subtrees that match a predicate.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front

of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Expression Tree Single-Leaf Mutation**: post-order evaluation (children first, then apply the operator); to test one leaf flip, re-evaluate with that leaf changed *[1Point3Acres (notes), PracHub (solution) as "Boolean Expression Tree with Leaf Flips"]*

H ☐ ★2 **Repeatedly Delete Leaves Matching a Target Value From a Binary Tree**: set node.left = dfs(node.left) and node.right likewise; return None when the node became a leaf with the target value *[PracHub (solution)]*

M ☐ ★3 **Given the root of a binary tree, return the length of the diameter of the tree.**: best = max(best, left + right); return 1 + max(left, right) (the diameter) *[Aced (solution), PracHub (solution) as "Compute binary tree diameter"]*

M ☐ ★4 **Longest Consecutive Path in Binary Tree**: pass the parent's value and the run length down; extend the run when value == parent + 1, else restart at 1 *[1Point3Acres (notes), PracHub (solution) as "Find the Longest Consecutive Path in a Binary Tree"]*

E ☐ **Find right-side view of binary tree**: pass depth down; the first node you reach at a new depth (visiting right before left) is the one you can see *[PracHub (solution)]*

E ☐ **Given two trees, determine if the pre-order traversals of two trees are same or not.**: walk both trees in preorder with an explicit stack (push right, then left) and compare value by value; stop at the first difference *[Aced, PracHub (solution) as "Compare Binary Trees at Matching Structural Positions"]*

M ☐ **Implement zigzag level-order traversal**: collect values by depth; reverse every odd level (or fill each level from the other end) *[PracHub (solution)]*

E ☐ ★ **Calculate the height of a binary tree.**: no update needed; return 1 + max(left, right) and read the value at the root *[Aced (solution)]*

31. Tree DP: build a parent order once, add up from the leaves

Binary trees and BSTs (app section 31) 5 questions: 0 easy, 3 medium, 2 hard

Spot it: "total size of a folder", "sum of distances from every node", "time to inform everyone", "cheapest cut".

Idea: Get an order where parents come before children (BFS or DFS from the root). Sweep it backwards to compute subtree values (size, sum, best) from children, then sweep it forwards to push information down. Distance sums for every root (rerooting): down[v] is the sum inside the subtree of v, and ans[child] = ans[parent] - size[child] + (n - size[child]).

```
order, parent = [root], {root: None}
for node in order:                          # BFS order, parents first
    for child in children[node]:
        parent[child] = node;  order.append(child)

size = {v: 1 for v in order};  down = {v: 0 for v in order}
for v in reversed(order):                   # leaves to root
    if parent[v] is not None:
        size[parent[v]] += size[v]
        down[parent[v]] += down[v] + size[v]

answer = {root: down[root]}
for v in order[1:]:                         # root to leaves (rerooting)
    answer[v] = answer[parent[v]] - size[v] + (n - size[v])
```

Cost: O(n), and no recursion so a long chain is safe. **Traps:** Prove the reroot formula with a two node tree. Hard links count a shared item once (seen set). Serve the child with the longest time first when informing in rounds.

Follow-ups: Weighted edges. Maximum distance instead of sum. Which edges to cut. The tree is queried often while files change (store subtree sizes and update upward).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★2 **Compute File System Total Size**: add each file's size while walking down from the root with a stack; a shared directory is counted once (seen set) *[1Point3Acres (notes), PracHub (solution) as "Compute Entity Size and Validate a Hierarchical File System"]*

M ☐ **Given a large n-ary tree where the root node can pass information to one child per iteration, and in subsequent iterations, nodes can pass information to one child while the root continues with others, find the minimum number of iterations to reach all nodes.**: t[v] = max over children (rank among children sorted by time, longest first, + t[child]); read t at the root *[Aced, PracHub (solution) as "Validate a Hierarchy and Compute Propagation Time"]*

H ☐ **Compute distance-sum from every tree node**: down[v] = sum of distances to v's subtree; then a forward sweep: ans[child] = ans[parent] - size[child] + (n - size[child]) *[PracHub (solution), 1Point3Acres (notes) as "Tree Distance Sum (Re-root DP)"]*

H ☐ **Tree Distance Sum (Re-root DP)**: exactly the same two sweeps (this is the re-rooting version of the question) *[1Point3Acres (notes), PracHub (solution) as "Compute distance-sum from every tree node"]*

M ☐ ★1 ★ **Minimum-Cost Tree Disconnection**: backwards sweep with two numbers per node: the best cost when the node is cut off, and the cost if it is free *[1Point3Acres (notes)]*

32. Trie: nested dicts, one character per level ★

Tries (app section 17) 8 questions: 0 easy, 7 medium, 1 hard

Spot it: "autocomplete", "words with a prefix", "longest match tokenizer", "URL router", "dictionary substrings".

Idea: Each node is a dict from character to child; a marker (END) says a word stops here. Insert walks and creates; search walks and checks END; prefix queries walk to the prefix node and collect everything below. For longest-match replacement, walk as far as

the trie goes from each text position, remember the last END seen, and jump past it. A router uses one level per path segment with a literal child and one wildcard child.

```
END = '$'
def insert(root, word):
    node = root
    for ch in word:
        node = node.setdefault(ch, {})
    node[END] = True

def longest_match(root, text, start):
    node, last = root, -1
    for i in range(start, len(text)):
        if text[i] not in node: break
        node = node[text[i]]
        if END in node: last = i
    return last                              # -1 when nothing matches
```

Cost: O(length of word) per operation, O(total characters) space. **Traps:** END must be distinguishable from a real character. Deleting needs pruning of empty nodes. For prefix lookups on a static list, sorting plus bisect is lighter than a trie.

Follow-ups: Top 3 completions (store them per node). Wildcard "." search (recurse over all children). Fuzzy match with one error. Huge dictionaries (sharding, compressed trie).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **URL Request Handler Router**: one trie level per URL segment; a node has literal children and one wildcard child; try the literal first *[1Point3Acres (notes), PracHub (solution) as "Implement an HTTP Router That Resolves Handlers for Paths with Named Parameters"]*

M ☐ ★2 **Build Prefix Lookup with a Trie**: build_trie, then walk to the prefix node and collect every word below it *[PracHub (solution), 1Point3Acres (notes) as "Trie Prefix Search / Autocomplete"]*

M ☐ ★3 **Trie Prefix Search / Autocomplete**: walk to the prefix node; return the words under it (autocomplete) *[1Point3Acres (notes), PracHub (solution) as "Build Prefix Lookup with a Trie"]*

M ☐ ★4 **Greedy Longest-Match Dictionary Tokenizer**: the same longest-match walk, with a priority rule to choose between equal-length matches *[1Point3Acres (notes), PracHub (solution) as "Implement Longest-Match Text Replacement"]*

M ☐ ★5 **Check if all substrings are dictionary words**: for every start walk the trie; a substring of length 3 or more that leaves the trie or lacks END fails the check *[PracHub (solution)]*

M ☐ ★6 **Return dictionary words matching a prefix**: no trie needed: sort the words and bisect to the first match, then read forward while startswith(prefix) *[PracHub (solution)]*

M ☐ ★7 **Implement Trie Insert and Exact-Word Search**: insert is build_trie; search is walk and then END in node *[PracHub (solution)]*

H ☐ **Implement Longest-Match Text Replacement**: at each text position walk the trie as far as it goes and remember the last END; replace the longest match and jump past it *[PracHub (solution), 1Point3Acres (notes) as "Greedy Longest-Match Dictionary Tokenizer"]*

33. A dict for lookup plus an ordered container for who is next ★

Data structure and object design (app section 18) 5 questions: 1 easy, 3 medium, 1 hard

Spot it: "LRU cache", "waitlist", "round robin scheduler", "remove the cheapest", "undo / redo".

Idea: A dict gives O(1) access by id; a second structure gives order: a deque for first-in-first-out, a doubly linked list or OrderedDict for recency, a heap for cheapest. Deleting from the ordered side is expensive, so use lazy deletion: only remove the id from the dict, and when you pop from the order structure, skip entries that are no longer alive. Undo and redo are two stacks; a batch is one stack entry.

```
class LRU:
    def __init__(self, capacity):
        self.capacity, self.data = capacity, OrderedDict()
    def get(self, key):
        if key not in self.data: return -1
        self.data.move_to_end(key)          # most recent at the end
        return self.data[key]
    def put(self, key, value):
        self.data[key] = value
        self.data.move_to_end(key)
        if len(self.data) > self.capacity:
            self.data.popitem(last=False)   # evict least recent

waiting = deque();  alive = {}              # lazy deletion waitlist
def next_party():
    while waiting and waiting[0] not in alive: waiting.popleft()
    return waiting.popleft() if waiting else None
```

Cost: O(1) per operation for LRU; O(log n) with a heap. **Traps:** Every access must refresh recency. Lazy deletion needs a unique ticket per insert (a name can rejoin). Say what the order is (arrival, price, deadline).

Follow-ups: Thread safety (a lock). TTL expiry. LFU instead of LRU (frequency buckets). Best fit instead of first fit (segment tree

over sizes). Limit the undo history.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

H ☐ ★1 **Implement LRU Cache.**: dict key -> node plus a doubly linked list; every get/put moves the node to the front; evict from the back (or use OrderedDict.move_to_end) *[Aced (solution), PracHub (solution) as "Implement an LRU cache"]*

M ☐ ★2 **Implement Python's itertools.tee**: one deque per consumer; whoever runs out pulls the next value from the source and appends it to every queue *[1Point3Acres (notes), PracHub (solution) as "Implement a Lazy Tee Iterator in Python"]*

M ☐ ★3 **Design a restaurant waitlist system**: a segment tree of party sizes gives the first entry that fits in O(log n); the dict maps ids to slots *[PracHub (solution), 1Point3Acres (notes) as "Restaurant Waitlist: join / delete / find_first_match"]*

E ☐ **Implement Batched Undo/Redo Layer**: two stacks (undo, redo) of batches; a batch records (old value, new value) per property; undo applies the old values and moves the batch *[PracHub (solution)]*

M ☐ **Restaurant Waitlist: join / delete / find_first_match**: a deque per party size with tickets; delete only removes the name from the dict; find_first_match skips dead tickets *[1Point3Acres (notes), PracHub (solution) as "Design a restaurant waitlist system"]*

34. Build or transform a grid with index arithmetic ★

Grids and simulation (app section 42) 3 questions: 0 easy, 3 medium, 0 hard

Spot it: "rotate a matrix", "range sum in a 2D matrix", "draw a diamond", "fill in snake order", "block minimum".

Idea: Derive the formula for where each cell goes and write it directly. Rotating clockwise: new[r][c] = old[rows - 1 - c][r], or in Python zip(*matrix[::-1]). A 2D prefix table p[r+1][c+1] = cell + p[r][c+1] + p[r+1][c] - p[r][c] answers any rectangle with four lookups by inclusion and exclusion. Block grouping uses (r // k, c // k) as the key.

```
prefix = [[0] * (cols + 1) for _ in range(rows + 1)]
for r in range(rows):
    for c in range(cols):
        prefix[r + 1][c + 1] = (matrix[r][c] + prefix[r][c + 1]
                                + prefix[r + 1][c] - prefix[r][c])
def rect_sum(r1, c1, r2, c2):              # inclusive corners
    return (prefix[r2 + 1][c2 + 1] - prefix[r1][c2 + 1]
            - prefix[r2 + 1][c1] + prefix[r1][c1])

rotated = [list(row) for row in zip(*matrix[::-1])]
```

Cost: O(rows * cols) to build, O(1) per query. **Traps:** Do not rotate non-square matrices in place. Check the formula on a 2 by 3 example by hand. Blocks that do not tile the grid exactly need a decision (pad or drop).

Follow-ups: Rotate by k quarter turns (k % 4). Cell updates between queries (2D Fenwick tree). Detect impossible recipes by cycle (topological sort).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Connected Crop Arrangement on a Grid**: fill the cells in snake order (left to right, then right to left) so equal crops stay in one connected strip *[1Point3Acres (notes), PracHub (solution) as "Construct connected crop layout and safe paths"]*

M ☐ **Solve matrix groups and recipe inventory**: block (r // k, c // k) holds the minimum of every cell that maps to it; the recipe part is a topological sort *[PracHub (solution)]*

M ☐ **Rotate Any m×n 2D Matrix**: clockwise: new[r][c] = old[rows - 1 - c][r] on a cols-by-rows result (or zip(*matrix[::-1])) *[1Point3Acres (notes), PracHub (solution) as "Rotate a Matrix In Place", Aced as "Rotate a 2D matrix."]*

35. Look at the input once and keep only a small summary ★

Streaming, scale and concurrency (app section 25) 6 questions: 0 easy, 4 medium, 2 hard

Spot it: "file does not fit in memory", "stream", "minimise the number of oracle calls", "two logs".

Idea: Decide what the smallest summary is that still answers the question: a running maximum, a set of ids from the smaller file, a size-k heap, a count per key. Stream the large input past it. When a test oracle is expensive, use group testing: test a whole group in one call and split only the groups that fail, or binary search a prefix to find the first failure.

```
best = None                                # maximum in a huge file
with open(path) as handle:
    for line in handle:
        value = int(line)
        if best is None or value > best: best = value

small_ids = {line.strip() for line in open(small_file)}
shared = [line for line in open(big_file) if line.strip() in small_ids]
```

```
def find_bad(group):                        # group testing
    if not run_test(group): return []
    if len(group) == 1: return group
    mid = len(group) // 2
    return find_bad(group[:mid]) + find_bad(group[mid:])
```

Cost: O(n) time, memory proportional to the summary, not the input. **Traps:** Never call `.read()` or `.readlines()` on a huge file. Say how you would handle both files being too big (partition by hash of the key, process partition by partition). Edge cases: empty file, blank lines, ties.

Follow-ups: Scale to many machines (map and reduce by key). Failures and retries. Top K instead of max (heap of size K). Is it I/O bound or CPU bound?

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

- M ☐ ★1 **Streaming Points: Emit K Within Distance**: keep the points in a sorted list; each new point is inserted with bisect and only windows around it are re-checked *[1Point3Acres (notes), PracHub (solution) as "Streaming Points: Remove Any Pair Within a Distance"]*
- M ☐ **Two-Day Log Intersection under Memory Limit**: dict object -> set of event codes for each day; keep the objects that appear on both days with at least two distinct codes *[1Point3Acres (notes), PracHub (solution) as "Find shared objects across two log files"]*
- M ☐ **Find shared objects across two log files**: load only the smaller file's object ids into a set, then stream the big file and keep ids found in the set *[PracHub (solution), 1Point3Acres (notes) as "Two-Day Log Intersection under Memory Limit"]*
- M ☐ **Minimize calls to find all bad test pairs**: the same group testing: put a whole group in one call and split only the groups that fail *[PracHub (solution), 1Point3Acres (notes) as "Find All Bad Pairs Using runTest Oracle"]*
- H ☐ **Find All Bad Pairs Using runTest Oracle**: for each test, binary search the prefix of earlier tests where runTest first fails; that test is a bad partner; repeat until none is left *[1Point3Acres, PracHub (solution) as "Minimize calls to find all bad test pairs"]*
- H ☐ **Find maximum value in huge text file**: the pattern template above; the file is never loaded as a whole *[PracHub (solution)]*

36. Simulate a machine: a class with state and one branch per event ★

Grids and simulation (app section 22) 5 questions: 0 easy, 4 medium, 1 hard

Spot it: "design tic-tac-toe", "elevator system", "iterator over lists", "game of life", "scorekeeper".

Idea: Name the state (board, floor, queues, counters), write one method per kind of input, and make every method update the state in one place. For tic-tac-toe, keep counters per player for rows, columns and both diagonals so a win check is O(1). For game of life in place, encode old and new state in two bits of each cell so every cell is computed from the old state, then shift.

```
class TicTacToe:
    def __init__(self, n):
        self.n = n
        self.rows = defaultdict(int);  self.cols = defaultdict(int)
        self.diag = defaultdict(int);  self.anti = defaultdict(int)
    def move(self, r, c, player):
        for table, key in ((self.rows, (player, r)), (self.cols, (player, c))):
            table[key] += 1
            if table[key] == self.n: return player
        if r == c:                self.diag[player] += 1
        if r + c == self.n - 1:   self.anti[player] += 1
        if self.n in (self.diag[player], self.anti[player]): return player
        return 0
```

Cost: O(1) per move for tic-tac-toe, O(cells) per step for life. **Traps:** Reject moves on a taken cell if the spec says so. Neighbour counts in life must read the old bit (value & 1). Elevator: define the cost of assigning a request to a car.

Follow-ups: Undo the last move (stack of moves). Infinite or sparse board (set of live cells). Wrap-around edges. Fairness, starvation and failures for schedulers.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

- M ☐ ★1 **Tic-Tac-Toe System with K Players and 3-in-a-Row Win**: store the board; after each move count the run through the new cell in four directions and compare with the target length *[1Point3Acres (notes), PracHub (solution) as "Implement Multi-Player Tic-Tac-Toe"]*
- H ☐ ★2 **Simulate In-Place Cellular State Updates**: encode old and new state in two bits of each cell so every cell is computed from the old state; shift at the end *[PracHub (solution)]*
- M ☐ ★3 **Design Tic-Tac-Toe With K Players**: per player Counters for rows, columns and both diagonals; a move wins when a counter reaches n *[PracHub (solution), 1Point3Acres (notes) as "Tic-Tac-Toe System with K Players and 3-in-a-Row Win"]*
- M ☐ **Implement iterator merging lists with filtering**: keep a position per list and a turn counter; advance() looks for the next unblocked, unseen value and stores it for next()/has_next() *[PracHub (solution)]*
- M ☐ **Design an elevator control system**: each car keeps floor, direction and stops; a request goes to the car with the lowest cost; step() moves every car one floor *[PracHub (solution)]*

37. Digits and characters: peel one at a time and carry ★

Math and bit tricks (app section 24) 4 questions: 0 easy, 3 medium, 1 hard

Spot it: "add two big numbers", "palindrome number", "monotone digits", "caesar cipher", "k-th permutation".

Idea: Work with one digit at a time using n % 10 (last digit) and n // 10 (drop it), or with ord/chr arithmetic for letters. Addition keeps a carry: total = a + b + carry, digit = total % 10, carry = total // 10. For the k-th permutation, divide k-1 by (n-1)!, (n-2)!, ... to choose each digit from the list of unused digits (factorial number system).

```
def add_digits(a, b):                       # lists, most significant first
    i, j, carry, out = len(a) - 1, len(b) - 1, 0, []
    while i >= 0 or j >= 0 or carry:
        total = carry + (a[i] if i >= 0 else 0) + (b[j] if j >= 0 else 0)
        out.append(total % 10);  carry = total // 10
        i -= 1;  j -= 1
    return out[::-1]

def shift(ch, k):                           # caesar
    return chr((ord(ch) - 97 + k) % 26 + 97)
```

Cost: O(number of digits). **Traps:** A final carry creates a new leading digit. Negative k works with % in Python. A number palindrome only needs the second half reversed.

Follow-ups: Numbers as linked lists (least significant first makes addition easy). Another base. Break a cipher given only ciphertext (frequency analysis). Position of a permutation (inverse of the factorial method).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Check Whether Digits Are Monotone Increasing**: peel digits from the right; each digit must not exceed the previous one you saw *[PracHub (solution)]*

H ☐ ★**2 Check Digits and Combine Ranges**: palindrome number: reverse only the second half with n % 10 and n // 10 and compare with the first half *[PracHub (solution)]*

M ☐ ★**3 Solve several coding interview problems**: the big-number addition on strings, with int(ch) per digit; the other parts are list-meeting and overlap sweeps *[PracHub (solution)]*

M ☐ **Add two big integers from digit lists**: the pattern template above *[PracHub (solution)]*

38. Split into pieces, fix each piece, join (and word wrap) ★

Strings and parsing (app section 43) 5 questions: 1 easy, 2 medium, 2 hard

Spot it: "snake_case to camelCase", "format a number", "justify text", "wrap at width".

Idea: Break the text at its separators, transform each piece with a tiny function, and join. For wrapping, add words greedily while current length + 1 + word length fits in the width, otherwise start a new line. For full justification, spread the spaces of a line over its gaps (the leftmost gaps get the remainder) and left-align the last line.

```
def to_camel(name):
    first, *rest = name.split('_')
    return first + ''.join(p.capitalize() for p in rest)

lines, current = [], []
for word in words:
    if current and len(' '.join(current + [word])) > width:
        lines.append(current);  current = []
    current.append(word)
lines.append(current)
```

Cost: O(total characters). **Traps:** Leading, trailing and repeated separators. A word longer than the width. Do not use += on strings in a long loop; collect pieces and join once.

Follow-ups: Convert back (camel to snake). Locale grouping (12,34,567). Words with pixel widths. Minimise raggedness instead of greedy packing (that is a DP over line breaks).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Snake Case to Camel Case with Repository-Wide Refactor**: the same function applied to every identifier found in the files *[1Point3Acres (notes), PracHub (solution) as "Convert Valid Snake Case to Lower Camel Case"]*

H ☐ ★**2 Convert Valid Snake Case to Lower Camel Case**: reformat(): split on '_', keep the first piece, capitalise the others *[PracHub (solution), Aced as "Define the scope and requirements for a snake_case to camelCase converter function.", 1Point3Acres (notes) as "Snake Case to Camel Case with Repository-Wide Refactor"]*

E ☐ ★**3 Implement string formatting and request tracking tasks**: split, fix each piece, join; the request-tracking part is a small dict *[PracHub (solution)]*

M ☐ **Google Doc Line Wrap (Word Wrap with Whitespace and Newlines)**: count lines: go paragraph by paragraph, add a word if it fits (current + 1 + length <= width), else start a new line *[1Point3Acres (notes), PracHub (solution) as "Format words into wrapped/justified lines"]*

H ☐ **Implement a function to justify a list of words within a given width.**: pack_lines for the words, then spread the extra spaces over the gaps (the left gaps get the remainder); the last line is left-aligned *[Aced, PracHub (solution) as "Format words into wrapped/justified lines"]*

39. Walk two strings at once with two indices

Strings and parsing (app section 21) 6 questions: 2 easy, 3 medium, 1 hard

Spot it: "is a subsequence", "lexicographic order in a custom alphabet", "can board A become board B", "similar sentences".

Idea: Keep one index per string and advance them by a rule: advance both on a match, advance only the text pointer on a mismatch (subsequence), skip blank squares in both (board transforms). A custom alphabet order is just a rank dict: compare the ranks at the first differing position. Compress words to (letter, run length) pairs when stretched letters matter.

```
def is_subsequence(needle, text):
    i = 0
    for ch in text:
        if i < len(needle) and ch == needle[i]:
            i += 1
    return i == len(needle)

rank = {c: r for r, c in enumerate(alphabet)}
def in_order(a, b):                          # alien dictionary order
    for x, y in zip(a, b):
        if x != y: return rank[x] < rank[y]
    return len(a) <= len(b)                  # a prefix comes first
```

Cost: O(n + m) per pair of strings. **Traps:** The prefix rule at the end of the comparison. Skipped squares must not shift the order of the remaining letters. Pieces that move in one direction only (L left, R right) cannot cross each other.

Follow-ups: Whole list of words in an alien order (compare neighbours). Case-insensitive or accented text. Minimum number of moves. Millions of words and many queries (compress once, index by shape).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

H ☐ ★**2 Find Dictionary Matches for a Jammed String**: compress every word to (letter, run length); a query matches when the letters agree and each run is equal or stretched to at least 3 *[PracHub (solution)]*

M ☐ ★**3 Can board states be transformed?**: skip the X squares in both strings; the remaining letters must match in order, and L may only move left, R only right *[PracHub (solution)]*

E ☐ **Determine if two words are sorted in lexicographic order.**: walk both words; at the first differing letter compare the letters' ranks in the given alphabet; if one word ends first, it must be the shorter *[Aced, PracHub (solution) as "Verify Ordering Under a Scrambled Alphabet"]*

M ☐ **Decide string transform with directional tokens and walls**: the same walk with dots as blanks and the movement rule from the statement (walls block) *[PracHub (solution)]*

E ☐ ★**1** ★ **Return all unique substrings of a given string.**: sort all suffixes; the unique substrings are all prefixes minus the shared prefixes of neighbouring suffixes *[Aced]*

M ☐ ★ **Determine if two sentences are similar.**: zip the two sentences; each pair of words must be equal or listed as similar (store pairs in both directions) *[Aced (solution)]*

40. Parent pointers: store who is above each node and climb ★

Data structure and object design (app section 41) 4 questions: 1 easy, 3 medium, 0 hard

Spot it: "lowest common ancestor with parent links", "validate a tree", "employee hierarchy", "is this a root".

Idea: Store parent[node] (and sometimes a set of children). Climbing is a loop following parents. Lowest common ancestor with parent links: advance two pointers, and when one falls off the top restart it at the other node; they meet at the ancestor after at most len(a) + len(b) steps. Validating a parent array: climb from every node with three states (unseen, on this walk, verified); meeting your own walk is a cycle, and there must be exactly one root.

```
def lca(a, b):
    p, q = a, b
    while p is not q:
        p = p.parent if p else b
        q = q.parent if q else a
    return p

state = {}                                   # 0 unseen, 1 on this walk, 2 verified
def has_cycle(start):
    walk = []
    node = start
    while node is not None and state.get(node, 0) == 0:
        state[node] = 1;  walk.append(node);  node = parent[node]
    bad = node is not None and state.get(node) == 1
```

```
    for n in walk: state[n] = 2
    return bad
```

Cost: O(depth) per climb, O(n) for validating all nodes. **Traps:** The root has no parent (None). Check for exactly one root and for nodes that are not reachable. Deleting a manager means re-attaching the reports.

Follow-ups: Many LCA queries on a static tree (binary lifting or Euler tour). Move an employee and reject cycles. Nodes might not share a tree (return None).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Lowest Common Ancestor of a Binary Tree III**: two pointers climb; when one falls off the top it restarts at the other node, so both meet at the common ancestor *[HackTheRounds #56 (solution), PracHub (solution) as "Find LCA With Parent Pointers"]*

M ☐ ★**2 Employee Hierarchy Add / Delete**: manager[emp] and a set of reports per manager; delete re-attaches the reports to the deleted person's manager *[1Point3Acres (notes), PracHub (solution) as "Implement Employee Management and Expression Evaluation", Aced (solution) as "Employee Hierarchy"]*

E ☐ ★**3 Tree Node.**: the SQL version of the same idea: p_id is the parent pointer, NULL means root, being someone's parent means inner, else leaf *[Aced, PracHub (solution) as "Classify Tree Nodes in SQL"]*

M ☐ ★**4 Validate parent array forms a tree**: climb from every node with three states (unseen, on this walk, verified); meeting your own walk is a cycle; there must be exactly one root *[PracHub (solution)]*

41. Random pick without repeats: swap with the last and pop ★

Data structure and object design (app section 37) 2 questions: 0 easy, 1 medium, 1 hard

Spot it: "random set with O(1) delete", "shuffle", "unique random generator", "no repeat within k songs".

Idea: To delete from the middle of a list in O(1), swap the victim with the last element, update the position dict, and pop. Random pick is then a random index into the list. A shuffle is Fisher-Yates: for i from the end, swap a[i] with a random a[j] where j <= i. A huge pool 0..n-1 can be lazy: keep a dict of only the swapped slots.

```
class RandomSet:
    def __init__(self):  self.items, self.where = [], {}
    def add(self, x):
        if x in self.where: return False
        self.where[x] = len(self.items);  self.items.append(x);  return True
    def remove(self, x):
        if x not in self.where: return False
        i, last = self.where[x], self.items[-1]
        self.items[i] = last;  self.where[last] = i
        self.items.pop();  del self.where[x];  return True
    def pick(self):  return random.choice(self.items)

for i in range(len(a) - 1, 0, -1):          # Fisher-Yates
    j = random.randint(0, i);  a[i], a[j] = a[j], a[i]
```

Cost: O(1) per add, remove and pick. **Traps:** Update the position of the moved element before deleting. A shuffle that picks j from the whole array is biased. Raise or return a sentinel when the pool is empty.

Follow-ups: Duplicates allowed (map value to a set of positions). Reset in O(1) (keep a counter, do not rebuild). Only a coin flip API exists (rejection sampling). Thread safety.

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★**1 Design set with O(1) random access**: add appends and records the position; delete swaps with the last item and pops; get_random picks a random index *[PracHub (solution)]*

H ☐ **Shuffle array using random integer API**: Fisher-Yates: for i from the end, j = rand_int(0, i), swap a[i] and a[j] *[PracHub (solution)]*

42. Tree recursion, tier 2: pass a rule down, return a node or yes/no up ★

Binary trees and BSTs (app section 36) 3 questions: 0 easy, 3 medium, 0 hard

Spot it: "valid BST", "lowest common ancestor", "in-order successor", "mode in a tree".

Idea: Pass information down as arguments (the allowed value window low and high for BST validation) and return a verdict or a node up. A BST in-order walk is sorted, so equal values are neighbours (running count gives the mode) and the successor of x is found by descending: go left after a hit, right otherwise. Lowest common ancestor: return the node if it is p or q; a node that receives an answer from both sides is the ancestor.

```
def is_bst(node, low=float('-inf'), high=float('inf')):
    if node is None: return True
```

```
    if not low < node.val < high: return False
    return is_bst(node.left, low, node.val) and is_bst(node.right, node.val, high)

def lca(node, p, q):
    if node is None or node is p or node is q: return node
    left, right = lca(node.left, p, q), lca(node.right, p, q)
    return node if left and right else left or right
```

Cost: O(n) for the walks, O(height) for a BST descent. **Traps:** Compare against the whole window, not only the parent. Decide how duplicates are placed. p or q may be missing from the tree (verify with a count).

Follow-ups: Not a BST (generic traversal). O(1) space (Morris traversal). Recover a BST with two swapped nodes. Deep trees (iterative).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Lowest Common Ancestor**: dfs returns the node if it is p or q; a node that gets an answer from both sides is the ancestor *[HackTheRounds #631 (solution), PracHub (solution) as "Lowest Common Ancestor in an Explicit Binary Tree", Aced (solution) as "Find the lowest common ancestor (LCA) of two nodes in a binary tree."]*

M ☐ **Find mode in a trinary search tree**: in-order walk (left, middle, right); equal values are neighbours, so one running count finds the mode(s) *[PracHub (solution)]*

M ☐ **Determine if a given binary tree is a binary search tree (BST).**: check(): pass lo and hi bounds down; fail when a value leaves its allowed window *[Aced (solution), PracHub (solution) as "Validate a Binary Search Tree"]*

43. Build and serialize: write out a structure, read it back with the same walk ★

Binary trees and BSTs (app section 38) 3 questions: 0 easy, 2 medium, 1 hard

Spot it: "serialize / deserialize", "construct tree from preorder and inorder", "print a directory tree", "encode a list of strings".

Idea: Pick a format that makes decoding unambiguous and decode with the same traversal order you used to encode. Strings: write length, a separator, then the text, so any character may appear inside. Preorder plus inorder: the first preorder item is the root and its position in the inorder list splits left from right. N-ary tree: write the value and the child count. Graph: write one line per node with the ids of neighbours, create all nodes first, then link.

```
def encode(strings):
    return ''.join(f'{len(s)}#{s}' for s in strings)
def decode(text):
    out, i = [], 0
    while i < len(text):
        j = text.index('#', i);  length = int(text[i:j])
        out.append(text[j + 1:j + 1 + length]);  i = j + 1 + length
    return out

def build(preorder, inorder):
    index = {v: i for i, v in enumerate(inorder)}
    it = iter(preorder)
    def make(low, high):
        if low > high: return None
        node = Node(next(it));  mid = index[node.val]
        node.left, node.right = make(low, mid - 1), make(mid + 1, high)
        return node
    return make(0, len(inorder) - 1)
```

Cost: O(n) time and space. **Traps:** Build left before right when consuming a preorder iterator. Values may repeat in preorder plus inorder construction (ambiguous). Graph: handle cycles with a seen map.

Follow-ups: Postorder plus inorder. Level order encoding. Binary encoding for compactness. Why not JSON (length prefix works for any bytes).

Questions in this section (E easy, M medium, H hard). Listed with PracHub solutions first, then 1Point3Acres, then the rest. Tick the small box in front of a question when you have done it (click it, then save the file). The numbered stars mark the 150 questions to learn in the whole book, chosen by how recently Google was reported asking them (every date found, from all sources): star 1 is the most recent in its pattern. Solve them first, in the order of the star number, because each carries an idea this pattern teaches and you can reuse it on a brand-new question. The big orange stars are those numbered ones. A small red star marks a question with no PracHub solution:

M ☐ ★1 **Construct a binary tree given its preorder and inorder traversals.**: the first preorder value is the root; its position in a value-to-index map of the inorder list splits left from right *[Aced (solution), PracHub (solution) as "Construct tree from inorder & postorder"]*

M ☐ **Write functions to serialize and deserialize a list of strings.**: write len + '#' + text for each string; read the length, then take exactly that many characters *[Aced (solution), PracHub (solution) as "Implement encode/decode for list of strings"]*

H ☐ **Implement a decode and encode function of a general N-ary tree.**: write value and child count for each node (preorder); read the count to know how many children to build *[Aced, PracHub (solution) as "Encode an N-ary tree as a string and decode it back to the identical tree"]*

Appendix: algorithms hiding inside single questions

These show up once or twice (often as a follow-up) but are worth knowing by name.

Fenwick tree (binary indexed tree)

Prefix sums with point updates in $O(\log n)$. Used for count of smaller numbers after self, dynamic range sums, and "array changes between queries" follow-ups. Index from 1; the lowest set bit of an index says how many items its cell covers.

```
class Fenwick:
    def __init__(self, n):  self.tree = [0] * (n + 1)
    def add(self, i, delta):                # i is 1-based
        while i < len(self.tree):
            self.tree[i] += delta;  i += i & -i
    def prefix(self, i):                    # sum of items 1..i
        total = 0
        while i > 0:
            total += self.tree[i];  i -= i & -i
        return total
# count of smaller numbers after self: go from the right, answer = prefix(rank - 1), then add(rank, 1)
```

Segment tree (first slot that fits, range minimum, range updates)

Store a summary of each range in a tree over positions. The waitlist question keeps the maximum party size per range, so "first group of at most s" descends left when the left half has a big enough slot. Cost $O(\log n)$ per update and query.

```
class MaxTree:
    def __init__(self, values):
        self.n = 1
        while self.n < len(values): self.n *= 2
        self.tree = [0] * (2 * self.n)
        for i, v in enumerate(values): self.tree[self.n + i] = v
        for i in range(self.n - 1, 0, -1):
            self.tree[i] = max(self.tree[2 * i], self.tree[2 * i + 1])
    def update(self, i, value):
        i += self.n;  self.tree[i] = value
        while i > 1:
            i //= 2;  self.tree[i] = max(self.tree[2 * i], self.tree[2 * i + 1])
    def first_at_least(self, need):         # leftmost index with value >= need
        if self.tree[1] < need: return -1
        i = 1
        while i < self.n:
            i = 2 * i if self.tree[2 * i] >= need else 2 * i + 1
        return i - self.n
```

Rolling hash (Rabin-Karp) for repeated substrings

Longest duplicated substring: binary search the length L (pattern 5), and for each L slide a hash over the string in $O(n)$, storing hashes in a set. Roll by removing the leading character and adding the next. Verify a hash hit to rule out collisions, or use two moduli.

```
def has_duplicate(s, length, base=131, mod=(1 << 61) - 1):
    if length == 0: return True
    power = pow(base, length - 1, mod);  h = 0;  seen = set()
    for i, ch in enumerate(s):
        if i >= length: h = (h - ord(s[i - length]) * power) % mod
        h = (h * base + ord(ch)) % mod
        if i >= length - 1:
            if h in seen: return True
            seen.add(h)
    return False
```

Sparse table (range minimum in $O(1)$)

Fence painting asks for the lowest stripe in a range again and again. Precompute `table[j][i]` = minimum of the 2^j items starting at i; a query takes two overlapping blocks.

```
def build(a):
    table = [a[:]];  j = 1
    while (1 << j) <= len(a):
        prev = table[-1]
        table.append([min(prev[i], prev[i + (1 << (j - 1))])
                      for i in range(len(a) - (1 << j) + 1)])
        j += 1
    return table
def range_min(table, l, r):                 # inclusive
    j = (r - l + 1).bit_length() - 1
```

```
    return min(table[j][l], table[j][r - (1 << j) + 1])
```

Huffman coding

Greedy with a heap (pattern 24): repeatedly pop the two lightest nodes and push a parent whose weight is their sum. Reading the final tree, going left is 0 and going right is 1. Canonical codes assign codes by length then symbol so only the lengths need storing.

```
def huffman_lengths(weights):               # weights: {symbol: weight}
    heap = [(w, i, {s: 0}) for i, (s, w) in enumerate(weights.items())]
    heapq.heapify(heap);  tiebreak = len(heap)
    while len(heap) > 1:
        w1, _, left = heapq.heappop(heap);  w2, _, right = heapq.heappop(heap)
        merged = {s: d + 1 for s, d in {**left, **right}.items()}
        heapq.heappush(heap, (w1 + w2, tiebreak, merged));  tiebreak += 1
    return heap[0][2]                       # symbol -> code length
```

Fast doubling and matrix power (huge n recurrences)

When n is 10^{18}, Fibonacci-like recurrences are computed in $O(\log n)$ by squaring. Fast doubling: $F(2k) = F(k)(2F(k+1)-F(k))$ and $F(2k+1) = F(k)^2 + F(k+1)^2$. Apply the modulus after every multiplication.

```
def fib(n, mod=10**9 + 7):                  # returns (F(n), F(n+1))
    if n == 0: return (0, 1)
    a, b = fib(n // 2, mod)
    c = a * ((2 * b - a) % mod) % mod
    d = (a * a + b * b) % mod
    return (d, (c + d) % mod) if n % 2 else (c, d)
```

Group testing, augmenting paths and other named ideas

- **Group testing** (find all bad pairs with few oracle calls): test a group in one call; if it passes, discard the whole group; if it fails, split it. Lower bound is about $k \log(n/k)$ calls for k bad items.
- **Augmenting path** (match people to questions): try each preferred question; if taken, recursively ask the owner to move to another; undo on failure. This is bipartite matching in $O(V \cdot E)$.
- **Fisher-Yates** and **reservoir sampling**: shuffle in place from the end; to pick one random item from a stream, replace the current pick with the i-th item with probability $1/i$.
- **Count-Min sketch** / **Misra-Gries**: approximate top-k over an unbounded stream in fixed memory.
- **Sprague-Grundy**: a sum of independent games is won by the player who makes the xor of the game values non-zero (the follow-up behind “games on several piles”).
- **Morris traversal**: in-order walk of a BST in $O(1)$ extra space by temporarily threading predecessor links.

Coverage by section and source

One table per source, in this order: PracHub, 1Point3Acres, HackTheRounds, Aced. In each table every pattern is listed, sorted by that source's own no column (no gaps first, then one gap, then two, and so on; ties keep the reading order), so the same pattern name appears in every table. A star marks the first pattern of each new no value (0, then 1, then 2, ...), with a thin line between the groups. Click a pattern name to jump to it.

PracHub: questions with a solution (the page has a Solution tab with code)

The pattern boxes here and in the contents are one linked box: tick it in either place and both tick. In Adobe Acrobat it also works with the question boxes in both directions: ticking the pattern box ticks every question of that pattern (unticking clears them), and ticking the last question ticks the pattern box. Preview does not run that automatic part, so there you tick each box yourself. Save the file (File, Save) to keep the ticks.

#	Pattern	Qs	PracHub yes	PracHub no
1	☐ ★ Union-find	7	7	0
2	☐ Frontier search (BFS or DFS), tier 1	6	6	0
4	☐ Backtracking	7	7	0
5	☐ Best-first search with a heap	8	8	0
6	☐ Topological sort (Kahn) and cycle detection	4	4	0
7	☐ Frontier search, tier 2	4	4	0
8	☐ Top-K by frequency	5	5	0
9	☐ Games	5	5	0
10	☐ K-way merge	4	4	0
13	☐ Grid DP	5	5	0
16	☐ Sweep over events in time order	9	9	0
18	☐ Sorted intervals	6	6	0
19	☐ Greedy	5	5	0
29	☐ Nested structures	5	5	0
32	☐ Trie	8	8	0
33	☐ A dict for lookup plus an ordered container for who is next	5	5	0
34	☐ Build or transform a grid with index arithmetic	3	3	0
35	☐ Look at the input once and keep only a small summary	6	6	0
36	☐ Simulate a machine	5	5	0
37	☐ Digits and characters	4	4	0
38	☐ Split into pieces, fix each piece, join (and word wrap)	5	5	0
40	☐ Parent pointers	4	4	0
41	☐ Random pick without repeats	2	2	0
42	☐ Tree recursion, tier 2	3	3	0
43	☐ Build and serialize	3	3	0
3	☐ ★ A heap of what is available	8	7	1
12	☐ Recurrence on a position or a range (memo or table)	8	7	1
15	☐ Dict of per-key state (update, expire, rank)	8	7	1
17	☐ Binary search on a yes/no condition	7	6	1
20	☐ Count pairs by looking up each partner	3	2	1
24	☐ Sets and counters across two collections	4	3	1
25	☐ Sliding window	4	3	1
26	☐ Two pointers moving toward each other (or along two lists)	6	5	1
28	☐ Prefix sums, tier 2	3	2	1
30	☐ Tree recursion, tier 1	8	7	1
31	☐ Tree DP	5	4	1
11	☐ ★ Rolling state	5	3	2
14	☐ Counting or reaching an amount	4	2	2
21	☐ Prefix sums, tier 1	6	4	2
22	☐ Runs and groups	4	2	2
23	☐ bisect on a sorted list	6	4	2
27	☐ Group by a canonical key	5	3	2
39	☐ Walk two strings at once with two indices	6	4	2
	All patterns	**228**	**203**	**25**

1Point3Acres: questions with notes (an approach write-up; this site has no solution code)

#	Pattern	Qs	1Point3Acres yes	1Point3Acres no
13	★ Grid DP	5	5	0
29	Nested structures	5	5	0
7	★ Frontier search, tier 2	4	3	1
20	Count pairs by looking up each partner	3	2	1
22	Runs and groups	4	3	1
31	Tree DP	5	4	1
34	Build or transform a grid with index arithmetic	3	2	1
9	★ Games	5	3	2
15	Dict of per-key state (update, expire, rank)	8	6	2
25	Sliding window	4	2	2
33	A dict for lookup plus an ordered container for who is next	5	3	2
35	Look at the input once and keep only a small summary	6	4	2
38	Split into pieces, fix each piece, join (and word wrap)	5	3	2
41	Random pick without repeats	2	0	2
6	★ Topological sort (Kahn) and cycle detection	4	1	3
12	Recurrence on a position or a range (memo or table)	8	5	3
19	Greedy	5	2	3
23	bisect on a sorted list	6	3	3
28	Prefix sums, tier 2	3	0	3
32	Trie	8	5	3
36	Simulate a machine	5	2	3
40	Parent pointers	4	1	3
42	Tree recursion, tier 2	3	0	3
43	Build and serialize	3	0	3
1	★ Union-find	7	3	4
2	Frontier search (BFS or DFS), tier 1	6	2	4
4	Backtracking	7	3	4
8	Top-K by frequency	5	1	4
10	K-way merge	4	0	4
14	Counting or reaching an amount	4	0	4
17	Binary search on a yes/no condition	7	3	4
21	Prefix sums, tier 1	6	2	4
24	Sets and counters across two collections	4	0	4
37	Digits and characters	4	0	4
3	★ A heap of what is available	8	3	5
11	Rolling state	5	0	5
18	Sorted intervals	6	1	5
26	Two pointers moving toward each other (or along two lists)	6	1	5
27	Group by a canonical key	5	0	5
5	★ Best-first search with a heap	8	2	6
16	Sweep over events in time order	9	3	6
30	Tree recursion, tier 1	8	2	6
39	Walk two strings at once with two indices	6	0	6
	All patterns	**228**	**90**	**138**

HackTheRounds: questions with a solution (the page shows solution code)

#	Pattern	Qs	HackTheRounds yes	HackTheRounds no
20	★ Count pairs by looking up each partner	3	1	2
41	Random pick without repeats	2	0	2
42	Tree recursion, tier 2	3	1	2
14	★ Counting or reaching an amount	4	1	3
25	Sliding window	4	1	3
28	Prefix sums, tier 2	3	0	3
34	Build or transform a grid with index arithmetic	3	0	3
40	Parent pointers	4	1	3
43	Build and serialize	3	0	3
6	★ Topological sort (Kahn) and cycle detection	4	0	4
7	Frontier search, tier 2	4	0	4
10	K-way merge	4	0	4
22	Runs and groups	4	0	4
24	Sets and counters across two collections	4	0	4
37	Digits and characters	4	0	4
8	★ Top-K by frequency	5	0	5
9	Games	5	0	5
11	Rolling state	5	0	5
13	Grid DP	5	0	5
19	Greedy	5	0	5
21	Prefix sums, tier 1	6	1	5
23	bisect on a sorted list	6	1	5
27	Group by a canonical key	5	0	5
29	Nested structures	5	0	5
31	Tree DP	5	0	5
33	A dict for lookup plus an ordered container for who is next	5	0	5
36	Simulate a machine	5	0	5
38	Split into pieces, fix each piece, join (and word wrap)	5	0	5
2	★ Frontier search (BFS or DFS), tier 1	6	0	6
3	A heap of what is available	8	2	6
18	Sorted intervals	6	0	6
26	Two pointers moving toward each other (or along two lists)	6	0	6
35	Look at the input once and keep only a small summary	6	0	6
39	Walk two strings at once with two indices	6	0	6
1	★ Union-find	7	0	7
4	Backtracking	7	0	7
15	Dict of per-key state (update, expire, rank)	8	1	7
17	Binary search on a yes/no condition	7	0	7
5	★ Best-first search with a heap	8	0	8
12	Recurrence on a position or a range (memo or table)	8	0	8
16	Sweep over events in time order	9	1	8
30	Tree recursion, tier 1	8	0	8
32	Trie	8	0	8
	All patterns	**228**	**11**	**217**

Aced: questions with a solution (the Solution tab shows code)

#	Pattern	Qs	Aced yes	Aced no
11	★ Rolling state	5	4	1
14	Counting or reaching an amount	4	3	1
28	Prefix sums, tier 2	3	2	1
42	Tree recursion, tier 2	3	2	1
43	Build and serialize	3	2	1
10	★ K-way merge	4	2	2
41	Random pick without repeats	2	0	2
6	★ Topological sort (Kahn) and cycle detection	4	1	3
20	Count pairs by looking up each partner	3	0	3
24	Sets and counters across two collections	4	1	3
26	Two pointers moving toward each other (or along two lists)	6	3	3
34	Build or transform a grid with index arithmetic	3	0	3
40	Parent pointers	4	1	3
7	★ Frontier search, tier 2	4	0	4
8	Top-K by frequency	5	1	4
19	Greedy	5	1	4
22	Runs and groups	4	0	4
25	Sliding window	4	0	4
27	Group by a canonical key	5	1	4
33	A dict for lookup plus an ordered container for who is next	5	1	4
37	Digits and characters	4	0	4
2	★ Frontier search (BFS or DFS), tier 1	6	1	5
4	Backtracking	7	2	5
9	Games	5	0	5
13	Grid DP	5	0	5
18	Sorted intervals	6	1	5
21	Prefix sums, tier 1	6	1	5
29	Nested structures	5	0	5
31	Tree DP	5	0	5
36	Simulate a machine	5	0	5
38	Split into pieces, fix each piece, join (and word wrap)	5	0	5
39	Walk two strings at once with two indices	6	1	5
12	★ Recurrence on a position or a range (memo or table)	8	2	6
17	Binary search on a yes/no condition	7	1	6
23	bisect on a sorted list	6	0	6
30	Tree recursion, tier 1	8	2	6
35	Look at the input once and keep only a small summary	6	0	6
1	★ Union-find	7	0	7
3	★ A heap of what is available	8	0	8
5	Best-first search with a heap	8	0	8
15	Dict of per-key state (update, expire, rank)	8	0	8
32	Trie	8	0	8
16	★ Sweep over events in time order	9	0	9
	All patterns	**228**	**36**	**192**

After the 150 starred questions: what to do next

Stop adding new problems once the starred questions are learned, and spend the remaining time on how you perform in the interview.

What to practise

1. **Mock interviews on the questions you have not seen.** Take a question from the unstarred ones. Within two minutes say the pattern, the data structure and the cost out loud, then code it in 25 to 30 minutes. This shows whether you can recognise a new question as a known pattern, which is what the real round tests.
2. **Practise talking while you code.** Ask clarifying questions, state a brute force first, then improve it, and say why you chose the structure. Practise with a listener or record yourself.
3. **Test every solution like you will in the interview.** Walk a small example by hand, then the edge cases: empty input, one item, duplicates, ties, and the case where there is no answer.
4. **Drill the follow-ups.** Google usually adds a twist after the first answer: a stream, a memory limit, many queries, an extra rule. Rehearse the follow-ups line of each pattern.
5. **Redo your misses.** Keep a short list of every problem you struggled with or peeked at, and redo each one from a blank file two days later. This beats starting new problems.
6. **Go through your friend's PDF** of reported questions at least 4 days before the interview, that is by Nov 2 for an interview on Nov 6. Those friend reports are no longer in this book, so read them there.
7. **Write the templates from memory** for your 10 weakest patterns the night before.
8. **Prepare the non-coding part.** Have two or three short stories ready for behavioural questions, know why Google, and sleep well before the interview. Sample answers written from your resume (page 45).

What to read in this book, in this order

1. Follow-up families that repeat everywhere (page 6): read it first and again the night before. The usual second questions Google asks.
2. Pattern finder (page 6): until you can pick the pattern from the words of a question without opening the book.
3. Python toolbox and complexity (page 5): skim for the tools you forget (heapq, deque, bisect, Counter) and the costs you must state.
4. The **Traps** and **Follow-ups** lines on every pattern page: a few minutes per pattern, they hold the mistakes and twists you are most likely to be caught on.
5. Priority order (page 4): if a high ranked pattern still is not comfortable, reread its page and do one more problem on it.
6. Appendix (page 38): skim last. Fenwick tree, segment tree and Huffman coding show up as follow-ups now and then, but you do not need to master them.

Leave for the very end: the questions with a red star, because you cannot check your answer against a PracHub solution.

Behavioural questions: sample answers from your resume

Every answer below uses only facts that are on your resume. Read each one aloud until you can say it in your own words, in about a minute and a half. Keep the numbers, because they are what the interviewer will remember.

Tell me about yourself.

I am a computer science student at Grambling State University with a 3.95 GPA, graduating in May 2028. What I enjoy most is building something from start to finish and then measuring whether it really works. I founded ZipMath, a mobile app that teaches mental math shortcuts, where I designed the lessons, the engagement system and a question engine that computes every answer in code. I have had three documentation pull requests merged into Apache Kafka, the Kubernetes Java client and PostgreSQL JDBC. In my GPU Robot Lab project I made a matrix multiply kernel 3.4 times faster by working out why my first version was slow. I also lead weekly data structures and algorithms workshops for more than 90 members as Academic Chair of ColorStack Grambling. I want to keep building at that depth next to strong engineers, which is why this internship at Google appeals to me.

Why Google?

Three parts of my own work lead me to Google. One of my open source contributions was to the Kubernetes Java client, which grew out of Google, and I wrote its release history from v15.0.0 to v25.0.0, including the breaking changes in v20.0.0, so I have read closely how a project like that explains change to its users. My GPU Robot Lab used JAX through MuJoCo MJX to run robot simulations up to 126 times faster than one CPU thread, so I already work with tools Google builds. And I like that Google engineers write code at a scale where a small inefficiency matters, which is the kind of problem I enjoyed chasing in my kernel work. I also teach algorithms to my peers, and I want to learn from people who are much better at it than I am.

Tell me about a hard technical problem you solved.

In my GPU Robot Lab I wrote Triton matrix multiply kernels and checked them against PyTorch on an NVIDIA T4. My first version was slow, and I did not guess at the cause. I traced it to register spilling and tuned the kernel from 1,096 to 3,727 GFLOPS, which is 3.4 times faster. I kept going and found that the compiled code uses no tensor cores, which is the likely reason for the gap that remains. I reported that as an open finding and did not claim the kernel was finished. To make sure my numbers could be trusted I ran 10 repeats of every version, because PyTorch itself varied by up to 44 percent between runs, and I changed one thing at a time to test whether the launch order mattered. The lesson I took is to measure before changing anything, and to say clearly what is still unexplained.

Tell me about a time you took initiative or dealt with an unclear problem.

With ZipMath nobody told me what to build, so I had to decide what mattered. I wanted learners to keep practising, so I designed an engagement system with XP, daily streaks, three hearts per practice round, and a one to three star rating where a round with no mistakes earns three stars. I also made one technical decision I care about: the question engine computes every answer in code, so a practice question can never have a wrong answer key. For an app that teaches math, one wrong answer would cost the learner's trust straight away. I shipped 26 levels across the Basic, Intermediate and Advanced tiers, with animated lessons and quizzes that save progress on the device, and there is a live demo.

Tell me about a time you led others.

As Academic Chair of ColorStack Grambling I lead weekly data structures and algorithms workshops for more than 90 members. The hard part is that people arrive at very different levels, so I teach each topic as a small set of patterns that they can recognise in new problems, instead of as a pile of separate questions. I am also starting as the AWS Student Builder Group Leader at Grambling in November 2026, and I plan to bring the same approach: make it easy to join, easy to learn from, and clear what to do next.

Tell me about a time you worked on a team.

In the AI4ALL Ignite project, our team of three built a model that predicts next day stock decreases from financial news headlines, using FinBERT sentiment features and more than 20 million stock price records. I took ownership of the evaluation. I wrote the script that runs 5-fold stratified cross validation on a sample of 50,000 articles and produces classification reports and confusion matrices for both models. That mattered because the Random Forest reached 54.4 percent accuracy against 50.8 percent for the Logistic Regression baseline on 64,599 held out headlines, which is a modest gap, and a careful evaluation was the only way to say honestly whether it was real. I learned that on a team, owning the part that keeps everyone honest is a useful role.

Tell me about a time you worked with feedback or a review.

My three documentation pull requests to Apache Kafka, the Kubernetes Java client and PostgreSQL JDBC each went through the project's own review process before they were merged, so I had to write for maintainers who know the code far better than I do. For Kafka I rewrote the Streams guide and the Javadoc of two core classes to explain that turning on topology optimizations takes two steps, with a code example and a warning that calling build() without Properties silently skips them. That came from spotting a mistake a reader could easily make. It taught me to write for the person who will misread the page, and to treat review as part of getting the change right.

Tell me about a time you had to learn something quickly.

Triton and MuJoCo MJX were both new to me when I started the GPU Robot Lab. I learned them by building: first the kernels, always checked against PyTorch so I knew at every step whether I was right, then the robot simulation, then a planner and a neural network controller trained with a curriculum until it could swing up and hold a pendulum. Having a reference to check against is how I learn quickly without fooling myself. I used the same habit when I wrote the Kubernetes client release history: I read each release in turn and checked every claim before writing it down.

Tell me about a time you made sure your work was high quality.

For InterviewPair, an app that pairs students for mock technical interviews, I did not stop at a working login. I wrote more than 30 automated tests, including cases for expired tokens, tampered tokens and tokens signed with the wrong key, because those are the cases where a login system fails in real life. I also set up continuous integration on GitHub Actions that builds the Docker image on every push, so a broken change shows up immediately and not days later. The GPU Robot Lab has the same habit, with more than 90 automated tests in continuous integration.

How do you manage your time with so many things going on?

I keep a 3.95 GPA while building ZipMath, running weekly workshops and contributing to open source, and the way I manage it is by working in small pieces with a clear next step, and by letting tests and automated builds remember what works so that I do not have to. That means I can come back to a project after a busy exam week and trust that it still runs. It also means I pick my projects deliberately: each one on my resume taught me something specific, instead of being one more thing to keep alive.

Tell me about a time you explained something complicated to someone.

Results like register spilling are easy to state and hard to feel, so for the GPU Robot Lab I built an interactive explainer site on GitHub Pages that lets a reader see the effect for themselves. I use the same approach in my algorithms workshops: start from one small example, show the pattern behind it, and then let the group apply it to a new problem on their own. When people can apply an idea to something they have not seen, I know they understood it.

What questions could you ask the interviewer?

I would ask what a strong intern on the team accomplishes in the first month, how interns get feedback on their code and how often, and what the hardest technical problem the team solved recently was and how they approached it. I would also ask what they would have wanted to know before they started at Google. These show that I care about doing good work and about learning, and each one is an opening for a conversation and not a yes or no answer.