GOOGLE | 1POINT3ACRES | TOPIC 8 OF 9

Greedy and Math

4 questions to study here, each on its own page, with the full interview answer and code

10 more fit this topic but are covered in another topic: they are listed as SKIP at the end of the contents

3	1
HARD	MEDIUM

How to use

- Questions are Hard first, then Medium, then Easy, newest last-asked date first.
- A question that fits several topics has its pages in only one PDF; the other topics list it as SKIP and name where it is.
- Every page has < Prev | Contents | Next > at the bottom. Tap a title in the Contents to jump to it.
- Select code and press Cmd/Ctrl+C to copy it.

Go to Contents >

Contents

Greedy and Math | 4 questions | tap a title to jump to it

H 1. Maximum-Area Rectangle from Planar Points **2026-09-29**

H 2. Wordle-Style 5-Letter Minimum-Guess Strategy **2025-12-15**

H 3. Find All Bad Pairs Using runTest Oracle **2025-11-27**

M 4. Bounded Knapsack Variant **2026-10-16**

SKIP: these fit this topic but have their pages in another topic (10)

H SKIP Huffman Tree Construction **topic 5: Trees and Tries**

H SKIP Count Perfect Wake Numbers up to N **topic 3: Dynamic Programming**

H SKIP Necklace Cut into Two Halves with Equal D/R Counts **topic 2: Strings, Parsing and Recursion**

M SKIP Connected Crop Arrangement on a Grid **topic 6: Grids and Simulation**

M SKIP Path Minimizing the Maximum Cell (Swim in Rising Water) **topic 4: Graphs and Search**

M SKIP Meeting Rooms / Interval Scheduling Variants **topic 1: Arrays, Hashing and Two Pointers**

M SKIP Cutting Ribbons (LC 1891) **topic 1: Arrays, Hashing and Two Pointers**

M SKIP Target Expressions from Ordered Digits **topic 1: Arrays, Hashing and Two Pointers**

M SKIP Greedy Longest-Match Dictionary Tokenizer **topic 2: Strings, Parsing and Recursion**

M SKIP Gate Entry / Exit Time Resolution with Tie-Break Rules **topic 6: Grids and Simulation**

STUDY HERE: topic 8, Greedy and Math

Also fits: Arrays, Hashing and Two Pointers (listed there as SKIP, no pages). Site tags: math, hashmap, hard

HARD | 1Point3Acres | onsite-coding | 1 of 4

Maximum-Area Rectangle from Planar Points

ASKED 2026-09-29 | 21 DAYS AGO | LAST 30 DAYS

STUDY: TALK NO LEETCODE FREQUENCY PRIORITY 177/228

ABOUT THIS QUESTION

Topic: Greedy and Math

Optimal answer from: written and tested here for this exact question

Related in your app (a neighbouring problem, not this exact question): LeetCode #963 Minimum Area Rectangle II

LeetCode match: LC 963 Minimum Area Free Rectangle (related)

Frequency: Low | Last asked: 2026-09-29 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/max-area-rectangle-from-points

QUESTION

Given 2D points, find the maximum area of a rectangle whose four corners are input points; sides need not be parallel to the axes. Candidates compare it to LC 963 (Minimum Area Rectangle II) but maximizing. Interviewers ask to improve beyond an O(N^3) approach.

Input: a list of 2D coordinate points.
Return the maximum area of any rectangle whose four vertices are all input points.
Rectangle sides do not need to be parallel to the coordinate axes.
Follow-up pressure: start from an O(N^3) idea, then optimize toward a better complexity.

APPROACH

Practice the diagonal-grouping idea: two point pairs form a rectangle when they share the same midpoint and the same length, so bucket all pairs by (midpoint, squared length) and compare pairs within each bucket.
Write the area computation from two adjacent side vectors and verify it on both an axis-aligned and a 45-degree rotated rectangle.
Prepare the complexity argument: O(N^2) pairs to bucket, with the worst case driven by bucket sizes.

HOW TO THINK ABOUT IT

Four points form a rectangle exactly when two of its diagonals share the same midpoint and have the same length. So you never need to pick four points: bucket every pair of points by (sum of the two points, squared distance). Two different pairs in one bucket are the two diagonals of a rectangle, and the area is the cross product of two adjacent sides.

HOW TO THINK ABOUT IT

Basis: LeetCode 963, Minimum Area Free Rectangle, with the objective flipped to maximum area.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def brute(points):
    best = 0
    for p, q, r in permutations(points, 3):
        if dot(q - p, r - p) != 0:
            continue  # no right angle at p
        s = q + r - p  # the fourth corner
        if s in points:
            area = abs(cross(q - p, r - p))
            best = max(best, area)
    return best
```

TIME	SPACE
O(n^3)	O(n)

WHY IT IS SLOW

Three nested loops over the points, O(n^3). Bucketing pairs by their diagonal midpoint and length finds rectangles from O(n^2) pairs.

OPTIMAL SOLUTION · Bucket pairs by diagonal midpoint and length

TIME	SPACE
O(n^2) pairs (worst case O(n^3) if every pair shares a bucket)	O(n^2)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 Can you do better than O(n^2) pairs?

Not in general: the answer can depend on all pairs, and even deciding whether any rectangle exists is believed to need about n^2 work. The practical gain is in the buckets: most buckets hold one pair, so the work after bucketing is small. State the worst case honestly (points on a circle put many pairs in few buckets).

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
from collections import defaultdict
from itertools import combinations


def cross(o, p, q):
    # twice the signed area of triangle o, p, q
    ax, ay = p[0] - o[0], p[1] - o[1]
```

```
    bx, by = q[0] - o[0], q[1] - o[1]
    return ax * by - ay * bx


# FOLLOW-UP 1: Can you do better than O(n^2) pairs?
def all_rectangle_areas(points):
    # duplicates cannot help, so drop them
    pts = list(set(map(tuple, points)))
    buckets = defaultdict(list)
    # every pair is a possible diagonal
    for p, q in combinations(pts, 2):
        # same centre (sum of corners) and same length
        key = (p[0] + q[0], p[1] + q[1],
               (p[0] - q[0]) ** 2 + (p[1] - q[1]) ** 2)
        buckets[key].append((p, q))
    # two diagonals in one bucket form a rectangle
    for pairs in buckets.values():
        for (p, q), (r, s) in combinations(pairs, 2):
            # area = |cross| of two sides leaving p
            yield abs(cross(p, r, s))


# BASE QUESTION (optimal): Bucket pairs by diagonal
# midpoint and length
def max_rect_area(points):
    # best area, or 0 when no rectangle exists
    return max(all_rectangle_areas(points), default=0)


# RELATED VARIANT: smallest area (LeetCode 963)
def min_rect_area(points):
    # LeetCode 963 asks for the smallest instead
    return min(all_rectangle_areas(points), default=0)


# BRUTE FORCE (say it first)
def max_rect_area_cubic(points):
    # first idea to say aloud: try every corner triple
```

```
    pts = set(map(tuple, points))
    best = 0
    for p in pts:
        for q in pts:
            for r in pts:
                if len({p, q, r}) < 3:
                    continue
                # right angle at p between p->q and p->r
                dot = ((q[0] - p[0]) * (r[0] - p[0])
                       + (q[1] - p[1]) * (r[1] - p[1]))
                # the fourth corner must also exist
                s = (q[0] + r[0] - p[0],
                     q[1] + r[1] - p[1])
                if dot == 0 and s in pts:
                    area = abs(cross(p, q, r))
                    best = max(best, area)
    return best


if __name__ == "__main__":
    # a tilted square of area 2
    sq = [(0, 0), (1, 1), (0, 2), (-1, 1)]
    assert max_rect_area(sq) == 2
    assert max_rect_area_cubic(sq) == 2
    # axis-aligned 2 x 1 rectangle plus a stray point
    rect = [(0, 0), (2, 0), (0, 1), (2, 1), (5, 5)]
    assert max_rect_area(rect) == 2
    # collinear points make no rectangle
    assert max_rect_area([(0, 0), (1, 1), (2, 2)]) == 0
    # the minimum version picks the small one
    two = rect + [(0, 3), (2, 3)]
    assert min_rect_area(two) == 2
    assert max_rect_area(two) == 6
```

NOTES FROM CANDIDATES

Candidates describe it as LC 963 (Minimum Area Rectangle II) with the objective flipped to maximum area.
Asked as a standalone L4 onsite coding round; in another L4 onsite it was the second problem after a CPU scheduling question, and the interviewer explicitly expected a better-than-cubic solution.
Clarify up front: whether duplicate points can appear, what to return when no rectangle exists (commonly 0), and whether integer or floating-point coordinates are used (area precision matters for rotated rectangles).

PREPARATION

Practice the diagonal-grouping idea: two point pairs form a rectangle when they share the same midpoint and the same length, so bucket all pairs by (midpoint, squared length) and compare pairs within each bucket.
Write the area computation from two adjacent side vectors and verify it on both an axis-aligned and a 45-degree rotated rectangle.
Prepare the complexity argument: O(N^2) pairs to bucket, with the worst case driven by bucket sizes.

STUDY HERE: topic 8, Greedy and Math

Site tags: algorithm-design, greedy, hard

HARD | 1Point3Acres | phone-screen, onsite-coding | 2 of 4

Wordle-Style 5-Letter Minimum-Guess Strategy

ASKED 2025-12-15 | 309 DAYS AGO | OLDER

STUDY: TALK | ALL TIME 42.2 | 30 DAYS 0 | PRIORITY 85/228

ABOUT THIS QUESTION

Topic: Greedy and Math

Optimal answer from: written and tested here for this exact question

Related in your app (a neighbouring problem, not this exact question): LeetCode #843 Guess the Word

LeetCode match: LC 843 Guess the Word (related)

Frequency: Single report | Last asked: 2025-12-15 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/wordle-min-guess-strategy

QUESTION

Phone screen coding: given a black-box oracle that tells you whether a specific letter is at a specific position in a hidden 5-letter word, design an algorithm to identify the word in as few oracle calls as possible. Follow-up assumes a dictionary of candidate words is available.

Hidden word: a 5-letter string from a..z.
Oracle: query(pos: int, letter: char) -> bool returns whether position pos (0-indexed) holds letter.
Return the hidden word.
Primary metric: minimize the number of query calls.

Brute-force baseline: 5 × 26 = 130 calls.
Follow-up

A dictionary of candidate 5-letter words is available; design the strategy to use bigram/letter-frequency information to cut the expected number of calls.

EXAMPLES

```
HiddenWord = "GRAPE"
query(0, 'A') -> false
query(0, 'G') -> true   // first letter found
...
```

APPROACH

Drill information-theoretic puzzles (egg drop, twenty questions) to internalize the entropy-split heuristic.
Practice maintaining a candidate set under positive/negative oracle feedback (filter on each query result).
Be ready with a one-liner on Wordle's known optimal first guess (SOARE / CRANE) for intuition; do not hard-code it.

HOW TO THINK ABOUT IT

Each call returns one bit, so you can never beat ceil(log2(number of candidates)) calls. Without a dictionary the best you can do is order the questions by likelihood and stop at the first yes for each position (and skip the 26th letter, because it is implied). With a dictionary, keep the list of words still possible and ask the (position, letter) question whose yes/no answer splits that list most evenly.

Basis: Written for this question; related to LeetCode 843, Guess the Word, but the oracle answers one position and one letter per call.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
word = ""
for i in range(5):
    for c in "abcdefghijklmnopqrstuvwxyz":
        if query(i, c):
            word += c
            break
```

TIME	SPACE
130 calls	O(1)

WHY IT IS SLOW

Always 130 calls. Asking likely letters first, skipping the implied last letter, or splitting a dictionary of candidates in half needs far fewer.

OPTIMAL SOLUTION · Frequency order, or entropy split when a dictionary is known

TIME	SPACE
at most 5 x 25 calls with... at most 5 x 25 calls without a dictionary; about log2(words) with one	O(words)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 What is the lower bound and when is it reached?

Every call returns at most one bit, so at least ceil(log2(candidates)) calls are needed. The entropy split reaches it only when a question can cut the candidates exactly in half; real dictionaries usually need a few more. A greedy scoring rule (closest to half) is enough to state; the interviewer mostly wants the entropy idea.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
import string

# letters from most to least common in English words
BY_FREQUENCY = "etaoinshrdlcumwfgypbvkjxqz"


# BRUTE FORCE: 130 calls
def brute_force(query):
    # 130 calls at most: every position, every letter
    return "".join(
        next(c for c in string.ascii_lowercase
             if query(i, c))
        for i in range(5))


# BASE QUESTION (optimal): Frequency order, or entropy
# split when a dictionary is known
def by_frequency(query):
    # no dictionary: common letters first, stop at the
    # first
    # yes, and skip the 26th letter (it must be the one
    # left)
    calls, word = 0, []
    for i in range(5):
        for k, c in enumerate(BY_FREQUENCY):
            if k == 25:
                word.append(c)
                break
            calls += 1
            if query(i, c):
                word.append(c)
                break
    return "".join(word), calls


# FOLLOW-UP 1: What is the lower bound and when is it
# reached?
def split_score(words, i, c):
    # how evenly does asking (position i, letter c) split?
    yes = sum(w[i] == c for w in words)
    return min(yes, len(words) - yes)


# BASE QUESTION (optimal): Frequency order, or entropy
# split when a dictionary is known
```

```
def with_dictionary(query, dictionary):
    # keep the words that still fit every answer so far
    words, calls = list(dictionary), 0
    while len(words) > 1:
        options = {(i, w[i]) for w in words
                   for i in range(5)}
        # ask the question that cuts the list closest to
        # half
        i, c = max(options, key=lambda o: split_score(
            words, o[0], o[1]))
        calls += 1
        hit = query(i, c)
        words = [w for w in words if (w[i] == c) == hit]
    return words[0], calls


if __name__ == "__main__":
    def oracle(word):
        return lambda i, c: word[i] == c
```

```
    secret = "grape"
    assert brute_force(oracle(secret)) == secret
    word, calls = by_frequency(oracle(secret))
    assert word == secret and calls <= 125
    words = ["grape", "grade", "trade", "crane", "plane",
             "slate", "stare", "share"]
    for target in words:
        ask = oracle(target)
        found, used = with_dictionary(ask, words)
        assert found == target
        # about log2(8) = 3 questions suffice; allow a few
        # more
        assert used <= 6
    # lower bound: one bit per call, so ceil(log2(8)) = 3
    # calls
    assert min(with_dictionary(oracle(w), words)[1]
               for w in words) >= 3
```

NOTES FROM CANDIDATES

Without the dictionary, query positions in order; per position, query letters in descending frequency (e, t, a, o, i, n, ...) and stop as soon as one returns true.
With the dictionary, maintain a candidate set; at each step pick the (position, letter) query whose true/false outcome splits the candidate set most evenly (entropy maximization).
Mention information theory: each query yields <= 1 bit, so the absolute lower bound on calls is ceil(log2(|candidates|)).
Don't waste minutes on a closed-form optimum — the interviewer mostly wants you to articulate the entropy heuristic and pick one reasonable scoring function.

PREPARATION

Drill information-theoretic puzzles (egg drop, twenty questions) to internalize the entropy-split heuristic.
Practice maintaining a candidate set under positive/negative oracle feedback (filter on each query result).
Be ready with a one-liner on Wordle's known optimal first guess (SOARE / CRANE) for intuition; do not hard-code it.

STUDY HERE: topic 8, Greedy and Math

Site tags: algorithm-design, math-reasoning, hard

HARD | 1Point3Acres | onsite-coding | 3 of 4

Find All Bad Pairs Using runTest Oracle

ASKED 2025-11-27 | 327 DAYS AGO | OLDER

STUDY: TALK | NO LEETCODE FREQUENCY | PRIORITY 192/228

ABOUT THIS QUESTION

Topic: Greedy and Math

Optimal answer from: written and tested here for this exact question

Frequency: Single report | Last asked: 2025-11-27 | Roles: SWE, RS

Original: https://www.1point3acres.com/interview/problems/company/google/find-bad-pairs-with-oracle

QUESTION

PhD intern coding round 2: given a black-box oracle **`runTest(S)`** that returns **`true`** iff a subset **`S`** contains no bad pair, identify all bad pairs across N test items with as few oracle calls as possible.

Items t_1, ..., t_N are given.
Oracle: runTest(S: Set[Item]) -> bool returns true iff S contains zero bad pairs; otherwise false.
Output: every bad pair (t_i, t_j).
Minimize total oracle calls.

EXAMPLES

```
N=4, bad pairs = {(1,2)}.
runTest({1,2,3,4}) -> false -> at least one bad pair
    somewhere.
runTest({1,3,4})   -> true  -> 2 is in every bad pair
    containing items from this set.
... binary search to isolate.
```

APPROACH

Read the Wikipedia article on group testing and Dorfman's two-stage scheme.
Drill recursive binary partitioning on a candidate set with positive/negative oracle feedback.
Have a back-of-envelope analysis ready: "Total calls O(K log N) for K bad pairs and N items."

HOW TO THINK ABOUT IT

Testing every pair costs N(N-1)/2 calls. A single call on a whole set answers 'is there any bad pair inside', so you can discard good sets fast and binary-search for the culprit inside a bad one. Add the tests one at a time to groups known to be clean: if a group plus the new test fails, binary search the group for a partner; otherwise the new test joins it.

Basis: Adaptive group testing; the same method as the 'bad pairs with runTest' question in the PracHub set.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
bad = []
for i in range(n):
    for j in range(i + 1, n):
        if not runTest({i, j}):
            bad.append((i, j))
```

TIME	SPACE
O(N^2) calls	O(1)

WHY IT IS SLOW

Always N(N - 1) / 2 calls. Testing whole groups and binary searching a failing group finds each bad pair in about log N calls.

OPTIMAL SOLUTION · Group the tests and binary search for partners

TIME	SPACE
O(K log N + N * G) about calls (K bad pairs, G groups)	O(N)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 How many calls are needed in the best case and the worst case?

With no bad pairs it takes N - 1 calls, one per test after the first. With K bad pairs each one costs about log N calls to locate, so roughly K log N plus the cost of building clean groups. The brute force is always N(N-1)/2, which is far larger as soon as bad pairs are rare.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```python
import itertools


# BRUTE FORCE: ask about every pair
def bad_pairs_brute(items, run_test):
    # every pair costs a call: N * (N - 1) / 2 calls
    pairs = itertools.combinations(items, 2)
    return {(a, b) for a, b in pairs
            if not run_test({a, b})}


def first_partner(group, t, run_test):
    # binary search: the shortest prefix of the group that
    # fails together with t; its last member is a partner
    lo, hi = 1, len(group)
    while lo < hi:
        mid = (lo + hi) // 2
        if run_test(set(group[:mid]) | {t}):
```

```
            lo = mid + 1
        else:
            hi = mid
    return group[lo - 1]


def partners(group, t, run_test):
    # all members of the group that are bad with t
    rest, found = list(group), []
    while rest and not run_test(set(rest) | {t}):
        x = first_partner(rest, t, run_test)
        found.append(x)
        # take it out and test the others again
        rest.remove(x)
    return found


# BASE QUESTION (optimal): Group the tests and binary
# search for partners
def find_bad_pairs(items, run_test):
    # groups of items that pass together (no bad pair
    # inside)
    groups, pairs = [], set()
    for t in items:
        placed = False
        for group in groups:
            found = partners(group, t, run_test)
```

```
            for x in found:
                pairs.add((min(x, t), max(x, t)))
            # t joins the first group with no partner for
            # it
            if not found and not placed:
                group.append(t)
                placed = True
        if not placed:
            groups.append([t])
    return pairs


if __name__ == "__main__":
    def make_oracle(bad):
        calls = [0]

        def run_test(subset):
            calls[0] += 1
            ordered = sorted(subset)
            pairs = itertools.combinations(ordered, 2)
            return not any(p in bad for p in pairs)
        return run_test, calls

    bad = {(3, 17), (5, 29)}
    run, calls = make_oracle(bad)
    assert find_bad_pairs(list(range(30)), run) == bad
    # far fewer than the 435 calls of the brute force
    assert calls[0] < 435 // 2
    # no bad pairs: about one call per item
    run, calls = make_oracle(set())
    assert find_bad_pairs(list(range(30)), run) == set()
    assert calls[0] <= 30
    run, _ = make_oracle({(1, 2)})
    assert bad_pairs_brute([1, 2, 3, 4], run) == {(1, 2)}
```

NOTES FROM CANDIDATES

Standard "group testing" / Dorfman testing problem. Walk through the analogy with COVID pooled tests for intuition.
Algorithm sketch:

NOTES FROM CANDIDATES

Test the full set. If true, no bad pairs.
Otherwise, binary-search to find one bad pair by partitioning the set and testing halves.
Remove the items that participate in the located bad pair; repeat.

Tighter bound: with up to K bad pairs, total queries are O(K log N) using adaptive group testing.
A naive $O(N^2)$ pairwise scheme works but the interviewer expects you to beat it with adaptive splitting.
Verbalize correctness: each binary-search step rules out half the candidates; you can always reduce the search space until exactly one bad pair is isolated.

PREPARATION

Read the Wikipedia article on group testing and Dorfman's two-stage scheme.
Drill recursive binary partitioning on a candidate set with positive/negative oracle feedback.
Have a back-of-envelope analysis ready: "Total calls O(K log N) for K bad pairs and N items."

STUDY HERE: topic 8, Greedy and Math

Also fits: Dynamic Programming, Heaps, Stacks, Sorting and Binary Search (listed there as SKIP, no pages). Site tags: greedy, dp, quickselect, sorting, medium

MEDIUM | 1Point3Acres | tech-screen | 4 of 4

Bounded Knapsack Variant

ASKED 2026-10-16 | 4 DAYS AGO | LAST 30 DAYS

STUDY: TALK | NO LEETCODE FREQUENCY | PRIORITY 180/228

ABOUT THIS QUESTION

Topic: Greedy and Math

Optimal answer from: written and tested here for this exact question

LeetCode match: LC 1710 Maximum Units on a Truck

Frequency: Single report | Last asked: 2026-10-16 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/unit-weight-bounded-knapsack

QUESTION

Each item type has a value and an available count; every item weighs 1. Given capacity W, maximize total value. DP or DFS + memo works at O(n·W), greedy by value after sorting is O(n log n), and the best answer uses quickselect for O(n).

Input: for each item type, its value and the number of copies available. Every item has weight 1.
Input: a maximum total weight W.
Output: the maximum total value achievable without exceeding W.

APPROACH

Write the sort-based greedy and argue its optimality in two sentences (unit weights mean every slot is interchangeable). Implement quickselect that partitions item types by value and tracks the total count on each side to locate where W runs out. Compare the O(n × W), O(n log n), and O(n) versions out loud with their complexities.

INTERVIEW-DAY PLAN

CLARIFY FIRST

Ask whether every item really weighs 1, whether values can be zero or negative (assume non-negative), whether W can exceed the total number of items, and how large the counts and W can get.

BRUTE FORCE TO SAY OUT LOUD

A knapsack table over the capacity, one pass per copy. O(total copies * W), correct for any weights.

OPTIMAL IDEA

With unit weights every slot is interchangeable, so sort the item types by value, highest first, and take min(copies, slots left) of each until the bag is full. O(n log n). To reach O(n), run quickselect on the value: split the types into higher, equal and lower, recurse into the side that contains the W-th slot, and keep a running total of what was taken.

INTERVIEW-DAY PLAN

DRY-RUN THESE CASES

W = 0, W larger than all copies together, one type, equal values, a type with a huge count, all values equal.

LIKELY FOLLOW-UPS

Different weights (bounded knapsack with binary splitting), a huge W (greedy still works), or negative values (never take them).

PITFALLS

Reaching for DP because the word knapsack appears. Forgetting that a type may only be partly taken. Letting the remaining capacity go negative.

HOW TO THINK ABOUT IT

The word knapsack points at DP, but the statement says every item weighs 1. With equal weights every slot in the bag is interchangeable, so an item that is worth more is never worse than one that is worth less: swapping a cheaper item for a dearer one always helps. Take item types by value, highest first, until W slots are used. You do not even need the full sort: quickselect finds where the W-th slot falls in O(n).

Basis: LeetCode 1710, Maximum Units on a Truck: every item weighs 1, so the knapsack collapses to 'take the most valuable items first'.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def knapsack(items, W):
    best = [0] * (W + 1)
    for value, count in items:
        for _ in range(count):
            for w in range(W, 0, -1):
                take = best[w - 1] + value
                best[w] = max(best[w], take)
    return best[W]
```

TIME	SPACE
O(total copies * W)	O(W)

WHY IT IS SLOW

Every copy is its own pass over the table, O(total copies * W). With unit weights the most valuable items are always the right ones, so sorting by value and taking greedily is O(n log n).

OPTIMAL SOLUTION · Sort by value and take greedily

TIME	SPACE
O(n log n)	O(1) extra

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 Why is greedy correct here but not for a general knapsack?

With unit weights each chosen item uses exactly one slot, so the only thing that matters is value per slot, which is just the value. If an optimal bag held a cheaper item while a dearer one was available, swapping them keeps the weight and raises the value, so the optimal bag can always be rearranged into the greedy one. With different weights a cheap item can leave room for several others, so value per weight and the leftover space interact and greedy fails.

FOLLOW-UP 2 What if the weights are not all 1?

Then it is a real bounded knapsack and greedy no longer works; use DP over the capacity. Splitting each count into 1, 2, 4, ... copies keeps the number of 0/1 items at about log(count) per type, so the cost is O(n * log(max count) * W) instead of O(total copies * W).

FOLLOW-UP 3 W is up to 10^9. Which version still works?

Only the greedy ones: sorting is O(n log n) and quickselect O(n), and neither depends on W, while any table over the capacity needs O(W) memory. The loop simply stops when the remaining capacity reaches zero.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```python
import random


# BASE QUESTION (optimal): Sort by value and take greedily
def max_value_sorted(items, W):
    """Unit weights: every slot is worth the same, so the
    best value always goes in first."""
    total = 0
    for value, count in sorted(items, reverse=True):
        # take as many copies as still fit in the bag
        take = min(count, W)
        total += value * take
        W -= take
        if W == 0:
            break
    return total


# SAME answer in O(n): quickselect
def max_value_select(items, W):
    """Same answer in O(n) expected: quickselect finds
    where the W-th slot falls without a full sort."""
    total, pool = 0, list(items)
    while pool and W > 0:
        pivot = random.choice(pool)[0]
        high = [it for it in pool if it[0] > pivot]
        mid = [it for it in pool if it[0] == pivot]
        low = [it for it in pool if it[0] < pivot]
        high_count = sum(c for _, c in high)
        if high_count >= W:
            # the bag fills up inside the dearer types
            pool = high
            continue
        # every dearer copy fits, then fill from the pivot
        total += sum(v * c for v, c in high)
        W -= high_count
        take = min(W, sum(c for _, c in mid))
        total += pivot * take
        W -= take
        pool = low
    return total


# FOLLOW-UP 2: What if the weights are not all 1?
def knapsack_bounded(items, W):
    """Weights differ: greedy fails, so use DP. Splitting
    each count into 1, 2, 4, ... copies keeps it small."""
    best = [0] * (W + 1)
    for value, weight, count in items:
        chunk = 1
        while count > 0:
            take = min(chunk, count)
            count -= take
            cost, gain = take * weight, take * value
            for w in range(W, cost - 1, -1):
                best[w] = max(best[w],
                              best[w - cost] + gain)
            chunk *= 2
    return best[W]


# CHECKER (brute force)
def knapsack_dp(items, W):
    # the oracle: every copy is its own 0/1 item
    best = [0] * (W + 1)
    for value, count in items:
        for _ in range(count):
            for w in range(W, 0, -1):
                take = best[w - 1] + value
                best[w] = max(best[w], take)
    return best[W]


if __name__ == "__main__":
    demo = [(3, 2), (5, 1), (4, 4)]
    assert max_value_sorted(demo, 4) == 17
    assert max_value_select(demo, 4) == 17
    assert max_value_sorted([], 5) == 0
    assert max_value_sorted([(7, 3)], 0) == 0
    assert max_value_sorted([(7, 3)], 10) == 21
    mixed = [(10, 5, 1), (6, 3, 2)]
    assert knapsack_bounded(mixed, 6) == 12
    # greedy is only valid when every weight is 1
    unit = [(3, 1, 2), (5, 1, 1), (4, 1, 4)]
    assert knapsack_bounded(unit, 4) == 17
    for _ in range(300):
        n = random.randint(0, 6)
        rnd = random.randint
        items = [(rnd(1, 9), rnd(1, 4)) for _ in range(n)]
        W = random.randint(0, 12)
        want = knapsack_dp(items, W)
        assert max_value_sorted(items, W) == want
        assert max_value_select(items, W) == want
```

NOTES FROM CANDIDATES

Pitfalls: Recognizing "knapsack" and reaching for DP is the trap; with unit weights the greedy argument is simpler and faster.

PREPARATION

Write the sort-based greedy and argue its optimality in two sentences (unit weights mean every slot is interchangeable). Implement quickselect that partitions item types by value and tracks the total count on each side to locate where W runs out. Compare the O(n × W), O(n log n), and O(n) versions out loud with their complexities.