GOOGLE | 1POINT3ACRES | TOPIC 8 OF 10 | LESS COMMON TOPIC

Grids and Simulation

3 questions to study here, each on its own page, with the full interview answer and code

13 more fit this topic but are covered in another topic: they are listed as SKIP at the end of the contents

1	2
HARD	MEDIUM

How to use

- Questions are Hard first, then Medium, then Easy, newest last-asked date first.
- PDFs run from the one with the most questions to the fewest. A question has pages in only one PDF; the others list it as SKIP
- Every page has < Prev | Contents | Next > at the bottom. Tap a title in the Contents to jump to it.
- Select code and press Cmd/Ctrl+C to copy it.

Go to Contents >

Contents

Grids and Simulation | 3 questions | tap a title to jump to it

H		1. Babylon Tile-Stack Merge Game	**2025-12-29**
M		2. Tic-Tac-Toe System with K Players and 3-in-a-Row Win	**2026-05-18**
M		3. Gate Entry / Exit Time Resolution with Tie-Break Rules	**2026-05-04**

SKIP: these fit this topic but have their pages in another topic (13)

H	SKIP	Matrix Flower Placement with House Adjacency	**topic 7: Strings, Parsing and Recursion**
M	SKIP	Mahjong Winning Hand Detection (14 Tiles)	**topic 7: Strings, Parsing and Recursion**
M	SKIP	Connected Crop Arrangement on a Grid	**topic 1: Arrays, Hashing and Two Pointers**
M	SKIP	Number of Islands	**topic 3: Graphs and Search**
M	SKIP	Grid Path Count with Diagonal Moves and Checkpoints	**topic 2: Dynamic Programming**
M	SKIP	Expression Tree Single-Leaf Mutation	**topic 5: Trees and Tries**
M	SKIP	Rotten Oranges / Multi-Source BFS (taxis)	**topic 3: Graphs and Search**
M	SKIP	Longest Same-Color Grid Sequence	**topic 1: Arrays, Hashing and Two Pointers**
M	SKIP	Piano Hand Movement Grouping	**topic 1: Arrays, Hashing and Two Pointers**
M	SKIP	Google Doc Line Wrap (Word Wrap with Whitespace and Newlines)	**topic 7: Strings, Parsing and Recursion**
M	SKIP	Rotate Any m×n 2D Matrix	**topic 1: Arrays, Hashing and Two Pointers**
M	SKIP	Count Same-Color Squares in Matrix	**topic 2: Dynamic Programming**
E	SKIP	Robot Status Message Deduplication Window	**topic 1: Arrays, Hashing and Two Pointers**

STUDY HERE: topic 8, Grids and Simulation

Also fits: Design, Streaming and Scale (listed there as SKIP, no pages). Site tags: simulation, data-structure, hard

HARD | 1Point3Acres | onsite-coding | 1 of 3

Babylon Tile-Stack Merge Game

ASKED 2025-12-29 | 295 DAYS AGO | OLDER

STUDY: TALK **NO LEETCODE FREQUENCY** **PRIORITY 191/228**

ABOUT THIS QUESTION

Topic: Grids and Simulation

Optimal answer from: written and tested here for this exact question

Frequency: Single report | Last asked: 2025-12-29 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/babylon-stack-merge-game

QUESTION

Onsite NG round 1 (Munich office): implement a two-player tile-stack merge game ("Babylon") given the rules. Start state is 12 height-1 stacks of 4 colors (Y / W / G / B), 3 of each; a move places one stack on top of another under a constraint; the player who cannot move loses.

12 tiles initially, partitioned into 12 height-1 stacks.
4 colors: Y / W / G / B; exactly 3 tiles per color.
Two players alternate moves.
A move = pick two stacks and place one entirely on top of the other (merge).

Legal iff at least one of: (a) the two stacks have the same height, OR (b) their top-most tiles share the same color.
After merge: new height = sum of the two heights, new top color = color of the stack placed on top.

The player who has no legal move on their turn loses.
Implement the full game so it can be played out (legal-move detection, state transition, terminal check).

EXAMPLES

Start: 12 stacks of height 1, with tops Y Y Y W W W G G G B B B. A legal opening move: stack the first Y on the second Y (same color top), producing one height-2 stack of color Y plus 10 singletons.

APPROACH

Drill clean game-state implementations: e.g. Connect 4, Othello move generators, in 30 minutes.
Practice writing a generic minimax framework so you can plug in this game's legalMoves(state) and terminal(state) quickly.
Watch the merge invariant carefully: "same height OR same top color" — fail mode is forgetting the alternative branch.

HOW TO THINK ABOUT IT

Model a state as the list of stacks, each as (height, top colour). A legal move puts stack i on stack j when the heights are equal or the top colours are equal; the new stack has the summed height and the colour of the stack that was placed on top. The player with no legal move loses, so a state is winning when some move leads to a losing state. The order of stacks does not matter, so sort the list to share states in a cache.

Basis: Written for this question: a two-player impartial game (the player who cannot move loses), solved by memoised search over the multiset of stacks.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def wins(position):
    for i, j in legal_moves(position):
        nxt = merge(position, i, j)
        if not wins(nxt):
            return True  # opponent loses
    return False
```

TIME	SPACE
exponential	O(depth)

WHY IT IS SLOW

Explores the same positions many times, exponential. Sorting the stacks and caching each position's result visits every distinct position once.

OPTIMAL SOLUTION · Negamax with memoisation on the sorted stack tuple

TIME	SPACE
O(states * moves^2)	O(states)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 How big is the real state space?

With 12 starting stacks the first position has at most 66 unordered pairs and the number of reachable multisets of (height, colour) stacks shrinks quickly because every merge removes a stack, so memoised search is feasible. In the interview build the clean state class, legal moves and play(move) first and add the solver only if there is time.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
import itertools
```

```
# BRUTE FORCE (say it first): Plain recursion over every
# move | time exponential
def babylon_brute(stacks):
    stacks = list(stacks)
    for i, j in itertools.permutations(
        range(len(stacks)), 2
    ):
        (h1, c1), (h2, c2) = stacks[i], stacks[j]
        if h1 == h2 or c1 == c2:
            rest = [
                s
                for k, s in enumerate(stacks)
                if k not in (i, j)
            ] + [(h1 + h2, c1)]
            if not babylon_brute(rest):
                return True
    return False


# HELPERS for the base question
class Babylon:
    # a stack is (height, colour of its top tile)
    def __init__(self, stacks=None):
        if stacks is None:
            # 12 tiles: Y, W, G, B three times each
            stacks = [(1, c) for c in "YWGB" * 3]
        self.stacks = list(stacks)

    def legal_moves(self):
        # put stack i on top of stack j
        moves = []
        for i, j in itertools.permutations(
                range(len(self.stacks)), 2):
            h1, c1 = self.stacks[i]
            h2, c2 = self.stacks[j]
            # same height, or the same colour on top
            if h1 == h2 or c1 == c2:
                moves.append((i, j))
        return moves

    def after(self, i, j):
        # the stacks after the move (the original is kept)
        (h1, c1), (h2, _) = self.stacks[i], self.stacks[j]
        rest = [s for k, s in enumerate(self.stacks)
                if k not in (i, j)]
        # heights add; the top colour is the moved stack's
        return rest + [(h1 + h2, c1)]

    def play(self, i, j):
        if (i, j) not in self.legal_moves():
            raise ValueError("illegal move")
        self.stacks = self.after(i, j)

    def is_over(self):
        # the player with no move loses
        return not self.legal_moves()


# BASE QUESTION (optimal): Negamax with memoisation on the
# sorted stack tuple
def wins(stacks):
    # does the player to move win with best play?
    cache = {}

    def solve(state):
        if state in cache:
            return cache[state]
        game = Babylon(state)
        result = False
        for i, j in game.legal_moves():
            nxt = tuple(sorted(game.after(i, j)))
            # a move to a position the opponent loses wins
            if not solve(nxt):
                result = True
                break
        cache[state] = result
        return result

    # sorted stacks: equal positions share one cache entry
    return solve(tuple(sorted(stacks)))


# PICK the winning move itself
def best_move(game):
    # pick a move that leaves the opponent lost, if any
    for i, j in game.legal_moves():
        if not wins(game.after(i, j)):
            return i, j
    moves = game.legal_moves()
    return moves[0] if moves else None


if __name__ == "__main__":
    # one stack: no moves, so the player to move loses
    assert Babylon([(1, "Y")]).is_over()
    assert not wins([(1, "Y")])
    # two equal stacks: one move ends the game, so we win
    assert wins([(1, "Y"), (1, "W")])
    g = Babylon([(1, "Y"), (1, "Y"), (1, "W")])
    assert (0, 1) in g.legal_moves()
    # play a whole game with the solver; the loser is
    # stuck
    while not g.is_over():
        g.play(*best_move(g))
    assert len(g.stacks) >= 1
    try:
        Babylon([(1, "Y"), (2, "W")]).play(0, 1)
        raise AssertionError("move should be illegal")
    except ValueError:
        pass
    # the real start has 12 stacks and up to 132 ordered
    # moves
    assert len(Babylon().legal_moves()) > 0
```

NOTES FROM CANDIDATES

This is closer to an implementation/simulation interview than a clever-algorithm problem — the interviewer wants a clean game state class, legal-move enumeration, and a play(move) API. Represent state as a List[Stack] where Stack = (height, topColor, fullSequence?). Whether you keep the full color sequence depends on whether the interviewer asks about "undo" or "re-color from top".

Add a minimax / negamax for "who wins from this state with optimal play" as the natural follow-up — branching factor is small (at most C(12, 2) = 66 initially, drops fast).

Reported as "too hard to finish coding in 45 minutes"; the interviewer gives hints throughout. Pacing matters.

PREPARATION

Drill clean game-state implementations: e.g. Connect 4, Othello move generators, in 30 minutes.
Practice writing a generic minimax framework so you can plug in this game's legalMoves(state) and terminal(state) quickly.
Watch the merge invariant carefully: "same height OR same top color" — fail mode is forgetting the alternative branch.

STUDY HERE: topic 8, Grids and Simulation

Also fits: Design, Streaming and Scale (listed there as SKIP, no pages). Site tags: object-design, grid, simulation, medium

MEDIUM | 1Point3Acres | phone-screen, tech-screen | 2 of 3

Tic-Tac-Toe System with K Players and 3-in-a-Row Win

ASKED 2026-05-18 | 155 DAYS AGO | LAST 6 MONTHS

STUDY: TALK **ALL TIME 18.1** **30 DAYS 0** **PRIORITY 154/228**

ABOUT THIS QUESTION

Topic: Grids and Simulation

Optimal answer from: written and tested here for this exact question

LeetCode match: LC 348 Design Tic-Tac-Toe (related)

Frequency: Single report | Last asked: 2026-05-18 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/tic-tac-toe-k-players-three-in-a-row

QUESTION

Design a Tic-Tac-Toe system for **k** players on an **n x n** board. After each move, return the game status; any player wins immediately after forming 3 consecutive marks in a straight line.

Input parameters: k players and an n x n board.
Players take turns placing their marks on empty cells.
A player wins as soon as they form 3 consecutive marks in a straight line, regardless of how large n is.
After each move, return the correct game status, such as Game Over or <player> won.
Discuss implementation details and time complexity before coding.

APPROACH

Implement the classic Tic-Tac-Toe design problem, then generalize from 2 players / 3x3 to k players / n x n.
Practice local directional win checks with boundary guards.
Write tests for corners, edges, duplicate moves, and simultaneous longer lines.

HOW TO THINK ABOUT IT

A move can only create a win through the cell just played, so after each move look only along the four lines through that cell (row, column, both diagonals) and count how far the same player's marks extend in both directions. With a target of 3 that is a handful of cells per move, instead of rescanning the whole board.

Basis: LeetCode 348, Design Tic-Tac-Toe, generalised to k players, an n x n board and a fixed run length of 3.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def status(player):
    for r, c in all_cells(board):
        for dr, dc in DIRECTIONS:
            cells = line_of_3(r, c, dr, dc)
            if all(board[x][y] == player
                   for x, y in cells):
                return "WIN"
    if is_full(board):
        return "DRAW"
    return "PLAYING"
```

TIME	SPACE
O(n^2) per move	O(n^2)

WHY IT IS SLOW

Scans O(n^2) cells per move. Only lines through the cell just played can have changed, so checking four directions around it is O(1).

OPTIMAL SOLUTION · Check the four lines through the new mark

TIME	SPACE
O(1) per move (O(target) when generalised)	O(n^2)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 What changes when the winning length is a parameter w?

Replace the constant 3 by w in the run check: the cost per move becomes O(w) because each of the four directions walks at most w - 1 cells each way. The board size does not matter any more, which is the point of checking locally.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
# HELPERS for the brute force
DIRECTIONS = ((0, 1), (1, 0), (1, 1), (1, -1))


# BRUTE FORCE (say it first): rescan the whole board
class TicTacToeBrute:
    def __init__(self, n, k, target=3):
        self.n, self.k, self.target = n, k, target
        self.board = [[0] * n for _ in range(n)]
        self.moves = 0
        self.over = False

    def winner(self):
```

```
        # every cell, every direction: is there a run?
        n, t, b = self.n, self.target, self.board
        for i in range(n):
            for j in range(n):
                for di, dj in DIRECTIONS:
                    cells = [(i + di * s, j + dj * s)
                             for s in range(t)]
                    inside = all(0 <= x < n and 0 <= y < n
                                 for x, y in cells)
                    if inside and b[i][j] and all(
                            b[x][y] == b[i][j]
                            for x, y in cells):
                        return b[i][j]
        return 0

    def move(self, player, r, c):
        turn = self.moves % self.k + 1
        if self.over or player != turn:
            raise ValueError("not this player's turn")
        if not (0 <= r < self.n and 0 <= c < self.n):
            raise ValueError("illegal cell")
        if self.board[r][c]:
            raise ValueError("illegal cell")
        self.board[r][c] = player
        self.moves += 1
        if self.winner():
            self.over = True
            return f"player {player} wins"
        if self.moves == self.n * self.n:
            self.over = True
            return "draw"
        return "continue"
```

```
# BASE QUESTION (optimal): Check the four lines through
# the new mark
# FOLLOW-UP 1: What changes when the winning length is a
# parameter w?
class TicTacToe:
    def __init__(self, n, k, target=3):
        self.n, self.k, self.target = n, k, target
        self.board = [[0] * n for _ in range(n)]
        self.moves = 0
        self.over = False

    def inside(self, r, c):
        return 0 <= r < self.n and 0 <= c < self.n

    def run_length(self, r, c, dr, dc):
        # marks of the same player in a line through (r,
        # c)
        player, run = self.board[r][c], 1
        for sign in (1, -1):
            a, b = r + sign * dr, c + sign * dc
            while self.inside(a, b) and \
                    self.board[a][b] == player:
                run += 1
                a, b = a + sign * dr, b + sign * dc
        return run

    def move(self, player, r, c):
        # whose turn is it? players go 1, 2, ..., k, 1,
        # ...
        turn = self.moves % self.k + 1
        if self.over or player != turn:
            raise ValueError("not this player's turn")
        if not self.inside(r, c) or self.board[r][c]:
            raise ValueError("illegal cell")
        self.board[r][c] = player
        self.moves += 1
        # only lines through the new mark can have changed
        for dr, dc in DIRECTIONS:
            run = self.run_length(r, c, dr, dc)
            # FOLLOW-UP 1: win with `target` in a row
            if run >= self.target:
                self.over = True
                return f"player {player} wins"
        if self.moves == self.n * self.n:
            self.over = True
            return "draw"
        return "continue"


if __name__ == "__main__":
    g = TicTacToe(3, 2)
    plan = [(1, 0, 0), (2, 1, 0), (1, 0, 1),
            (2, 1, 1), (1, 0, 2)]
    results = [g.move(p, r, c) for p, r, c in plan]
    assert results[:-1] == ["continue"] * 4
    assert results[-1] == "player 1 wins"
    # three players on a 4 x 4 board: player 2 wins down
    h = TicTacToe(4, 3)
    seq = [(1, 0, 0), (2, 0, 3), (3, 3, 3), (1, 1, 0),
           (2, 1, 3), (3, 3, 0), (1, 2, 1), (2, 2, 3)]
    out = [h.move(p, r, c) for p, r, c in seq]
    assert out[-1] == "player 2 wins"
    # wrong turn and used cell are rejected
    for bad in ((2, 0, 0), (1, 5, 5)):
        try:
            TicTacToe(3, 2).move(*bad)
            raise AssertionError(bad)
        except ValueError:
            pass
    # the win length is a parameter: 4 in a row
    w = TicTacToe(5, 2, target=4)
    for p, r, c in [(1, 0, 0), (2, 1, 0), (1, 0, 1),
                    (2, 1, 1), (1, 0, 2), (2, 1, 2)]:
        assert w.move(p, r, c) == "continue"
    # the brute force agrees with the fast version
    import random
    for _ in range(100):
        fast, slow = TicTacToe(4, 3), TicTacToeBrute(4, 3)
        cells = [(r, c) for r in range(4)
                 for c in range(4)]
        random.shuffle(cells)
        for i, (r, c) in enumerate(cells):
            p = i % 3 + 1
            got = fast.move(p, r, c)
            assert got == slow.move(p, r, c)
            if got != "continue":
                break
```

NOTES FROM CANDIDATES

Check four directions through the new mark: horizontal, vertical, main diagonal, anti-diagonal. Count contiguous marks on both sides; win if total reaches 3.

Avoid scanning the whole board after every move unless n is tiny. The local directional check is O(1) for the fixed 3-in-a-row target, or O(target) if generalized.

Clarify invalid moves, turn order, draw state, and whether 3-in-a-row must be exactly 3 or at least 3.

PREPARATION

Implement the classic Tic-Tac-Toe design problem, then generalize from 2 players / 3x3 to k players / n x n.
Practice local directional win checks with boundary guards.
Write tests for corners, edges, duplicate moves, and simultaneous longer lines.

STUDY HERE: topic 8, Grids and Simulation

Also fits: Greedy and Math (listed there as SKIP, no pages). Site tags: simulation, algorithm, medium

MEDIUM | 1Point3Acres | onsite-coding | 3 of 3

Gate Entry / Exit Time Resolution with Tie-Break Rules

ASKED 2026-05-04 | 169 DAYS AGO | LAST 6 MONTHS

STUDY: TALK **ALL TIME 27.4** **30 DAYS 0** **PRIORITY 131/228**

ABOUT THIS QUESTION

Topic: Grids and Simulation

Optimal answer from: written and tested here for this exact question

LeetCode match: LC 2534 Time Taken to Cross the Door

Frequency: Single report | Last asked: 2026-05-04 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/gate-entry-exit-scheduler

QUESTION

Onsite coding round 2: given an array indexed by person id, each containing **`(timestamp, action in {enter, exit})`**, return per-person the actual time they passed through the gate after applying tie-break rules between simultaneous enters and exits.SWEsimulationalgorithmmediumFrequencySingle reportLast asked2026-05-04Stageonsite-coding

APPROACH

Translate the four priority rules into an explicit state machine (prev_state in {enter, exit, none}) before writing any code.
Practice scheduling problems that pair a per-tick simulation with an event-driven optimization (CPU scheduler, traffic light, calendar conflict resolver).
Stress-test corner cases: same timestamp containing both an enter and an exit for the same person, all-enters or all-exits batches, and very large gaps between two events.

INTERVIEW-DAY PLAN

CLARIFY FIRST

Ask for the exact tie-break rules in the four cases (previous second unused, entered, exited, and which wins), whether one person passes per second, and how ties between equal arrival times break (by id).

BRUTE FORCE TO SAY OUT LOUD

Simulate every second between the first and last timestamp. Works, but its cost grows with the size of the timestamps.

OPTIMAL IDEA

INTERVIEW-DAY PLAN

Sort people by (time, id) and keep two queues. At each step apply the previous-state rule to choose a queue. When both queues are empty jump to the next arrival time and reset the previous state to idle if at least one second was skipped. Linear after the sort.

DRY-RUN THESE CASES

Same-time enter and exit, an idle gap of exactly one second and of many seconds, all enters, all exits.

LIKELY FOLLOW-UPS

Several gates, priority classes, or answering online.

PITFALLS

Forgetting to reset the state across an idle gap. Letting the clock go backward.

HOW TO THINK ABOUT IT

Write the rule as a small state machine first: the state is what the gate did in the previous second (entered, exited or idle), and the rule picks who goes next from that state. Then replace the second-by-second loop with a loop over people: when nobody is waiting jump the clock forward, and remember that an idle gap resets the state to idle.

Basis: LeetCode 2534, Time Taken to Cross the Door. Assumption: one person passes per second; if an enter and an exit compete, the door's previous-second use decides (enter if it was an enter, otherwise exit, and exit if the door was idle).

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def brute(arrivals):
    t, prev = 0, None
    enter_q, exit_q = deque(), deque()
    while pending(arrivals, enter_q, exit_q):
        add_arrivals(t, arrivals, enter_q,
                     exit_q)
        queue = choose(prev, enter_q, exit_q)
        if queue:
            person = queue.popleft()
            passed[person] = t
            prev = name_of(queue)
        t += 1
    return passed
```

TIME	SPACE
O(T + n)	O(n)
where T is the last timestamp	

WHY IT IS SLOW

It runs through every second between the first and last timestamp, so a timestamp of 10^9 means a billion iterations. Jumping to the next arrival when the queues are empty makes it O(n log n).

OPTIMAL SOLUTION · Event-driven simulation with two queues

TIME	SPACE
O(n log n) for the sort, O(n) simulation	**O(n)**

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 Why does the optimization need care across an empty gap?

If the door sits idle for even one second the previous state becomes 'none', which favours exits again. Jumping time forward without resetting the state would wrongly keep 'enter' priority from before the gap. Compare t with last_used + 1 to know whether the gap really was idle.

FOLLOW-UP 2 What is the complexity, and what if events come online?

Offline it is O(n log n) for the sort, then O(n) for the simulation, independent of how large the timestamps are. Online, keep the two queues and the previous-state variable and advance the clock on every arrival or tick; the work per person stays O(1).

FOLLOW-UP 3 Timestamps go up to 10^9.

That is why the event-driven loop jumps to the next arrival: its running time is O(n log n) and does not depend on the size of the timestamps. Remember to reset the previous state to idle across a gap.

FOLLOW-UP 4 The tie-break when the gate was idle differs in this version of the question.

Keep the simulation and make the idle winner a parameter: only the line that picks the queue changes. With 'exit first' it equals the standard rule; with 'enter first' an enter and an exit arriving together swap places.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
from collections import deque


# BRUTE FORCE (say it first): Simulate every second | time
# O(T + n) where T is the last timestamp
def gate_ticks(events):
    n = len(events)
    order = sorted(
        range(n), key=lambda i: (events[i][0], i)
    )
    res, ent, ext, prev, t, k = (
        [None] * n,
        deque(),
        deque(),
        None,
        0,
        0,
    )
    while any(r is None for r in res):
        while k < n and events[order[k]][0] <= t:
            (
                ent
                if events[order[k]][1] == "enter"
                else ext
            ).append(order[k])
            k += 1
        if ent and ext:
            q = ent if prev == "enter" else ext
        else:
            q = ent or ext or None
        if q is None:
            prev = None
        else:
            p = q.popleft()
            res[p], prev = t, events[p][1]
        t += 1
    return res


# BASE QUESTION (optimal): Event-driven simulation with
# two queues
# FOLLOW-UP 4: The tie-break when the gate was idle
# differs in this version of the ...
def gate_pass_times(events, idle_first="exit"):
    # events[i] = (arrival time, "enter" or "exit")
    n = len(events)
    # people sorted by arrival, then by id
    order = sorted(range(n),
                   key=lambda i: (events[i][0], i))
    queues = {"enter": deque(), "exit": deque()}
    result = [0] * n
    t, i, prev, last_used = 0, 0, None, -10
    while i < n or queues["enter"] or queues["exit"]:
        if not queues["enter"] and not queues["exit"]:
            # nobody waits: jump the clock to the next
            # arrival
            t = max(t, events[order[i]][0])
            # a gap of idle time forgets the previous
            # action
            if t > last_used + 1:
                prev = None
        # everyone who arrived by now joins their line
        while i < n and events[order[i]][0] <= t:
            person = order[i]
            queues[events[person][1]].append(person)
            i += 1
        both = queues["enter"] and queues["exit"]
        if both:
            # FOLLOW-UP 4: idle tie-break set by
            # idle_first
            # tie: continue the direction used last second
            # (when idle, use idle_first)
            want = prev or idle_first
        else:
            want = "enter" if queues["enter"] else "exit"
        person = queues[want].popleft()
        result[person] = t
        prev, last_used = want, t
        t += 1
    return result
```

```
if __name__ == "__main__":
    # one enter and one exit arrive together: exit first
    ev = [(0, "enter"), (0, "exit")]
    assert gate_pass_times(ev) == [1, 0]
    # with "enter" winning the idle tie it swaps
    assert gate_pass_times(ev, "enter") == [0, 1]
    # the same direction continues while people keep
    # coming
```

```
    ev2 = [(0, "enter"), (0, "enter"), (1, "exit")]
    assert gate_pass_times(ev2) == [0, 1, 2]
    # a long idle gap resets the state
    ev3 = [(0, "enter"), (10, "enter"), (10, "exit")]
    assert gate_pass_times(ev3) == [0, 11, 10]
    # huge timestamps cost nothing extra
    far = [(10 ** 9, "enter"), (10 ** 9, "exit")]
    assert gate_pass_times(far) == [10 ** 9 + 1, 10 ** 9]
```

NOTES FROM CANDIDATES

The clean implementation uses two queues (pending enters, pending exits) plus a prev_state variable; at each tick peek both queues, apply the four rules, advance, and update prev_state.
A naive tick-by-tick simulation walks every integer time unit from min to max timestamp; the expected optimization is to jump directly to the next event timestamp in the input, dropping complexity to O(n).
The optimization has to carry the "prior moment" state across empty gaps correctly (a 100-tick gap still counts as prev_state = none).
Walk through one or two examples on the whiteboard before coding — interviewers in this family prefer candidates who fully verbalize the precedence rules before writing.

PREPARATION

Translate the four priority rules into an explicit state machine (prev_state in {enter, exit, none}) before writing any code.
Practice scheduling problems that pair a per-tick simulation with an event-driven optimization (CPU scheduler, traffic light, calendar conflict resolver).
Stress-test corner cases: same timestamp containing both an enter and an exit for the same person, all-enters or all-exits batches, and very large gaps between two events.