GOOGLE | 1POINT3ACRES | TOPIC 9 OF 10 | LESS COMMON TOPIC

Greedy and Math

2 questions to study here, each on its own page, with the full interview answer and code

12 more fit this topic but are covered in another topic: they are listed as SKIP at the end of the contents

2
HARD

How to use

- Questions run from Easy to Medium to Hard, newest last-asked date first inside each level.
- PDFs run from the one with the most questions to the fewest. A question has pages in only one PDF; the others list it as SKIP
- Every page has < Prev | Contents | Next > at the bottom. Tap a title in the Contents to jump to it.
- Select code and press Cmd/Ctrl+C to copy it.

Go to Contents >

Contents

Greedy and Math | 2 questions | tap a title to jump to it

H	1. Wordle-Style 5-Letter Minimum-Guess Strategy		**2025-12-15**
H	2. Find All Bad Pairs Using runTest Oracle		**2025-11-27**

SKIP: these fit this topic but have their pages in another topic (12)

M	SKIP	Bounded Knapsack Variant	**topic 2: Dynamic Programming**
M	SKIP	Connected Crop Arrangement on a Grid	**topic 1: Arrays, Hashing and Two Pointers**
M	SKIP	Path Minimizing the Maximum Cell (Swim in Rising Water)	**topic 3: Graphs and Search**
M	SKIP	Meeting Rooms / Interval Scheduling Variants	**topic 4: Heaps, Stacks and Queues**
M	SKIP	Cutting Ribbons (LC 1891)	**topic 6: Binary Search and Sorting**
M	SKIP	Target Expressions from Ordered Digits	**topic 2: Dynamic Programming**
M	SKIP	Greedy Longest-Match Dictionary Tokenizer	**topic 5: Trees and Tries**
M	SKIP	Gate Entry / Exit Time Resolution with Tie-Break Rules	**topic 8: Grids and Simulation**
H	SKIP	Maximum-Area Rectangle from Planar Points	**topic 1: Arrays, Hashing and Two Pointers**
H	SKIP	Huffman Tree Construction	**topic 5: Trees and Tries**
H	SKIP	Count Perfect Wake Numbers up to N	**topic 2: Dynamic Programming**
H	SKIP	Necklace Cut into Two Halves with Equal D/R Counts	**topic 1: Arrays, Hashing and Two Pointers**

STUDY HERE: topic 9, Greedy and Math

Site tags: algorithm-design, greedy, hard

HARD | 1Point3Acres | phone-screen, onsite-coding | 1 of 2

Wordle-Style 5-Letter Minimum-Guess Strategy

ASKED 2025-12-15 | 309 DAYS AGO | OLDER

STUDY: TALK **ALL TIME 42.2** **30 DAYS 0** **PRIORITY 85/228**

ABOUT THIS QUESTION

Topic: Greedy and Math

Optimal answer from: written and tested here for this exact question

Related in your app (a neighbouring problem, not this exact question): LeetCode #843 Guess the Word

LeetCode match: LC 843 Guess the Word (related)

Frequency: Single report | Last asked: 2025-12-15 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/wordle-min-guess-strategy

QUESTION

Phone screen coding: given a black-box oracle that tells you whether a specific letter is at a specific position in a hidden 5-letter word, design an algorithm to identify the word in as few oracle calls as possible. Follow-up assumes a dictionary of candidate words is available.

Hidden word: a 5-letter string from a..z.
Oracle: query(pos: int, letter: char) -> bool returns whether position pos (0-indexed) holds letter.
Return the hidden word.
Primary metric: minimize the number of query calls.

Brute-force baseline: 5 × 26 = 130 calls.
Follow-up

A dictionary of candidate 5-letter words is available; design the strategy to use bigram/letter-frequency information to cut the expected number of calls.

EXAMPLES

```
HiddenWord = "GRAPE"
query(0, 'A') -> false
query(0, 'G') -> true   // first letter found
...
```

APPROACH

Drill information-theoretic puzzles (egg drop, twenty questions) to internalize the entropy-split heuristic.
Practice maintaining a candidate set under positive/negative oracle feedback (filter on each query result).
Be ready with a one-liner on Wordle's known optimal first guess (SOARE / CRANE) for intuition; do not hard-code it.

HOW TO THINK ABOUT IT

Each call returns one bit, so you can never beat ceil(log2(number of candidates)) calls. Without a dictionary the best you can do is order the questions by likelihood and stop at the first yes for each position (and skip the 26th letter, because it is implied). With a dictionary, keep the list of words still possible and ask the (position, letter) question whose yes/no answer splits that list most evenly.

Basis: Written for this question; related to LeetCode 843, Guess the Word, but the oracle answers one position and one letter per call.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
word = ""
for i in range(5):
    for c in "abcdefghijklmnopqrstuvwxyz":
        if query(i, c):
            word += c
            break
```

TIME	SPACE
130 calls	O(1)

WHY IT IS SLOW

Always 130 calls. Asking likely letters first, skipping the implied last letter, or splitting a dictionary of candidates in half needs far fewer.

OPTIMAL SOLUTION · Frequency order, or entropy split when a dictionary is known

TIME	SPACE
at most 5 x 25 calls with... at most 5 x 25 calls without a dictionary; about log2(words) with one	O(words)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 What is the lower bound and when is it reached?

Every call returns at most one bit, so at least ceil(log2(candidates)) calls are needed. The entropy split reaches it only when a question can cut the candidates exactly in half; real dictionaries usually need a few more. A greedy scoring rule (closest to half) is enough to state; the interviewer mostly wants the entropy idea.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
# MAIN QUESTION: Wordle-Style 5-Letter Minimum-Guess
# Strategy
#
# Phone screen coding: given a black-box oracle that tells
# you whether a specific letter is at a specific position
# in a hidden 5-letter word, design an algorithm to
# identify the word in as few oracle calls as possible.
# Follow-up assumes a dictionary of candidate words is
# available. Hidden word: a 5-letter string from a..z.
# Oracle: query(pos: int, letter: char) -> bool returns
# whether position pos (0-indexed) holds letter. Return
# the hidden word. Primary metric: minimize the number of
# query calls. Brute-force baseline: 5 × 26 = 130 calls.
# Follow-up A dictionary of candidate 5-letter words is
# available; design the strategy to use bigram/letter-
# frequency information to cut the expected number of
# calls. Examples: HiddenWord = "GRAPE" query(0, 'A') ->
# false query(0, 'G') -> true // first letter found ...
#
# FOLLOW-UPS (answered below, in this order):
# 1. What is the lower bound and when is it reached?
#
# INTERVIEW NOTES
# Each call returns one bit, so you can never beat
# ceil(log2(number of candidates)) calls. Without a
# dictionary the best you can do is order the questions by
# likelihood and stop at the first yes for each position
# (and skip the 26th letter, because it is implied). With
# a dictionary, keep the list of words still possible and
# ask the (position, letter) question whose yes/no answer
# splits that list most evenly.
#
# BRUTE FORCE (say it first): Ask every position and
# letter
# Query all 5 x 26 pairs and read the word off the
# answers. Always 130 calls, no matter how lucky the word
# is. Time 130 calls. Space O(1).


import string


# BRUTE FORCE: 130 calls
def brute_force(query):
    # 130 calls at most: every position, every letter
```

```
    return "".join(
        next(c for c in string.ascii_lowercase
             if query(i, c))
        for i in range(5))


# HELPERS for the base question
# letters from most to least common in English words
BY_FREQUENCY = "etaoinshrdlcumwfgypbvkjxqz"


# FOLLOW-UP 1: What is the lower bound and when is it
# reached?
# Answer: Every call returns at most one bit, so at least
# ceil(log2(candidates)) calls are needed. The entropy
# split reaches it only when a question can cut the
# candidates exactly in half; real dictionaries usually
# need a few more.
# NEEDED BY THE BASE QUESTION: it is also the answer to
# the follow-up below
```

```
def split_score(words, i, c):
    # how evenly does asking (position i, letter c) split?
    yes = sum(w[i] == c for w in words)
    return min(yes, len(words) - yes)


# BASE QUESTION (optimal): Frequency order, or entropy
# split when a dictionary is known
# Idea: Without a dictionary: for each position ask
# letters from most to least common and stop at the first
# yes; if 25 letters said no the 26th is implied and needs
# no call.
def by_frequency(query):
    # no dictionary: common letters first, stop at the
    # first
    # yes, and skip the 26th letter (it must be the one
    # left)
    calls, word = 0, []
    for i in range(5):
        for k, c in enumerate(BY_FREQUENCY):
            if k == 25:
                word.append(c)
                break
            calls += 1
            if query(i, c):
```

```
                word.append(c)
                break
    return "".join(word), calls


# BASE QUESTION (optimal): Frequency order, or entropy
# split when a dictionary is known
# Idea: Without a dictionary: for each position ask
# letters from most to least common and stop at the first
# yes; if 25 letters said no the 26th is implied and needs
# no call.
def with_dictionary(query, dictionary):
    # keep the words that still fit every answer so far
    words, calls = list(dictionary), 0
    while len(words) > 1:
        options = {(i, w[i]) for w in words
                   for i in range(5)}
        # ask the question that cuts the list closest to
        # half
        i, c = max(options, key=lambda o: split_score(
            words, o[0], o[1]))
        calls += 1
        hit = query(i, c)
        words = [w for w in words if (w[i] == c) == hit]
    return words[0], calls


if __name__ == "__main__":
    def oracle(word):
        return lambda i, c: word[i] == c

    secret = "grape"
    assert brute_force(oracle(secret)) == secret
    word, calls = by_frequency(oracle(secret))
    assert word == secret and calls <= 125
    words = ["grape", "grade", "trade", "crane", "plane",
             "slate", "stare", "share"]
    for target in words:
        ask = oracle(target)
        found, used = with_dictionary(ask, words)
        assert found == target
        # about log2(8) = 3 questions suffice; allow a few
        # more
        assert used <= 6
    # lower bound: one bit per call, so ceil(log2(8)) = 3
```

```
# calls
assert min(with_dictionary(oracle(w), words)[1]
           for w in words) >= 3
```

NOTES FROM CANDIDATES

Without the dictionary, query positions in order; per position, query letters in descending frequency (e, t, a, o, i, n, ...) and stop as soon as one returns true.
With the dictionary, maintain a candidate set; at each step pick the (position, letter) query whose true/false outcome splits the candidate set most evenly (entropy maximization).
Mention information theory: each query yields <= 1 bit, so the absolute lower bound on calls is ceil(log2(|candidates|)).
Don't waste minutes on a closed-form optimum — the interviewer mostly wants you to articulate the entropy heuristic and pick one reasonable scoring function.

PREPARATION

Drill information-theoretic puzzles (egg drop, twenty questions) to internalize the entropy-split heuristic.
Practice maintaining a candidate set under positive/negative oracle feedback (filter on each query result).
Be ready with a one-liner on Wordle's known optimal first guess (SOARE / CRANE) for intuition; do not hard-code it.

STUDY HERE: topic 9, Greedy and Math

Site tags: algorithm-design, math-reasoning, hard

HARD | 1Point3Acres | onsite-coding | 2 of 2

Find All Bad Pairs Using runTest Oracle

ASKED 2025-11-27 | 327 DAYS AGO | OLDER

STUDY: TALK **NO LEETCODE FREQUENCY** **PRIORITY 192/228**

ABOUT THIS QUESTION

Topic: Greedy and Math

Optimal answer from: written and tested here for this exact question

Frequency: Single report | Last asked: 2025-11-27 | Roles: SWE, RS

Original: https://www.1point3acres.com/interview/problems/company/google/find-bad-pairs-with-oracle

QUESTION

PhD intern coding round 2: given a black-box oracle `runTest(S)` that returns `true` iff a subset `S` contains no bad pair, identify all bad pairs across N test items with as few oracle calls as possible.

Items t_1, ..., t_N are given.
Oracle: runTest(S: Set[Item]) -> bool returns true iff S contains zero bad pairs; otherwise false.
Output: every bad pair (t_i, t_j).
Minimize total oracle calls.

EXAMPLES

```
N=4, bad pairs = {(1,2)}.
runTest({1,2,3,4}) -> false -> at least one bad pair
    somewhere.
runTest({1,3,4})   -> true  -> 2 is in every bad pair
    containing items from this set.
... binary search to isolate.
```

APPROACH

Read the Wikipedia article on group testing and Dorfman's two-stage scheme.
Drill recursive binary partitioning on a candidate set with positive/negative oracle feedback.
Have a back-of-envelope analysis ready: "Total calls O(K log N) for K bad pairs and N items."

HOW TO THINK ABOUT IT

Testing every pair costs N(N-1)/2 calls. A single call on a whole set answers 'is there any bad pair inside', so you can discard good sets fast and binary-search for the culprit inside a bad one. Add the tests one at a time to groups known to be clean: if a group plus the new test fails, binary search the group for a partner; otherwise the new test joins it.

Basis: Adaptive group testing; the same method as the 'bad pairs with runTest' question in the PracHub set.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
bad = []
for i in range(n):
    for j in range(i + 1, n):
        if not runTest({i, j}):
            bad.append((i, j))
```

TIME	SPACE
O(N^2) calls	O(1)

WHY IT IS SLOW

Always N(N - 1) / 2 calls. Testing whole groups and binary searching a failing group finds each bad pair in about log N calls.

OPTIMAL SOLUTION · Group the tests and binary search for partners

TIME	SPACE
O(K log N + N * G) about calls (K bad pairs, G groups)	O(N)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 How many calls are needed in the best case and the worst case?

With no bad pairs it takes N - 1 calls, one per test after the first. With K bad pairs each one costs about log N calls to locate, so roughly K log N plus the cost of building clean groups. The brute force is always N(N-1)/2, which is far larger as soon as bad pairs are rare.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
# MAIN QUESTION: Find All Bad Pairs Using runTest Oracle
#
# PhD intern coding round 2: given a black-box oracle
# runTest(S) that returns true iff a subset S contains no
# bad pair, identify all bad pairs across N test items
# with as few oracle calls as possible. Items t_1, ...,
# t_N are given. Oracle: runTest(S: Set[Item]) -> bool
# returns true iff S contains zero bad pairs; otherwise
# false. Output: every bad pair (t_i, t_j). Minimize total
# oracle calls. Examples: N=4, bad pairs = {(1,2)}.
# runTest({1,2,3,4}) -> false -> at least one bad pair
# somewhere. runTest({1,3,4}) -> true -> 2 is in every bad
# pair containing items from this set. ... binary search
# to isolate.
#
# FOLLOW-UPS (answered below, in this order):
# 1. How many calls are needed in the best case and the
# worst case?
```

```
#
# INTERVIEW NOTES
# Testing every pair costs N(N-1)/2 calls. A single call
# on a whole set answers 'is there any bad pair inside',
# so you can discard good sets fast and binary-search for
# the culprit inside a bad one. Add the tests one at a
# time to groups known to be clean: if a group plus the
# new test fails, binary search the group for a partner;
# otherwise the new test joins it.
#
# BRUTE FORCE (say it first): Call runTest on every pair
# Test every pair of items with runTest and keep the
# failing pairs. Always N(N-1)/2 calls. Time O(N^2) calls.
# Space O(1).


import itertools


# BRUTE FORCE: ask about every pair
def bad_pairs_brute(items, run_test):
    # every pair costs a call: N * (N - 1) / 2 calls
    pairs = itertools.combinations(items, 2)
    return {(a, b) for a, b in pairs
            if not run_test({a, b})}
```

```
# HELPERS for the base question
def first_partner(group, t, run_test):
    # binary search: the shortest prefix of the group that
    # fails together with t; its last member is a partner
    lo, hi = 1, len(group)
    while lo < hi:
        mid = (lo + hi) // 2
        if run_test(set(group[:mid]) | {t}):
            lo = mid + 1
        else:
            hi = mid
    return group[lo - 1]


def partners(group, t, run_test):
    # all members of the group that are bad with t
    rest, found = list(group), []
    while rest and not run_test(set(rest) | {t}):
        x = first_partner(rest, t, run_test)
        found.append(x)
        # take it out and test the others again
        rest.remove(x)
    return found


# BASE QUESTION (optimal): Group the tests and binary
# search for partners
# Idea: Keep groups of tests that pass together.
def find_bad_pairs(items, run_test):
    # groups of items that pass together (no bad pair
    # inside)
    groups, pairs = [], set()
    for t in items:
        placed = False
        for group in groups:
            found = partners(group, t, run_test)
            for x in found:
                pairs.add((min(x, t), max(x, t)))
            # t joins the first group with no partner for
            # it
            if not found and not placed:
                group.append(t)
                placed = True
        if not placed:
            groups.append([t])
    return pairs


if __name__ == "__main__":
    def make_oracle(bad):
        calls = [0]

        def run_test(subset):
            calls[0] += 1
            ordered = sorted(subset)
            pairs = itertools.combinations(ordered, 2)
            return not any(p in bad for p in pairs)
        return run_test, calls

    bad = {(3, 17), (5, 29)}
    run, calls = make_oracle(bad)
    assert find_bad_pairs(list(range(30)), run) == bad
    # far fewer than the 435 calls of the brute force
    assert calls[0] < 435 // 2
    # no bad pairs: about one call per item
    run, calls = make_oracle(set())
    assert find_bad_pairs(list(range(30)), run) == set()
    assert calls[0] <= 30
    run, _ = make_oracle({(1, 2)})
    assert bad_pairs_brute([1, 2, 3, 4], run) == {(1, 2)}
```

NOTES FROM CANDIDATES

Standard "group testing" / Dorfman testing problem. Walk through the analogy with COVID pooled tests for intuition. Algorithm sketch:

Test the full set. If true, no bad pairs.
Otherwise, binary-search to find one bad pair by partitioning the set and testing halves.
Remove the items that participate in the located bad pair; repeat.

Tighter bound: with up to K bad pairs, total queries are O(K log N) using adaptive group testing.
A naive O(N²) pairwise scheme works but the interviewer expects you to beat it with adaptive splitting.
Verbalize correctness: each binary-search step rules out half the candidates; you can always reduce the search space until exactly one bad pair is isolated.

PREPARATION

Read the Wikipedia article on group testing and Dorfman's two-stage scheme.
Drill recursive binary partitioning on a candidate set with positive/negative oracle feedback.
Have a back-of-envelope analysis ready: "Total calls O(K log N) for K bad pairs and N items."