More practice - Grids and simulation

3 questions to study here, each on its own page, with the full interview answer and code

Nothing to skip: every question that fits this topic is covered here

3
HARD

How to use

- Questions run from Easy to Medium to Hard, newest last-asked date first inside each level.
- PDFs run from the one with the most questions to the fewest. A question has pages in only one PDF; the others list it as SKIP
- Every page has < Prev | Contents | Next > at the bottom. Tap a title in the Contents to jump to it.
- Select code and press Cmd/Ctrl+C to copy it.

Go to Contents >

Contents

More practice - Grids and simulation | 3 questions | tap a title to jump to it

H	1. Simulate In-Place Cellular State Updates	PracHub	2026-05-04
H	2. Count Overlapping Rectangle Updates on a Grid	PracHub	2026-07-10
H	3. Babylon Tile-Stack Merge Game	1Point3Acres	2025-12-29

STUDY PATH: question 157 of 167 (more practice)

Priority score 2.44 | Grids and simulation | asked 169 days ago

HARD | IN APP LC 289 | PracHub | Technical Screen | 1 of 3

Simulate In-Place Cellular State Updates

POSTED 2026-05-04 | 169 DAYS AGO | LAST 6 MONTHS

STUDY: TALK | ALL TIME 23.4 | 30 DAYS 0 | PRIORITY 143/228

ABOUT THIS QUESTION

Topic: More practice - Grids and simulation

Optimal answer from: the interview site's own solution

Also in your app: LeetCode #289 Game of Life

LeetCode match: LC 289 Game of Life (inferred from overview)

Posted on PracHub: 2026-05-04 (updated 2026-05-07)

Original: https://prachub.com/coding-questions/simulate-in-place-cellular-state-updates

QUESTION

Overview: This question evaluates a candidate's ability to implement grid-based state transitions and manage in-place updates with strict space constraints, assessing skills in matrix manipulation, neighbor traversal, and correctness under simultaneous update rules.

APPROACH

Encode old and new state together in each cell (extra bits or special values like 2 and 3) so neighbors still read the old state, then finalize in a second pass. O(mn) time, O(1) extra space.

HOW TO THINK ABOUT IT

'Update every cell at once, in place' means a cell's new value must not disturb its neighbors' reads. The trick is to store both the old and the new state in one cell (two bits), then decode in a second pass. Whenever you need old and new values together with no extra grid, think 'pack two states into one number'.

Basis: LeetCode 289, Game of Life.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def step(board):
    old = [row[:] for row in board]
    for r, c in all_cells(old):
        live = live_neighbours(old, r, c)
        alive = old[r][c] == 1
        if alive and live in (2, 3):
            board[r][c] = 1
        elif not alive and live == 3:
            board[r][c] = 1
        else:
            board[r][c] = 0
```

BRUTE FORCE

TIME	SPACE
O(m * n)	O(m * n)

WHY IT IS SLOW

Time O(m * n) is already optimal; the cost is a second full board, O(m * n) extra space. Storing the next state in a spare bit of each cell removes the copy.

OPTIMAL SOLUTION · Two-bit encoding in the same board

TIME	SPACE
O(m * n)	O(1)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 The board is infinite, or huge and mostly empty.

Store only the live cells in a set. A cell's next state depends on its live neighbor count, so count how many live neighbors every candidate cell has with a Counter and keep cells with 3, or with 2 if already live. Work is proportional to the number of live cells.

FOLLOW-UP 2 Why is bit 0 used when counting neighbors?

Neighbors that were already processed in this pass carry their next state in bit 1. Masking with & 1 always reads the original generation, which is the whole point of the encoding.

FOLLOW-UP 3 Cells wrap around the edges (a torus).

Replace the bounds check with modular indexing: (r + dr) % m and (c + dc) % n. Be careful on very small boards where the same neighbor can be reached through two offsets.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
# MAIN QUESTION: Simulate In-Place Cellular State Updates
#
# This question evaluates a candidate's ability to
# implement grid-based state transitions and manage in-
# place updates with strict space constraints, assessing
# skills in matrix manipulation, neighbor traversal, and
# correctness under simultaneous update rules.
#
# FOLLOW-UPS (answered below, in this order):
# 1. The board is infinite, or huge and mostly empty.
# 2. Why is bit 0 used when counting neighbors?
# 3. Cells wrap around the edges (a torus).
#
# INTERVIEW NOTES
# 'Update every cell at once, in place' means a cell's new
```

```
# value must not disturb its neighbors' reads. The trick
# is to store both the old and the new state in one cell
# (two bits), then decode in a second pass. Whenever you
# need old and new values together with no extra grid,
# think 'pack two states into one number'.
#
# BRUTE FORCE (say it first): Copy the board, read from
# the copy
# Compute the next generation from an untouched copy of
# the board. It is the simplest correct answer, and the
# extra copy is exactly what the in-place version removes.
# Time O(m * n). Space O(m * n).


from collections import Counter


# BRUTE FORCE (say it first): Copy the board, read from
# the copy | time O(m * n)
def _life_neighbors(board, r, c, live):
    return sum(
        live(board[nr][nc])
        for nr in range(r - 1, r + 2)
        for nc in range(c - 1, c + 2)
        if (nr, nc) != (r, c)
        and 0 <= nr < len(board)
        and 0 <= nc < len(board[0])
    )
```

```
def life_brute(board):
    old = [row[:] for row in board]
    for r in range(len(board)):
        for c in range(len(board[0])):
            n = _life_neighbors(
                old, r, c, lambda v: v == 1
            )
            board[r][c] = (
                1
                if (old[r][c] == 1 and n in (2, 3))
                or (old[r][c] == 0 and n == 3)
                else 0
            )
    return board


# HELPERS for the base question
AROUND = [(dr, dc) for dr in (-1, 0, 1)
          for dc in (-1, 0, 1) if (dr, dc) != (0, 0)]


def live_neighbors(board, r, c, wrap=False):
    m, n = len(board), len(board[0])
    count = 0
    for dr, dc in AROUND:
        a, b = r + dr, c + dc
        if wrap:
            # a torus: the edges are glued together
            a, b = a % m, b % n
        elif not (0 <= a < m and 0 <= b < n):
            continue
        # bit 0 is the OLD state, even for updated cells
        count += board[a][b] & 1
    return count


# BASE QUESTION (optimal): Two-bit encoding in the same
# board
# Idea: Bit 0 holds the current state and bit 1 holds the
# next state. First pass: count neighbors using only bit 0
# (value & 1) and set bit 1 by the rules.
def step_in_place(board, wrap=False):
    m, n = len(board), len(board[0])
```

```
    for r in range(m):
        for c in range(n):
```

```
            live = live_neighbors(board, r, c, wrap)
            alive = board[r][c] & 1
            # bit 1 stores the NEW state
            if live == 3 or (alive and live == 2):
                board[r][c] |= 2
    for r in range(m):
        for c in range(n):
            # shift the new state into place
            board[r][c] >>= 1
    return board


# FOLLOW-UP 1: The board is infinite, or huge and mostly
# empty.
# Answer: Store only the live cells in a set. A cell's
# next state depends on its live neighbor count, so count
# how many live neighbors every candidate cell has with a
# Counter and keep cells with 3, or with 2 if already
# live.
def step_sparse(live):
    # infinite or huge board: only the live cells are
    # stored
    counts = Counter()
    for r, c in live:
        for dr, dc in AROUND:
            counts[(r + dr, c + dc)] += 1
    return {cell for cell, k in counts.items()
            if k == 3 or (k == 2 and cell in live)}


if __name__ == "__main__":
    board = [[0, 1, 0], [0, 0, 1], [1, 1, 1], [0, 0, 0]]
    want = [[0, 0, 0], [1, 0, 1], [0, 1, 1], [0, 1, 0]]
    assert step_in_place(board) == want
    # a blinker flips between a row and a column
    blinker = {(1, 0), (1, 1), (1, 2)}
    flipped = {(0, 1), (1, 1), (2, 1)}
    assert step_sparse(blinker) == flipped
    assert step_sparse(step_sparse(blinker)) == blinker
    # a block never changes
    block = {(0, 0), (0, 1), (1, 0), (1, 1)}
    assert step_sparse(block) == block
    # the sparse and dense versions agree on a glider
    glider = {(0, 1), (1, 2), (2, 0), (2, 1), (2, 2)}
    grid = [[0] * 6 for _ in range(6)]
```

```
    for r, c in glider:
        grid[r][c] = 1
    dense = step_in_place(grid)
    assert {(r, c) for r in range(6) for c in range(6)
            if dense[r][c]} == step_sparse(glider)
    # on a torus the left and right edges touch
    ring = [[1, 0, 1]]
    assert live_neighbors(ring, 0, 0, wrap=True) >= 1
```

FUNCTION AND TEST CASES

Signature

```
def solution(grid: list[list[int]]) -> list[list[int]]:
```

Sample call

```
solution([[0, 1, 0], [0, 0, 1], [1, 1, 1], [0, 0, 0]])
```

Test cases

Input	Expected
`([[0, 1, 0], [0, 0, 1], [1, 1, 1], [0, 0, 0]],)`	`[[0, 0, 0], [1, 0, 1], [0, 1, 1], [0, 1, 0]]`
`([[1, 1], [1, 1]],)`	`[[1, 1], [1, 1]]`

STUDY PATH: question 158 of 167 (more practice)

Priority score 2.22 | Grids and simulation | asked 102 days ago

HARD | IN APP LC 2536 | PracHub | Technical Screen | 2 of 3

Count Overlapping Rectangle Updates on a Grid

POSTED 2026-07-10 | 102 DAYS AGO | LAST 6 MONTHS

STUDY: TALK | ALL TIME 18.1 | 30 DAYS 0 | PRIORITY 151/228

ABOUT THIS QUESTION

Topic: More practice - Grids and simulation

Optimal answer from: the interview site's own solution

Also in your app: LeetCode #2536 Increment Submatrices by One

LeetCode match: LC 2536 Increment Submatrices by One

Posted on PracHub: 2026-07-10 (updated 2026-07-18)

Original: https://prachub.com/coding-questions/count-overlapping-rectangle-updates-on-a-grid

QUESTION

Count Overlapping Rectangle Updates on a Grid Implement `range_add_counts(n, rectangles)`. Start with an `n x n` matrix of zeroes. Each rectangle is represented as `[row1, col1, row2, col2]` with inclusive, zero-based coordinates. For every rectangle, add one to each matrix cell inside that rectangle. Return the final matrix, where each cell therefore contains the number of input rectangles covering it.

Example

For `n = 3` and `rectangles = [[0, 0, 1, 1], [1, 1, 2, 2]]`, return:

```
[ [1, 1, 0], [1, 2, 1], [0, 1, 1] ]
```

Constraints

- `1 <= n <= 500`
- `0 <= len(rectangles) <= 100_000`
- `0 <= row1 <= row2 < n`
- `0 <= col1 <= col2 < n`

APPROACH

2D difference array: for each rectangle do +1 at (r1,c1), -1 at (r1,c2+1) and (r2+1,c1), +1 at (r2+1,c2+1); then take 2D prefix sums once. O(n^2 + q) instead of O(n^2*q).

HOW TO THINK ABOUT IT

Adding 1 to every cell of up to 100,000 rectangles one cell at a time is too slow. Record only the corners of each rectangle (+1 at the top-left, -1 just beyond the right and bottom edges, +1 at the bottom-right corner beyond), and rebuild the counts with one pass of 2D prefix sums. Each rectangle then costs O(1) instead of its area.

HOW TO THINK ABOUT IT

Basis: LeetCode 2536, Increment Submatrices by One (2D difference array).

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def range_add_counts(n, rectangles):
    grid = [[0] * n for _ in range(n)]
    for r1, c1, r2, c2 in rectangles:
        for r in range(r1, r2 + 1):
            for c in range(c1, c2 + 1):
                grid[r][c] += 1
    return grid
```

TIME	SPACE
O(R * area)	O(n^2)

WHY IT IS SLOW

O(R * area): 100,000 rectangles of area up to 250,000 is billions of cell updates. A 2D difference array touches four corners per rectangle and one prefix-sum pass rebuilds the counts, O(R + n^2).

OPTIMAL SOLUTION · 2D difference array and prefix sums

TIME	SPACE
O(R + n^2)	O(n^2)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 How many cells are covered by at least k rectangles?

Build the count matrix with the difference array and count the cells with value at least k. O(R + n^2).

FOLLOW-UP 2 The grid is huge (10^9 x 10^9) but there are few rectangles.

Coordinate-compress the corner coordinates and apply the same difference array on the compressed grid (at most (2R)^2 cells), or use a sweep line with a segment tree over columns. **Cost:** O(R log R) to sort the corner coordinates plus up to O(R^2) compressed cells, instead of 10^18 cells.

FOLLOW-UP 3 Rectangles are added and the count of a cell is queried online.

Use a 2D Fenwick tree with range update and point query (the difference-array idea turned into a BIT): O(log^2 n) per operation.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
# MAIN QUESTION: Count Overlapping Rectangle Updates on a
# Grid
#
# # Count Overlapping Rectangle Updates on a Grid
# Implement range_add_counts(n, rectangles). Start with an
# n x n matrix of zeroes. Each rectangle is represented as
# [row1, col1, row2, col2] with inclusive, zero-based
# coordinates. For every rectangle, add one to each matrix
# cell inside that rectangle. Return the final matrix,
# where each cell therefore contains the number of input
# rectangles covering it. ## Example For n = 3 and
# rectangles = [[0, 0, 1, 1], [1, 1, 2, 2]], return: text
# [ [1, 1, 0], [1, 2, 1], [0, 1, 1] ] ## Constraints - 1
# <= n <= 500 - 0 <= len(rectangles) <= 100_000 - 0 <=
# row1 <= row2 < n - 0 <= col1 <= col2 < n
#
# FOLLOW-UPS (answered below, in this order):
# 1. How many cells are covered by at least k rectangles?
# 2. The grid is huge (10^9 x 10^9) but there are few
# rectangles.
# 3. Rectangles are added and the count of a cell is
# queried online.
#
# INTERVIEW NOTES
# Adding 1 to every cell of up to 100,000 rectangles one
# cell at a time is too slow. Record only the corners of
# each rectangle (+1 at the top-left, -1 just beyond the
# right and bottom edges, +1 at the bottom-right corner
# beyond), and rebuild the counts with one pass of 2D
# prefix sums. Each rectangle then costs O(1) instead of
# its area.
#
# BRUTE FORCE (say it first): Add one to every cell of
# every rectangle
# Loop over each rectangle and increment all of its cells.
# Correct and the oracle; with 100,000 rectangles of area
# up to 250,000 it is far too slow. Time O(R * area).
# Space O(n^2).


# BRUTE FORCE (say it first): Add one to every cell of
# every rectangle | time O(R * area)
def range_add_brute(n, rects):
    g = [[0] * n for _ in range(n)]
    for r1, c1, r2, c2 in rects:
        for r in range(r1, r2 + 1):
            for c in range(c1, c2 + 1):
                g[r][c] += 1
    return g


# BASE QUESTION (optimal): 2D difference array and prefix
# sums
# Idea: For each rectangle add +1 at (r1, c1), -1 at (r1,
# c2 + 1), -1 at (r2 + 1, c1) and +1 at (r2 + 1, c2 + 1)
# in an (n + 1) x (n + 1) array.
def range_add_counts(n, rects):
    # rects are [r1, c1, r2, c2], inclusive. Mark only the
    # 4
    # corners of each rectangle in a difference array ...
    d = [[0] * (n + 1) for _ in range(n + 1)]
    for r1, c1, r2, c2 in rects:
        d[r1][c1] += 1
        d[r1][c2 + 1] -= 1
        d[r2 + 1][c1] -= 1
        d[r2 + 1][c2 + 1] += 1
    # ... then two prefix sums spread them over the cells
    for r in range(n + 1):
        for c in range(1, n + 1):
            d[r][c] += d[r][c - 1]
    for c in range(n + 1):
        for r in range(1, n + 1):
            d[r][c] += d[r - 1][c]
    return [row[:n] for row in d[:n]]


# FOLLOW-UP 1: How many cells are covered by at least k
# rectangles?
# Answer: Build the count matrix with the difference array
# and count the cells with value at least k. O(R + n^2).
def cells_at_least(n, rects, k):
    # how many cells are covered by k or more rectangles?
    counts = range_add_counts(n, rects)
    return sum(1 for row in counts for v in row if v >= k)


# FOLLOW-UP 2: The grid is huge (10^9 x 10^9) but there
# are few rectangles.
# Answer: Coordinate-compress the corner coordinates and
# apply the same difference array on the compressed grid
# (at most (2R)^2 cells), or use a sweep line with a
# segment tree over columns.
def cells_at_least_compressed(rects, k):
    # a huge grid, few rectangles: compress the
    # coordinates
    rows = sorted({r for r1, _, r2, _ in rects
                   for r in (r1, r2 + 1)})
    cols = sorted({c for _, c1, _, c2 in rects
                   for c in (c1, c2 + 1)})
    ri = {r: i for i, r in enumerate(rows)}
    ci = {c: i for i, c in enumerate(cols)}
    shape = (len(rows) + 1, len(cols) + 1)
    d = [[0] * shape[1] for _ in range(shape[0])]
    for r1, c1, r2, c2 in rects:
        a, b = ri[r1], ci[c1]
        e, f = ri[r2 + 1], ci[c2 + 1]
        d[a][b] += 1
        d[a][f] -= 1
        d[e][b] -= 1
        d[e][f] += 1
    total = 0
    run_col = [0] * (len(cols) + 1)
    for i in range(len(rows) - 1):
        row, run = [0] * (len(cols) + 1), 0
        for j in range(len(cols) - 1):
            run_col[j] += d[i][j]
            run += run_col[j]
            # each compressed cell stands for a block of
            # cells
            if run >= k:
                total += (rows[i + 1] - rows[i]) * (
                    cols[j + 1] - cols[j])
    return total


# FOLLOW-UP 3: Rectangles are added and the count of a
# cell is queried online.
# Answer: Use a 2D Fenwick tree with range update and
# point query (the difference-array idea turned into a
# BIT): O(log^2 n) per operation.
class Fenwick2D:
    """Online updates and point reads: O(log^2 n)."""

    def __init__(self, n):
        self.n = n
```

```
        self.t = [[0] * (n + 2) for _ in range(n + 2)]

    def _add(self, r, c, v):
        i = r + 1
        while i <= self.n + 1:
            j = c + 1
            while j <= self.n + 1:
                self.t[i][j] += v
                j += j & -j
            i += i & -i

    def add_rect(self, r1, c1, r2, c2):
        corners = ((r1, c1, 1), (r1, c2 + 1, -1),
                   (r2 + 1, c1, -1), (r2 + 1, c2 + 1, 1))
        for r, c, v in corners:
            self._add(r, c, v)

    def get(self, r, c):
        total, i = 0, r + 1
        while i > 0:
            j = c + 1
            while j > 0:
                total += self.t[i][j]
                j -= j & -j
            i -= i & -i
        return total


if __name__ == "__main__":
    rects = [[0, 0, 1, 1], [1, 1, 2, 2]]
    assert range_add_counts(3, rects) == [
        [1, 1, 0], [1, 2, 1], [0, 1, 1]]
    assert cells_at_least(3, rects, 2) == 1
    assert cells_at_least(3, rects, 1) == 7
    # the compressed version gives the same area
    assert cells_at_least_compressed(rects, 2) == 1
    assert cells_at_least_compressed(rects, 1) == 7
    big = [[0, 0, 10 ** 9, 10 ** 9], [5, 5, 10, 10]]
    assert cells_at_least_compressed(big, 2) == 36
    # online updates and point queries
    f = Fenwick2D(3)
    for r in rects:
        f.add_rect(*r)
    assert f.get(1, 1) == 2 and f.get(0, 2) == 0
```

FUNCTION AND TEST CASES

Signature

```
def range_add_counts(n: int, rectangles: list[list[int]])
-> list[list[int]]:
```

Sample call

```
range_add_counts(1, [])
```

Test cases

Input	Expected
`(1, [])`	`[[0]]`
`(1, [[0, 0, 0, 0]])`	`[[1]]`

STUDY PATH: question 159 of 167 (more practice)

Priority score 1.0 | Grids and simulation | asked 295 days ago | frequency Single report

HARD | 1Point3Acres | onsite-coding | 3 of 3

Babylon Tile-Stack Merge Game

ASKED 2025-12-29 | 295 DAYS AGO | OLDER

STUDY: TALK | NO LEETCODE FREQUENCY | PRIORITY 191/228

ABOUT THIS QUESTION

Topic: More practice - Grids and simulation

Optimal answer from: written and tested here for this exact question

Frequency: Single report | Last asked: 2025-12-29 | Roles: SWE

Original: https://www.1point3acres.com/interview/problems/company/google/babylon-stack-merge-game

QUESTION

Onsite NG round 1 (Munich office): implement a two-player tile-stack merge game ("Babylon") given the rules. Start state is 12 height-1 stacks of 4 colors (Y / W / G / B), 3 of each; a move places one stack on top of another under a constraint; the player who cannot move loses.

12 tiles initially, partitioned into 12 height-1 stacks.
4 colors: Y / W / G / B; exactly 3 tiles per color.
Two players alternate moves.
A move = pick two stacks and place one entirely on top of the other (merge).

Legal iff at least one of: (a) the two stacks have the same height, OR (b) their top-most tiles share the same color.
After merge: new height = sum of the two heights, new top color = color of the stack placed on top.

The player who has no legal move on their turn loses.
Implement the full game so it can be played out (legal-move detection, state transition, terminal check).

EXAMPLES

Start: 12 stacks of height 1, with tops Y Y Y W W W G G G B B B. A legal opening move: stack the first Y on the second Y (same color top), producing one height-2 stack of color Y plus 10 singletons.

APPROACH

Drill clean game-state implementations: e.g. Connect 4, Othello move generators, in 30 minutes.
Practice writing a generic minimax framework so you can plug in this game's legalMoves(state) and terminal(state) quickly.
Watch the merge invariant carefully: "same height OR same top color" — fail mode is forgetting the alternative branch.

HOW TO THINK ABOUT IT

Model a state as the list of stacks, each as (height, top colour). A legal move puts stack i on stack j when the heights are equal or the top colours are equal; the new stack has the summed height and the colour of the stack that was placed on top. The player with no legal move loses, so a state is winning when some move leads to a losing state. The order of stacks does not matter, so sort the list to share states in a cache.

Basis: Written for this question: a two-player impartial game (the player who cannot move loses), solved by memoised search over the multiset of stacks.

BRUTE FORCE

BRUTE FORCE · PSEUDO-CODE

```
def wins(position):
    for i, j in legal_moves(position):
        nxt = merge(position, i, j)
        if not wins(nxt):
            return True  # opponent loses
    return False
```

TIME	SPACE
exponential	O(depth)

WHY IT IS SLOW

Explores the same positions many times, exponential. Sorting the stacks and caching each position's result visits every distinct position once.

OPTIMAL SOLUTION · Negamax with memoisation on the sorted stack tuple

TIME	SPACE
O(states * moves^2)	O(states)

Code: see FULL INTERVIEW CODE below (the optimal solution is its main function).

FOLLOW-UP 1 How big is the real state space?

With 12 starting stacks the first position has at most 66 unordered pairs and the number of reachable multisets of (height, colour) stacks shrinks quickly because every merge removes a stack, so memoised search is feasible. In the interview build the clean state class, legal moves and play(move) first and add the solver only if there is time.

FULL INTERVIEW CODE · base question + every follow-up

One program for the whole question: the base answer and every follow-up above, built from small helper functions, with a comment above each step. Type it top to bottom in the interview; the demo at the bottom runs it.

```
# MAIN QUESTION: Babylon Tile-Stack Merge Game
#
```

```
# Onsite NG round 1 (Munich office): implement a two-
# player tile-stack merge game ("Babylon") given the
# rules. Start state is 12 height-1 stacks of 4 colors (Y
# / W / G / B), 3 of each; a move places one stack on top
# of another under a constraint; the player who cannot
# move loses. 12 tiles initially, partitioned into 12
# height-1 stacks. 4 colors: Y / W / G / B; exactly 3
# tiles per color. Two players alternate moves. A move =
# pick two stacks and place one entirely on top of the
# other (merge). Legal iff at least one of: (a) the two
# stacks have the same height, OR (b) their top-most tiles
# share the same color. After merge: new height = sum of
# the two heights, new top color = color of the stack
# placed on top. The player who has no legal move on their
# turn loses. Implement the full game so it can be played
# out (legal-move detection, state transition, terminal
# check).
#
# FOLLOW-UPS (answered below, in this order):
# 1. How big is the real state space?
#
# INTERVIEW NOTES
# Model a state as the list of stacks, each as (height,
# top colour). A legal move puts stack i on stack j when
# the heights are equal or the top colours are equal; the
# new stack has the summed height and the colour of the
# stack that was placed on top. The player with no legal
# move loses, so a state is winning when some move leads
# to a losing state. The order of stacks does not matter,
# so sort the list to share states in a cache.
#
# BRUTE FORCE (say it first): Plain recursion over every
# move
# Try every legal move recursively with no caching.
# Correct and fine for a few stacks; it recomputes the
# same positions over and over. Time exponential. Space
# O(depth).


import itertools


# BRUTE FORCE (say it first): Plain recursion over every
```

```
# move | time exponential
def babylon_brute(stacks):
    stacks = list(stacks)
    for i, j in itertools.permutations(
        range(len(stacks)), 2
    ):
        (h1, c1), (h2, c2) = stacks[i], stacks[j]
        if h1 == h2 or c1 == c2:
            rest = [
                s
                for k, s in enumerate(stacks)
                if k not in (i, j)
            ] + [(h1 + h2, c1)]
            if not babylon_brute(rest):
                return True
    return False


# HELPERS for the base question
class Babylon:
    # a stack is (height, colour of its top tile)
    def __init__(self, stacks=None):
        if stacks is None:
            # 12 tiles: Y, W, G, B three times each
            stacks = [(1, c) for c in "YWGB" * 3]
        self.stacks = list(stacks)

    def legal_moves(self):
        # put stack i on top of stack j
```

```
        moves = []
        for i, j in itertools.permutations(
                range(len(self.stacks)), 2):
            h1, c1 = self.stacks[i]
            h2, c2 = self.stacks[j]
            # same height, or the same colour on top
            if h1 == h2 or c1 == c2:
                moves.append((i, j))
        return moves

    def after(self, i, j):
        # the stacks after the move (the original is kept)
        (h1, c1), (h2, _) = self.stacks[i], self.stacks[j]
        rest = [s for k, s in enumerate(self.stacks)
                if k not in (i, j)]
        # heights add; the top colour is the moved stack's
```

```
        return rest + [(h1 + h2, c1)]

    def play(self, i, j):
        if (i, j) not in self.legal_moves():
            raise ValueError("illegal move")
        self.stacks = self.after(i, j)

    def is_over(self):
        # the player with no move loses
        return not self.legal_moves()


# BASE QUESTION (optimal): Negamax with memoisation on the
# sorted stack tuple
# Idea: legal_moves lists the ordered pairs that may
# merge, apply builds the next state, and wins(state) is
# true when some move reaches a state where the opponent
# loses.
def wins(stacks):
    # does the player to move win with best play?
    cache = {}

    def solve(state):
        if state in cache:
            return cache[state]
        game = Babylon(state)
        result = False
        for i, j in game.legal_moves():
            nxt = tuple(sorted(game.after(i, j)))
            # a move to a position the opponent loses wins
            if not solve(nxt):
                result = True
                break
        cache[state] = result
        return result

    # sorted stacks: equal positions share one cache entry
    return solve(tuple(sorted(stacks)))


# PICK the winning move itself
def best_move(game):
    # pick a move that leaves the opponent lost, if any
    for i, j in game.legal_moves():
        if not wins(game.after(i, j)):
```

```
            return i, j
    moves = game.legal_moves()
    return moves[0] if moves else None


if __name__ == "__main__":
    # one stack: no moves, so the player to move loses
    assert Babylon([(1, "Y")]).is_over()
    assert not wins([(1, "Y")])
    # two equal stacks: one move ends the game, so we win
```

```
    assert wins([(1, "Y"), (1, "W")])
    g = Babylon([(1, "Y"), (1, "Y"), (1, "W")])
    assert (0, 1) in g.legal_moves()
    # play a whole game with the solver; the loser is
    # stuck
    while not g.is_over():
        g.play(*best_move(g))
    assert len(g.stacks) >= 1
    try:
        Babylon([(1, "Y"), (2, "W")]).play(0, 1)
        raise AssertionError("move should be illegal")
    except ValueError:
        pass
    # the real start has 12 stacks and up to 132 ordered
    # moves
    assert len(Babylon().legal_moves()) > 0
```

NOTES FROM CANDIDATES

This is closer to an implementation/simulation interview than a clever-algorithm problem — the interviewer wants a clean game state class, legal-move enumeration, and a play(move) API.
Represent state as a List[Stack] where Stack = (height, topColor, fullSequence?). Whether you keep the full color sequence depends on whether the interviewer asks about "undo" or "re-color from top".
Add a minimax / negamax for "who wins from this state with optimal play" as the natural follow-up — branching factor is small (at most C(12, 2) = 66 initially, drops fast).
Reported as "too hard to finish coding in 45 minutes"; the interviewer gives hints throughout. Pacing matters.

PREPARATION

Drill clean game-state implementations: e.g. Connect 4, Othello move generators, in 30 minutes.
Practice writing a generic minimax framework so you can plug in this game's legalMoves(state) and terminal(state) quickly.
Watch the merge invariant carefully: "same height OR same top color" — fail mode is forgetting the alternative branch.